

Apache Toree: A Jupyter Kernel for Spark

Marius van Niekerk
Apache Toree Contributor



Apache Toree



A **Jupyter** kernel to connect to **Spark** and create
interactive applications

Ecosystem



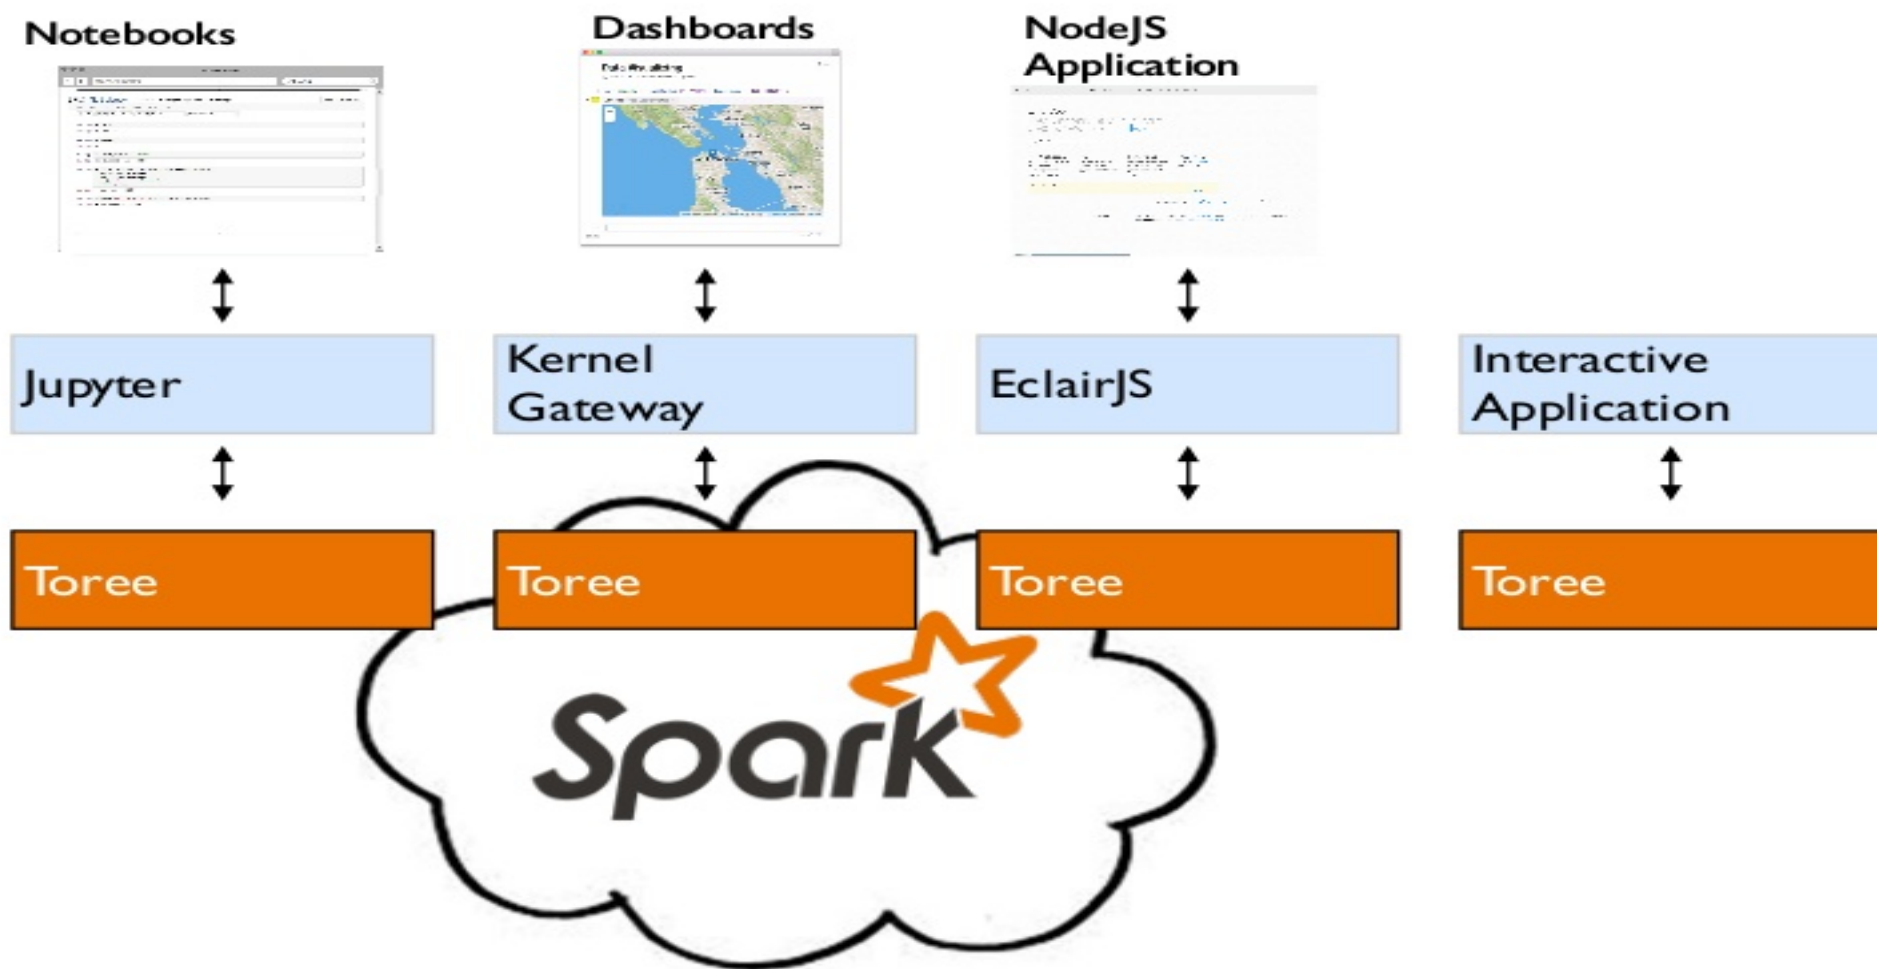
Jupyter



- Open source, interactive data science and scientific computing across over 40 programming languages.
- Torette is an implementation of the Jupyter Kernel Protocol

Apache Toree History

- Started as ibm/spark-kernel in 2013
- Started Apache incubation in late 2015
- First Apache Release is coming soon.



Compatibility

- Toree 0.1.x supports Spark 1.6.x
- Toree 0.2.x supports Spark 2.x

Features

- Kernel languages: Scala, Python, R
- Magics
- Tab completion
- Plugin system

Try it out

- Welcome to Spark with Scala:
 - <https://tmpnb.org>
- Docker
 - `docker run -it --rm -p 8888:8888 \`
`jupyter/all-spark-notebook`
 - This is a very large container ~4gb

In []:

Interrupt

Restart

Restart & Clear Output

Restart & Run All

Reconnect

Change kernel

Examples

- github.com/Lull3rSkat3r/apache-toree-demos
- github.com/apache/incubator-toree/tree/master/etc/examples/notebooks

Extending on top of Toree

- github.com/Brunel-Visualization/Brunel
- github.com/jupyter/declarativewidgets

Other Alternatives

- sparkmagic (livy)
- jupyter-scala
- Apache Zeppelin
- sylon-kernel
- databricks

Help out

Looking for contributors

- Web - toree.apache.org
- Mailing - dev@toree.incubator.apache.org
- Chat - gitter.im/apache/toree

Thank You.

Github: @mariusvniekerk

Twitter: @__mvn__

