

# Bulletproof Jobs @ Scale

## Patterns for Spark Magic

Simeon Simeonov

Founder & CTO, Swoop, @simeons



# Bulletproof jobs

Reliable operations  
over long periods of time  
and fast recovery from failures



petabyte scale data engineering & ML  
run our business

# Spark vs. magic

great magic comes  
from the consistent application  
of smart conventions

(often quite easy to add on top of Spark)

# Problem: job failure response

Your job failed.  
What should you do?

# 90+% of job failures are transient

Does that help?

# If I had a magic wand...

I'd like my system to **auto-retry** the job.

Would that be OK?



Automatic retry works  
only for **idempotent** jobs

(if multiple executions don't change state)

# Idempotence in action

```
DROP TABLE IF EXISTS my_table
```

# Idempotence in Spark

Spark has one **semi**-idempotent way to save data

```
dataset.write.mode("overwrite").save(...)
```

After **standardizing on idempotency**, Swoop  
now auto-recovers from 98+% of failures

(see <http://bit.ly/spark-idempotency> talk)

idempotence is the **second most valuable** property of your Spark jobs

(after correctness)

# Problem: root cause analysis

While ingesting 100M pieces of data, your code throws 1,000 exceptions of 5 different types.


How will you **find & fix all** the problems?

# Exception handling for jobs

- Spark's default: stop after first exception
- Catch & log exceptions
- Use structured logging service

# If I had a magic wand...

- Errors part of job output




Hmm, we'll have to separate data & errors...



# If I had a magic wand...

- Errors part of job output
- Detailed exception info



# If I had a magic wand...

- Errors part of job output
- Detailed exception info
- **Flexible error summaries**

category_type	id_message	cnt	messages
E	number out of range	5	▶ ["E01001: input outside [0, 100]"]
E	null	1	▶ ["java.lang.ArrayIndexOutOfBoundsException: 100"]

# If I had a magic wand...

- Errors part of job output
- Detailed exception info
- Flexible error summaries
- **Drill into errors**

`dataset.errorDetails()`

`dataset.unknownErrorDetails()`

# If I had a magic wand...

- Errors part of job output
- Detailed exception info
- Flexible error summaries
- Drill into errors
- **Errors tied to source data**


A	B	C

source

A	X	Y	Z


# If I had a magic wand...

- Errors part of job output
- Detailed exception info
- Flexible error summaries
- Drill into errors
- Errors tied to source data
- **Errors tied to job flight**



**job flight**

# Smart convention: envelopes





# bulletproof Spark jobs

[View on GitHub](#)

# Spark Records

```
import com.swoop.spark.records._

case class MyData(/* whatever your data is */)

case class MyDataRecord(
  features: Int,                // fast categorization
  data: Option[MyData] = None, // your data
  source: Option[MyInput] = None, // provenance tracking
  flight: Option[String] = None, // related jobs
  → issues: Option[Seq[Issue]] = None // row-level "log file"
) extends Record[MyData, MyInput]
```



# Record builders

```
case class Builder(input: MyInput, jc: JobContext)
  extends RecordBuilder[...](input, jc) {

  def buildData = { /* your existing code here */ }

  def dataRecord(data: MyData, issues: Seq[Issue]) =
    MyDataRecord(...)

  def errorRecord(issues: Seq[Issue]) =
    MyDataRecord(...)
}
```

# Transparent data pattern

```
// Get the data, skip all error records  
records.recordData  
    .createOrReplaceGlobalView("my_data")
```

any data persistence. near-zero overhead.

# Demo: root cause analysis

# Smart conventions & 15 lines of code

- Bulletproof job execution
- Automatic row-level structured logging
- Automatic metric collection & data quality checks
- Identifiable errors & lightning-fast RCA
- Fully customizable yet fully transparent
- Near-zero overhead

any complex data production using  
custom code would benefit from the  
Spark Records pattern

(use the schema pattern if you can't use the code)



# Swoop's Spark Magic Toolkit

- Spark Records (bundled with Spark Metrics)
- Idempotent I/O to files systems & databases
  - Resilient Partitioned Tables
  - Semi-permanent tables
- Parallel DAG execution in Spark notebooks
- Smart Data Warehouse

# Swoop's magic has given us

5-10x faster exploration & ML

15x more reliable job execution

10-100x faster root cause analysis

# Get Swoop's Magic & Share Yours

<http://bit.ly/spark-records>

Email [spark@swoop.com](mailto:spark@swoop.com) for more  
or tweet [@simeons](https://twitter.com/simeons)



# Thank You.

Get Swoop's magic & share yours.

<http://bit.ly/spark-records>

[spark@swoop.com](mailto:spark@swoop.com)

