



Keeping Spark on Track: Productionizing Spark for ETL

Kyle Pistor

kyle@databricks.com

Miklos Christine

mwc@databricks.com



\$ whoami

Kyle Pistor

- SA @ Databricks
- 100s of Customers
- Focus on ETL and big data warehousing using Apache Spark
- BS/MS - EE

Miklos Christine

- SA @ Databricks!
- Previously: Systems Engineer @ Cloudera
- Deep Knowledge of Big Data Stack
- Apache Spark Enthusiast

Agenda

- ① **ETL:** Why ETL?
- ② **Schemas**
- ③ **Metadata:** Best Practices FS Metadata
- ④ **Performance:** Tips & Tricks
- ⑤ **Error Handling**

Agenda

- 1 **ETL:** Why ETL?
- 2 **Schemas**
- 3 **Metadata:** Best Practices FS Metadata
- 4 **Performance:** Tips & Tricks
- 5 **Error Handling**

Why ETL?

Goal

- Transform raw files into more efficient binary format

Benefits

- Performance improvements, statistics
- Space savings
- Standard APIs to different sources (CSV, JSON, etc.)
- More robust queries

Agenda

① **ETL:** Why ETL?

② **Schemas**

③ **Metadata:** Best Practices FS Metadata

④ **Performance:** Tips & Tricks

⑤ **Error Handling**

Schema Handling

Common Raw Data

- Delimited (CSV, TSV, etc)
- JSON

Infer Schema?

- Easy and quick

Specify Schema

- Faster
- More robust

JSON Record #1

```
{  
  "time": 1486166400,  
  "host": "my.app.com",  
  "eventType": "test",  
  "event": {  
    "message": "something happened",  
    "type": "INFO"  
  }  
}
```


JSON Record #1

```
{  
  "time": 1486166400,  
  "host": "my.app.com",  
  "eventType": "test",  
  "event": {  
    "message": "something happened",  
    "type": "INFO"  
  }  
}
```

```
spark.read.json("record1.json").schema
```

```
root  
|-- event: struct (nullable = true)  
|   |-- message: string (nullable = true)  
|   |-- type: string (nullable = true)  
|-- host: string (nullable = true)  
|-- source: string (nullable = true)  
|-- time: long (nullable = true)
```

JSON Record #2

```
{  
  "time": 1486167800,  
  "host": "my.app.com",  
  "eventType": "test",  
  "event": {  
    "message": "Something else happened",  
    "type": "INFO",  
    "ip" : "59.16.12.0",  
    "browser" : "chrome",  
    "os" : "mac",  
    "country" : "us"  
  }  
}
```

JSON Record #2

```
{  
  "time": 1486166400,  
  "host": "my.app.com",  
  "eventType": "test",  
  "event": {  
    "message": "something happened",  
    "type": "INFO"  
  }  
}  
  
→  
  
{  
  "time": 1486167800,  
  "host": "my.app.com",  
  "eventType": "test",  
  "event": {  
    "message": "Something else happened",  
    "type": "INFO",  
    {  
      "ip" : "59.16.12.0",  
      "browser" : "chrome",  
      "os" : "mac",  
      "country" : "us"  
    }  
  }  
}
```

JSON Record #1 & #2

```
{  
  "time": 1486167800,  
  "host": "my.app.com",  
  "eventType": "test",  
  "event": {  
    "message": "Something else happened",  
    "type": "INFO",  
    "ip" : "59.16.12.0",  
    "browser" : "chrome",  
    "os" : "mac",  
    "country" : "us"  
  }  
}
```

spark.read.json("record*.json").schema

```
root  
|-- event: struct (nullable = true)  
|   |-- message: string (nullable = true)  
|   |-- type: string (nullable = true)  
|   |-- ip: string (nullable = true)  
|   |-- browser: string (nullable = true)  
|   |-- os: string (nullable = true)  
|   |-- country: string (nullable = true)  
|-- host: string (nullable = true)  
|-- source: string (nullable = true)  
|-- time: long (nullable = true)
```

JSON Generic Specified Schema

Print Schema

```
spark.read.json("record*.json").printSchema
```



StructType => MapType

```
customSchema = StructType([  
    StructField("time", TimestampType(), True),  
    StructField("host", StringType(), True),  
    StructField("source", StringType(), True),  
    StructField("event",  
        MapType(StringType(), StringType()))  
])
```



Specify Schema

```
spark.read.schema(customSchema).json("record*.json")
```

Specify Schemas!

Faster

- No scan to infer the schema

More Flexible

- Easily handle evolving or variable length fields

More Robust

- Handle type errors on ETL vs query

Agenda

- ① **ETL:** Why ETL?
- ② **Schemas**
- ③ **Metadata:** Best Practices FS Metadata
- ④ **Performance:** Tips & Tricks
- ⑤ **Error Handling**

Filesystem Metadata Management

Common Source FS Metadata

- Partitioned by arrival time (“path/yyyy/MM/dd/HH”)
- Many small files (kBs - 10s MB)
- Too little data per partition

```
...myRawData/2017/01/29/01  
...myRawData/2017/01/29/02
```

Backfill - Naive Approach

Backfill

- Use only wildcards

```
df = spark.read.json("myRawData/2016/**/**/**")  
:  
:  
:  
df.write.partitionBy($"date").parquet("myParquetData")
```

Why this is a poor approach

- Spark (currently) is not aware of the existing partitions (yyyy/MM/dd/HH)
- Expensive full shuffle

Backfill - Naive Approach

Backfill

- Use only wildcards

```
df = spark.read.json("myRawData/2016/**/**/**")  
:  
:  
:  
df.write.partitionBy($"date").parquet("myParquetData")
```

Why this is a poor approach

- Spark (currently) is not aware of the existing partitions (yyyy/MM/dd/HH)
- Expensive full shuffle

Backfill - Scalable Approach

List of Paths

- Create a list of paths to backfill
- Use FS ls commands

Iterate over list to backfill

- Backfill each path

Operate in parallel

- Use Multiple Threads

Scala .par

Python multithreading

```
def convertToParquet (path:String) = {  
  val df = spark.read.json(path)  
  .  
  .  
  df.coalesce(20).write.mode("overwrite").save()  
}
```

```
dirList.foreach(x => convertToParquet(x))
```

```
dirList.par.foreach(x => convertToParquet(x))
```

Source Path Considerations

Directory

- Specify Date instead of Year, Month, Day

Block Sizes

- Read: Parquet Larger is Okay
- Write: 500MB-1GB

Blocks / GBs per Partition

- Beware of Over Partitioning
- ~30GB per Partition

Agenda

- ① **ETL:** Why ETL?
- ② **Schemas**
- ③ **Metadata:** Best Practices FS Metadata
- ④ **Performance:** Tips & Tricks
- ⑤ **Error Handling**

Performance Optimizations

- Understand how Spark interprets Null Values
 - `nullValue`: specifies a string that indicates a null value, any fields matching this string will be set as nulls in the DataFrame

```
df = spark.read.options(header='true', inferschema='true', \
                        nullValue='\\N').csv('people.csv')
```



- Spark can understand it's own null data type
 - Users must translate their null types to Spark's native null type

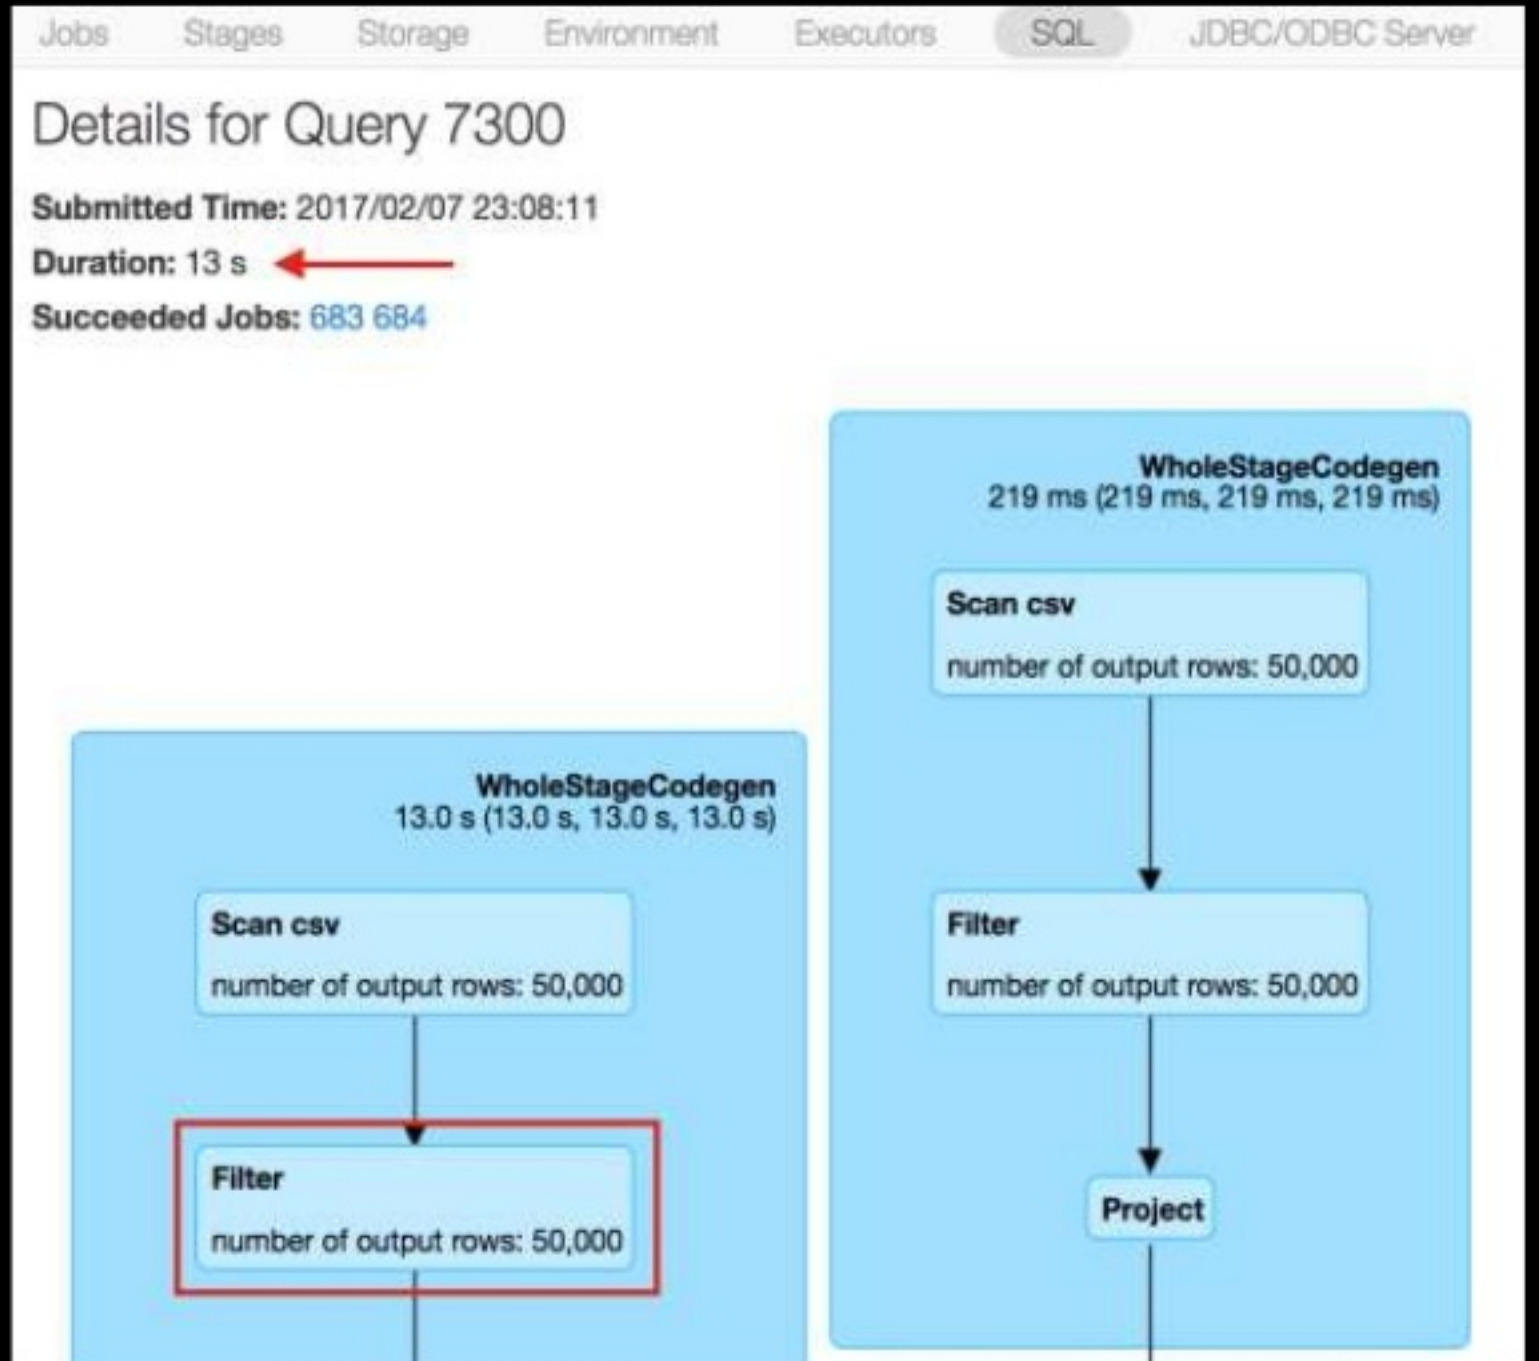
Test Data

| <i>id</i> | <i>name</i> |
|-----------|-------------|
| | foo |
| 2 | bar |
| | bar |
| 11 | foo |
| | bar |

| <i>id</i> | <i>name</i> |
|-----------|-------------|
| null | bar |
| null | bar |
| 3 | foo |
| 15 | foo |
| 2 | foo |

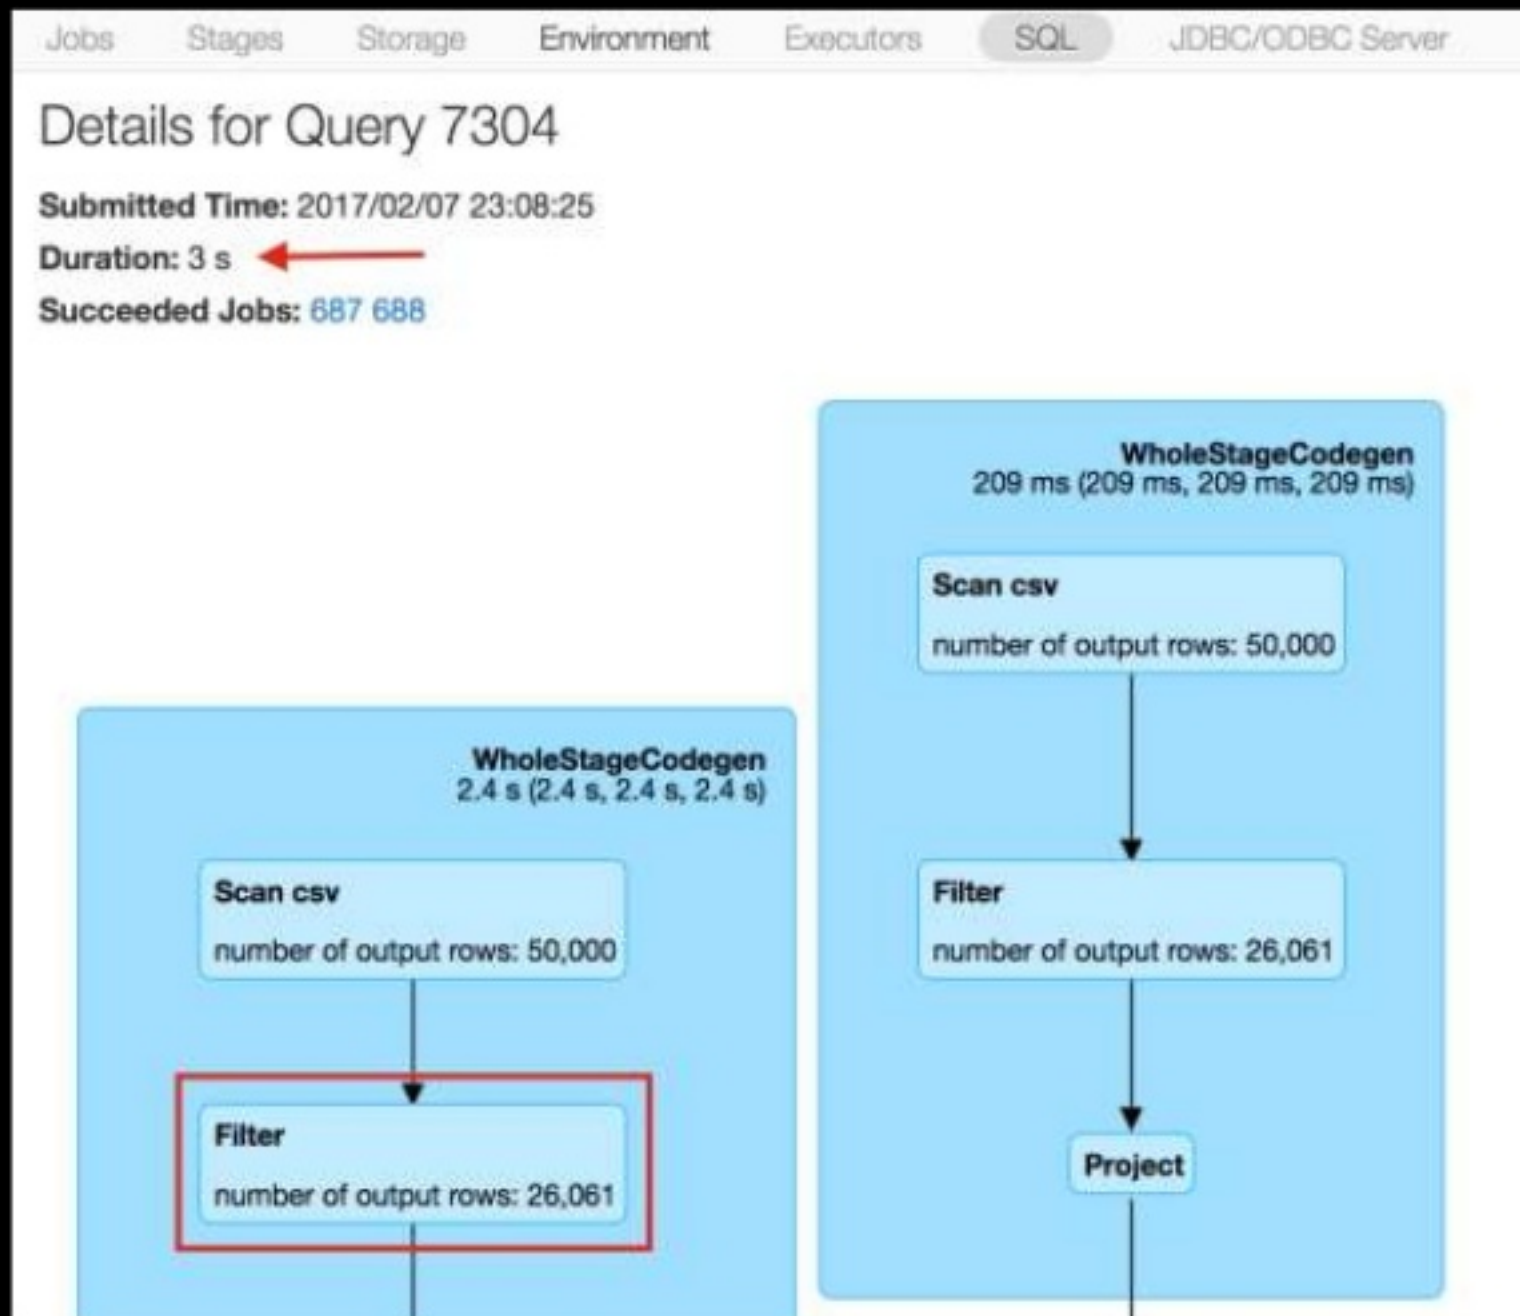
Performance: Join Key

- Spark's WholeStageCodeGen will attempt to filter null values within the join key
- Ex: Empty String Test Data



Performance: Join Key

- The run time decreased from 13 seconds to 3 seconds
- Ex: Null Values for Empty Strings



DataFrame Joins

```
df = spark.read.options(header='true').csv('/mnt/mwc/csv_join_empty')  
  
df_duplicate = df.join(df, df['id'] == df['id'])  
  
df_duplicate.printSchema()
```

(2) Spark Jobs

root

```
|-- id: string (nullable = true)  
|-- name1: string (nullable = true)  
|-- id: string (nullable = true)  
|-- name2: string (nullable = true)
```

DataFrame Joins

```
df = spark.read.options(header='true').csv('/mnt/mwc/csv_join_empty')  
df_id = df.join(df, 'id')  
df_id.printSchema()
```

(2) Spark Jobs

```
root  
|-- id: string (nullable = true)  
|-- name1: string (nullable = true)  
|-- name2: string (nullable = true)
```

Optimizing Apache Spark SQL Joins:

<https://spark-summit.org/east-2017/events/optimizing-apache-spark-sql-joins/>

Agenda

- ① **ETL:** Why ETL?
- ② **Schemas**
- ③ **Metadata:** Best Practices FS Metadata
- ④ **Performance:** Tips & Tricks
- ⑤ **Error Handling**

Error Handling: Corrupt Records

```
display(json_df)
```

| <code>_corrupt_records</code> | <code>eventType</code> | <code>host</code> | <code>time</code> |
|--------------------------------------------------------------------|------------------------|--------------------------|-------------------------|
| <code>null</code> | <code>test1</code> | <code>my.app.com</code> | <code>1486166400</code> |
| <code>null</code> | <code>test2</code> | <code>my.app.com</code> | <code>1486166402</code> |
| <code>null</code> | <code>test3</code> | <code>my.app2.com</code> | <code>1486166404</code> |
| <code>{"time": 1486166406, "host": "my2.app.com", "event_"}</code> | <code>null</code> | <code>null</code> | <code>null</code> |
| <code>null</code> | <code>test5</code> | <code>my.app2.com</code> | <code>1486166408</code> |

Error Handling: UDFs

```
Py4JJavaError: An error occurred while calling o270.showString.
```

```
: org.apache.spark.SparkException: Job aborted due to stage failure: Task 0 in stage 59.0 failed 4 times, most recent failure: Lost task 0.3 in stage 59.0 (TID 607, 172.128.245.47, executor 1):  
org.apache.spark.api.python.PythonException: Traceback (most recent call last):
```

```
...
```

```
File "/databricks/spark/python/pyspark/worker.py", line 92, in <lambda>
```

```
    mapper = lambda a: udf(*a)
```

```
...
```

```
File "<ipython-input-10-b7bf56c9b155>", line 7, in add_minutes
```

```
AttributeError: type object 'datetime.datetime' has no attribute 'timedelta'
```

```
    at org.apache.spark.api.python.PythonRunner$$anon$1.read(PythonRDD.scala:193)
```

Error Handling: UDFs

- Spark UDF Example

```
from datetime import date

from datetime import datetime, timedelta

def add_minutes(start_date, minutes_to_add):

    return datetime.combine(b, datetime.min.time()) + timedelta(minutes=long(minutes_to_add))
```

- Best Practices

- Verify the input data types from the DataFrame
- Sample data to verify the UDF
- Add test cases

Error Handling: InputFiles

- Identify the records input source file

```
from pyspark.sql.functions import *  
  
df = df.withColumn("fname", input_file_name())  
  
display(spark.sql("select *, input_file_name() from companies"))
```

| city | zip | fname |
|-------------|-------|--------------------------------------------------------------------------------------------------|
| LITTLETON | 80127 | dbfs:/user/hive/warehouse/companies/part-r-00000-8cf2a23f-f3b5-4378-b7b0-72913fbb7414.gz.parquet |
| BOSTON | 02110 | dbfs:/user/hive/warehouse/companies/part-r-00000-8cf2a23f-f3b5-4378-b7b0-72913fbb7414.gz.parquet |
| NEW YORK | 10019 | dbfs:/user/hive/warehouse/companies/part-r-00000-8cf2a23f-f3b5-4378-b7b0-72913fbb7414.gz.parquet |
| SANTA CLARA | 95051 | dbfs:/user/hive/warehouse/companies/part-r-00000-8cf2a23f-f3b5-4378-b7b0-72913fbb7414.gz.parquet |



Thanks

Enjoy The Snow

<https://databricks.com/try-databricks>

