

Monitoring the Dynamic Resource Usage of Scala and Python Spark Jobs in Yarn

Ed Barnes, Ruslan Vaulin and Chris McCubbin
Sqrri Data





DATA SOURCES

SECURITY DATA



Alerts



Threat Intelligence



NETWORK DATA



Proxy



NetFlow



ENDPOINT/IDENTITY DATA



HR

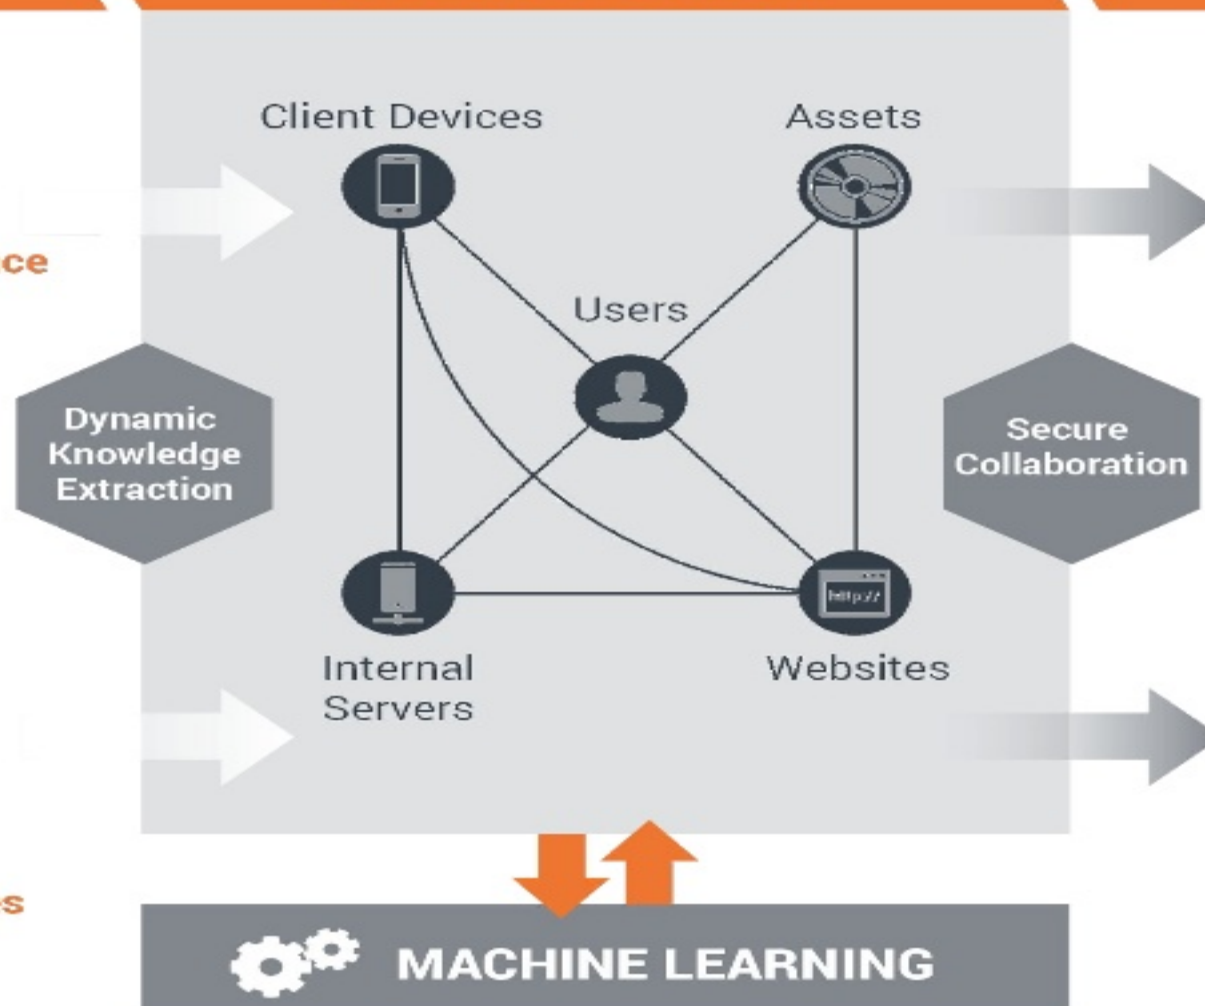


Authentication



Processes

BEHAVIOR GRAPH



THREAT HUNTING

SEARCH



EXPLORATION



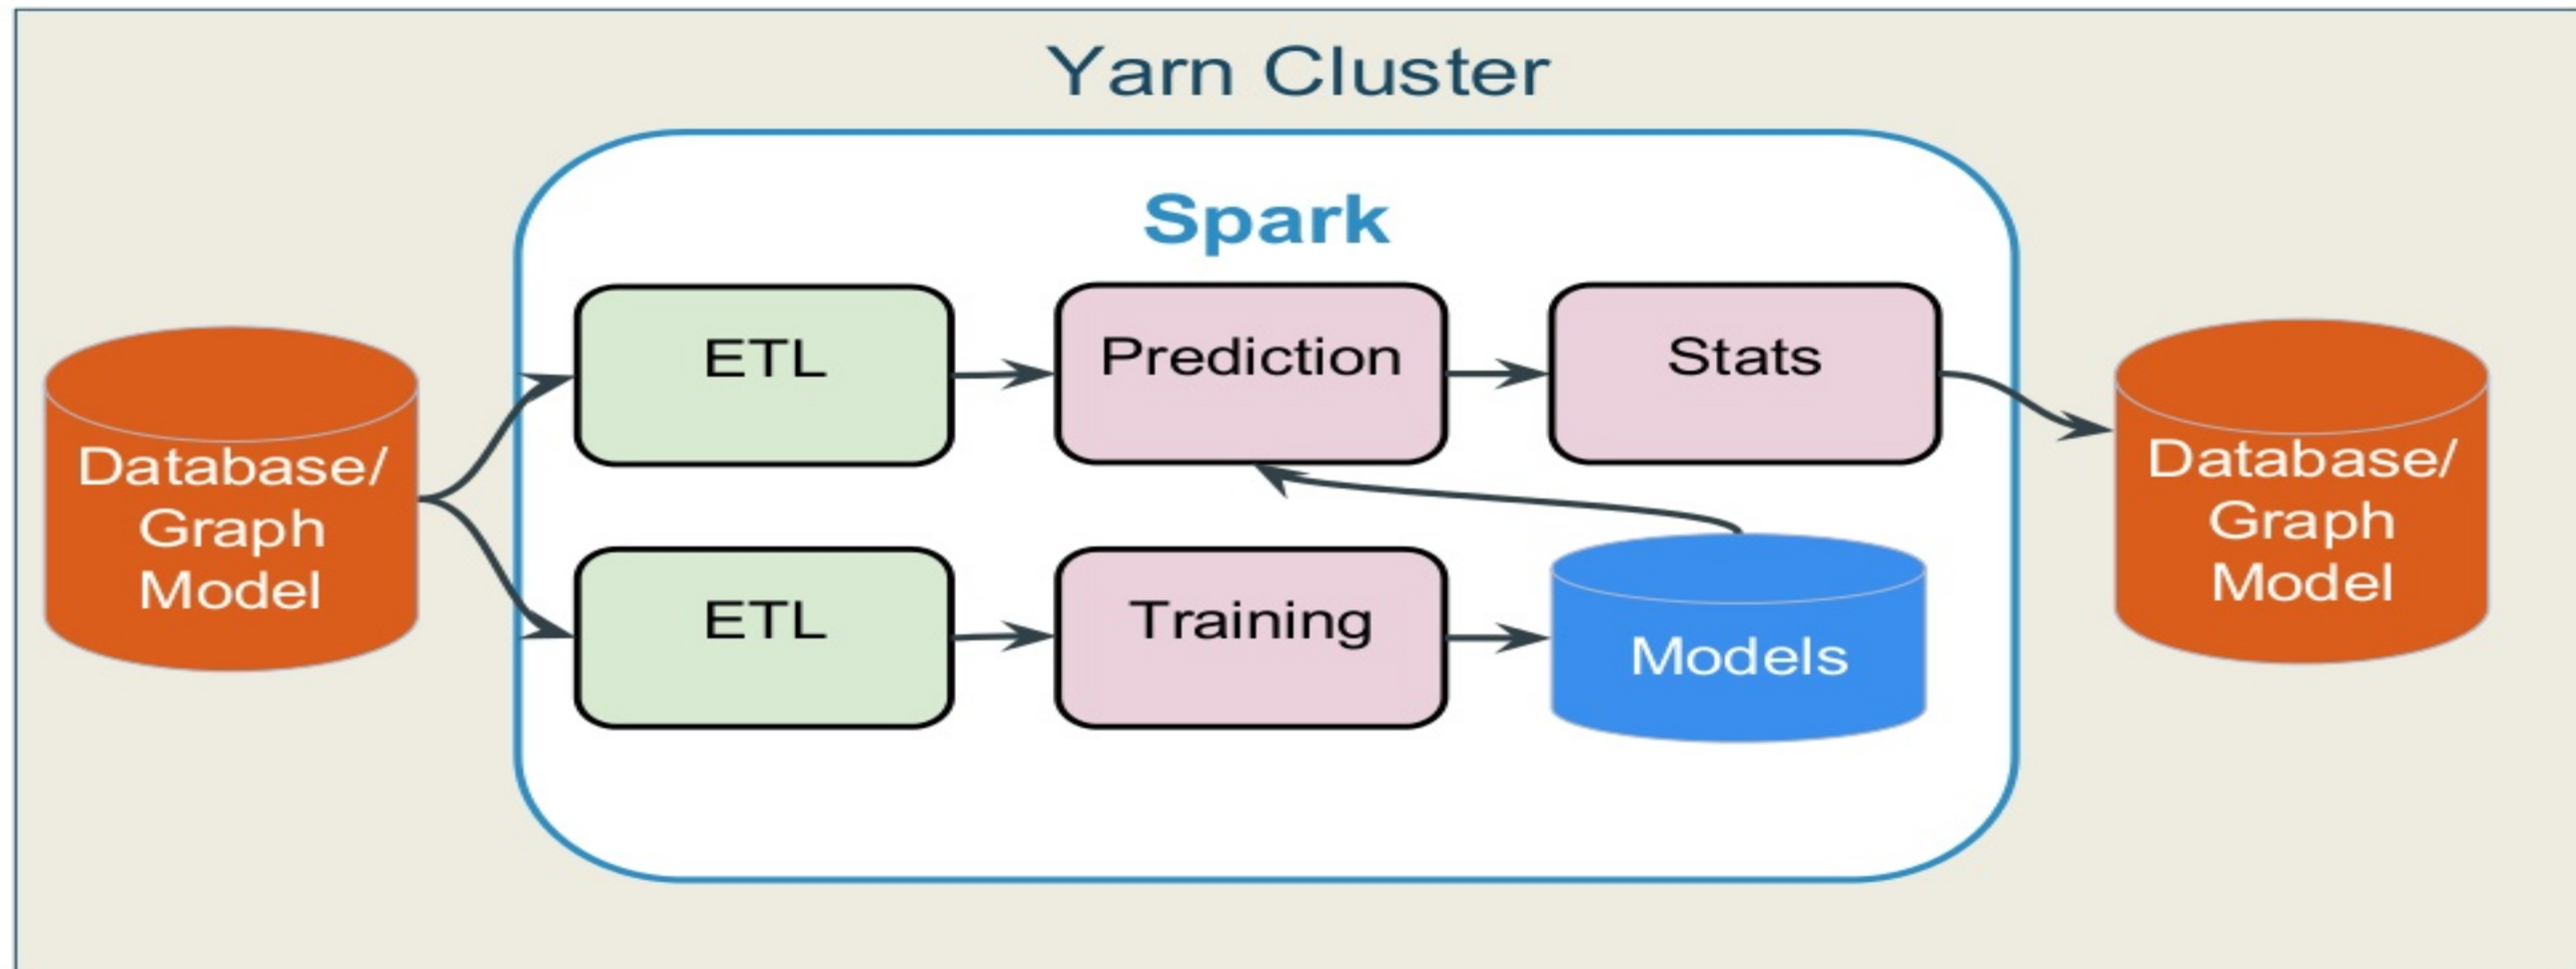
REPORTS



ANOMALIES



ML Application Workflow



Python

Scala

Taking Spark Applications into Production

- Requires execution framework
- Scalable, Robust, Tested
- Test at scale
- Many issues show up only at scale
 - Performance
 - Memory requirements
 - Failures
 - Scaling
- Debugging distributed applications is really hard!

Spark UI: job level

 1.5.1

Jobs

Stages

Storage

Environment

Executors

SPARK: DGA - training for att dn... application UI

Spark Jobs (?)

Total Uptime: 1.6 h
Scheduling Mode: FIFO
Completed Jobs: 10
Failed Jobs: 1

▶ Event Timeline

Completed Jobs (10)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
9	reduce at /data7/yarn/local/usercache/yarn/appcache/application_1485380275144_0028/container...	2017/01/25 23:06:31	9.6 min	1/1	40/40
8	reduce at /data7/yarn/local/usercache/yarn/appcache/application_1485380275144_0028/container...	2017/01/25 22:58:14	8.0 min	2/2	80/80
7	count at /data7/yarn/local/usercache/yarn/appcache/application_1485380275144_0028/container...	2017/01/25 22:50:39	7.6 min	1/1	40/40
6	collect at /data7/yarn/local/usercache/yarn/appcache/application_1485380275144_0028/container...	2017/01/25 22:50:21	18 s	1/1	40/40
5	collect at /data7/yarn/local/usercache/yarn/appcache/application_1485380275144_0028/container...	2017/01/25 22:31:21	19 min	2/2	80/80
4	collect at /data7/yarn/local/usercache/yarn/appcache/application_1485380275144_0028/container...	2017/01/25 22:24:50	6.5 min	1/1	40/40
3	count at /data7/yarn/local/usercache/yarn/appcache/application_1485380275144_0028/container...	2017/01/25 22:18:18	6.5 min	1/1	40/40
2	stats at /data7/yarn/local/usercache/yarn/appcache/application_1485380275144_0028/container_1...	2017/01/25 22:11:30	6.8 min	1/1	40/40
1	runJob at PythonRDD.scala:393	2017/01/25 22:11:25	5 s	1/1	1/1
0	take at SerDeUtil.scala:201	2017/01/25 22:11:22	3 s	1/1	1/1

Failed Jobs (1)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
10	reduce at /data7/yarn/local/usercache/yarn/appcache/application_1485380275144_0028/container...	2017/01/25 23:16:24	30 min	0/1 (1 failed)	39/40 (4 failed)

Spark UI: task level



1.6.1

Jobs

Stages

Storage

Environment

Executors

SPARK: DGA - training for att dn... application UI

Details for Stage 12 (Attempt 0)

Total Time Across All Tasks: 2.4 h

Locality Level Summary: Node local: 2; Process local: 41

Input Size / Records: 2.1 GB / 50932

▼ DAG Visualization



Tasks

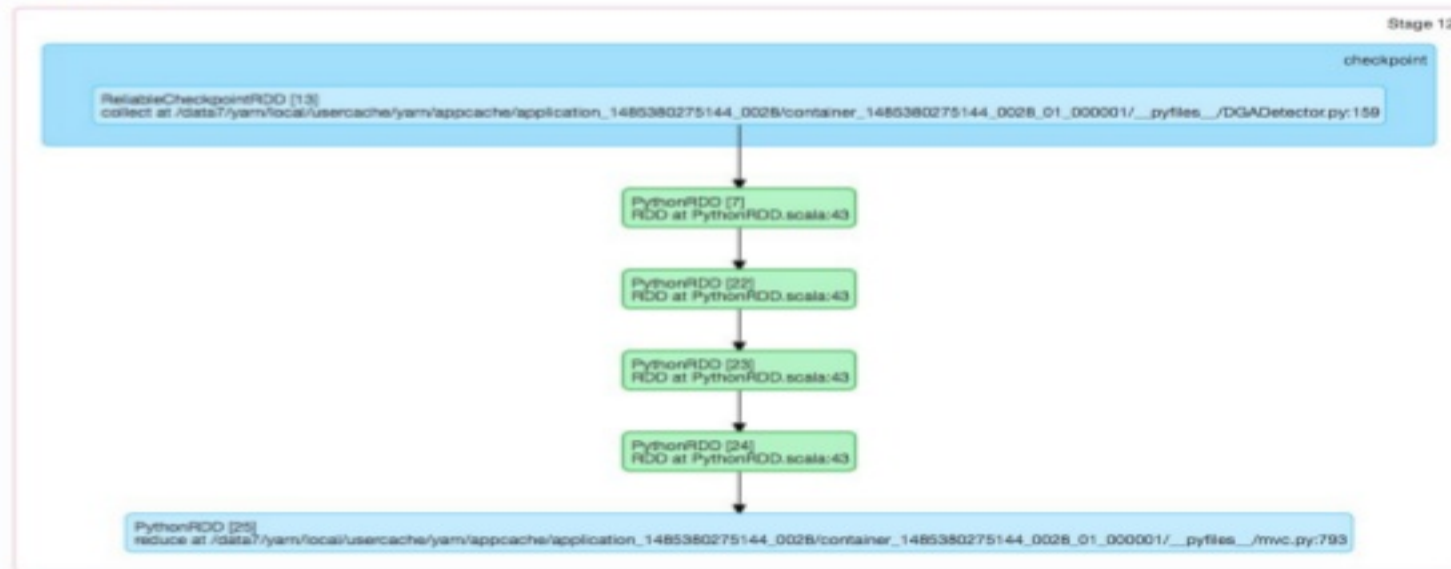
Index ▲	ID	Attempt	Status	Locality Level	Executor ID / Host	Launch Time	Duration	GC Time	Input Size / Records	Errors
0	407	0	SUCCESS	PROCESS_LOCAL	8 /	2017/01/25	8.2 min	90	104.9	
32	441	0	FAILED	PROCESS_LOCAL	6 / n12.sqrrl- lab.net	2017/01/25 23:25:08	21 min	2 s	362.9 MB (memory) / 8211	java.lang.OutOfMemoryError: Java heap space +details

Spark UI: task level

Details for Stage 12 (Attempt 0)

Total Time Across All Tasks: 2.4 h
Locality Level Summary: Node local: 2; Process local: 41
Input Size / Records: 2.1 GB / 50932

▼ DAG Visualization



Tasks

Index ▲	ID	Attempt	Status	Locality Level	Executor ID / Host	Launch Time	Duration	GC Time	Input Size / Records	Errors
0	407	0	SUCCESS	PROCESS_LOCAL	8 /	2017/01/25	8.2 min	90	104.9	
32	441	0	FAILED	PROCESS_LOCAL	6 / n12.sqrrl-lab.net	2017/01/25 23:25:08	21 min	2 s	362.9 MB (memory) / 8211	java.lang.OutOfMemoryError: Java heap space +details

OOM
Exception

Case Study: Py4J Issue

- Testing engineer - “Your code blows up with OOM when processing X amount of data!”


```
17/01/30 12:01:18 INFO ContextCleaner: Cleaned accumulator 2
Exception in thread "Thread-7" java.lang.OutOfMemoryError: Java heap space
    at java.util.Arrays.copyOf(Arrays.java:2367)
    at java.lang.AbstractStringBuilder.expandCapacity(AbstractStringBuilder.java:130)
    at java.lang.AbstractStringBuilder.ensureCapacityInternal(AbstractStringBuilder.java:114)
    at java.lang.AbstractStringBuilder.append(AbstractStringBuilder.java:587)
    at java.lang.StringBuilder.append(StringBuilder.java:214)
    at py4j.Protocol.getOutputCommand(Protocol.java:322)
    at py4j.commands.CallCommand.execute(CallCommand.java:82)
    at py4j.GatewayConnection.run(GatewayConnection.java:209)
    at java.lang.Thread.run(Thread.java:745)
```

- Why?!!!!
- Nothing obvious in Spark UI!

Case Study: Py4J Issue

- Testing engineer - “Your code blows up with OOM when processing X amount of data!”

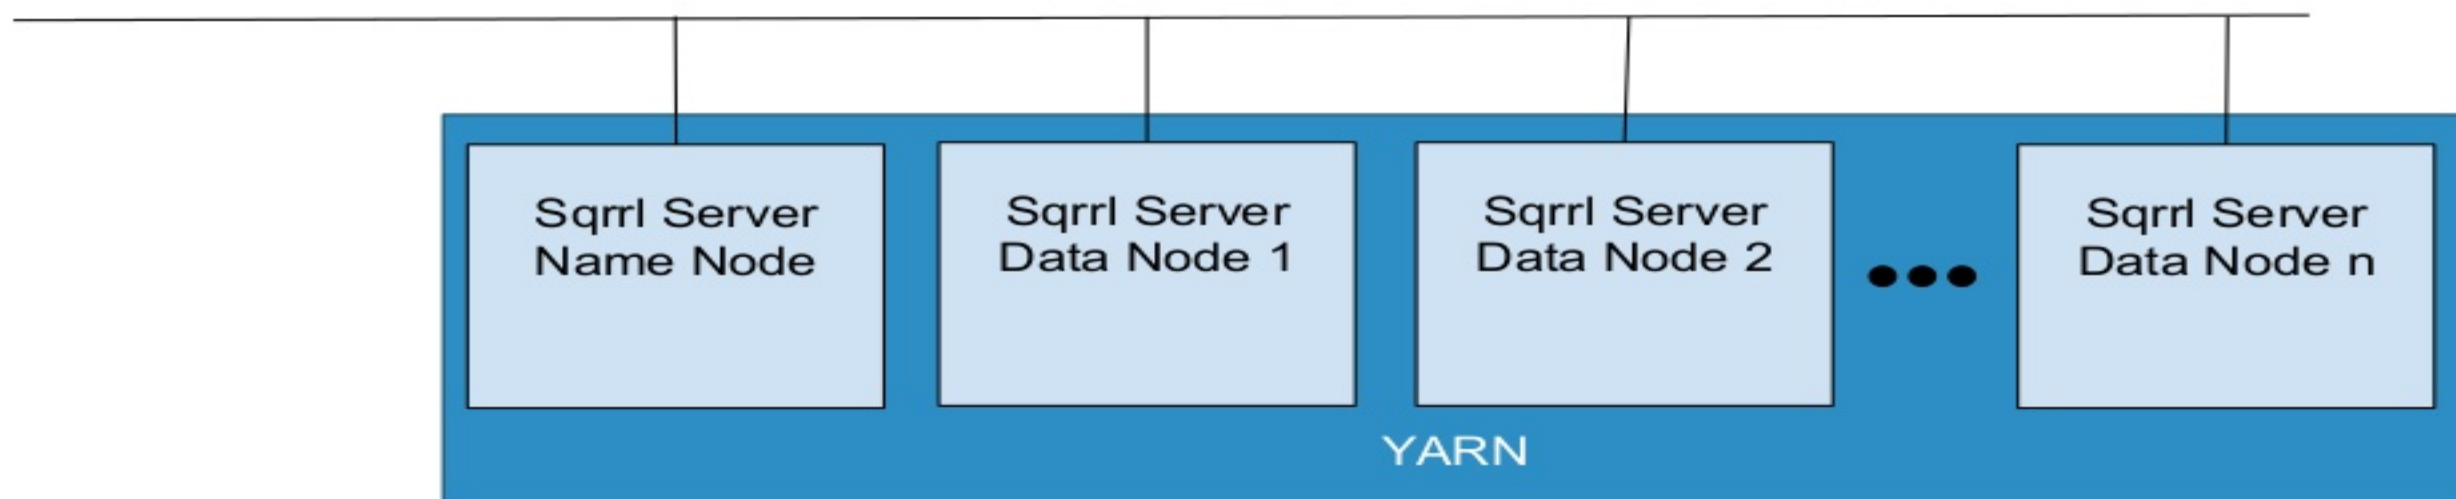
```
17/01/30 12:01:18 INFO ContextCleaner: Cleaned accumulator 2  
Exception in thread "Thread-7" java.lang.OutOfMemoryError: Java heap space  
    at java.util.Arrays.copyOf(Arrays.java:2367)  
    at java.lang.AbstractStringBuilder.expandCapacity(AbstractStringBuilder.java:130)  
    at java.lang.AbstractStringBuilder.ensureCapacityInternal(AbstractStringBuilder.java:114)  
    at java.lang.AbstractStringBuilder.append(AbstractStringBuilder.java:587)  
    at java.lang.StringBuilder.append(StringBuilder.java:214)  
    at py4j.Protocol.getOutputCommand(Protocol.java:322)  
    at py4j.commands.CallCommand.execute(CallCommand.java:82)  
    at py4j.GatewayConnection.run(GatewayConnection.java:209)  
    at java.lang.Thread.run(Thread.java:745)
```



- Why?!!!!
- Nothing obvious in Spark UI!

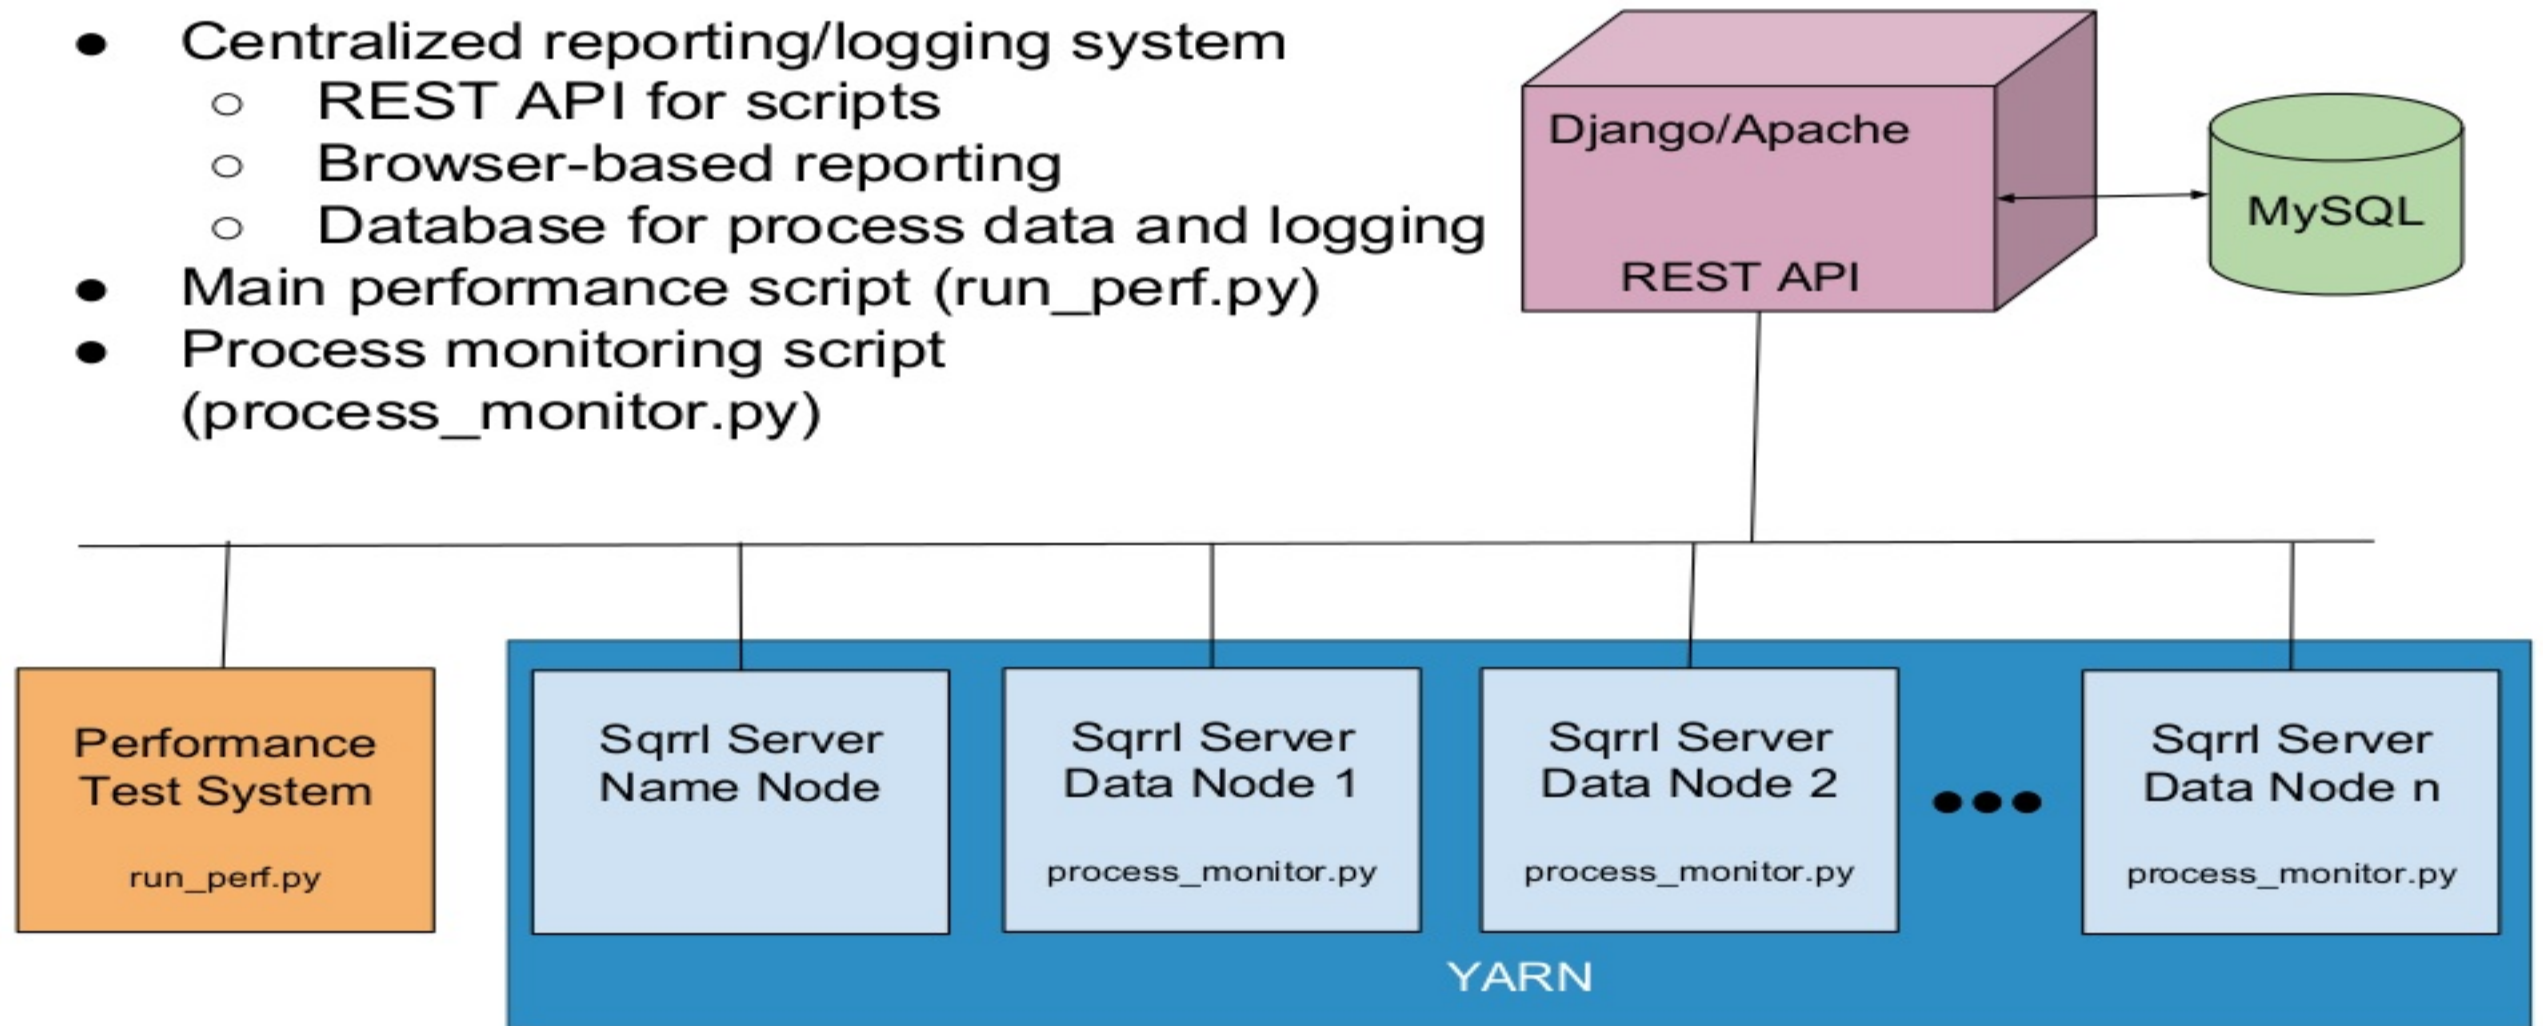
Tooling Approach Requirements

- Cluster-wide process monitoring
- Per-node, per-process statistics
- Identification of high CPU, memory usage
- Record process hierarchy/timing of Spark jobs under YARN
- Characterize scaling behaviors for production/releases
- Integrate with internal processes/harnesses for general development and test use



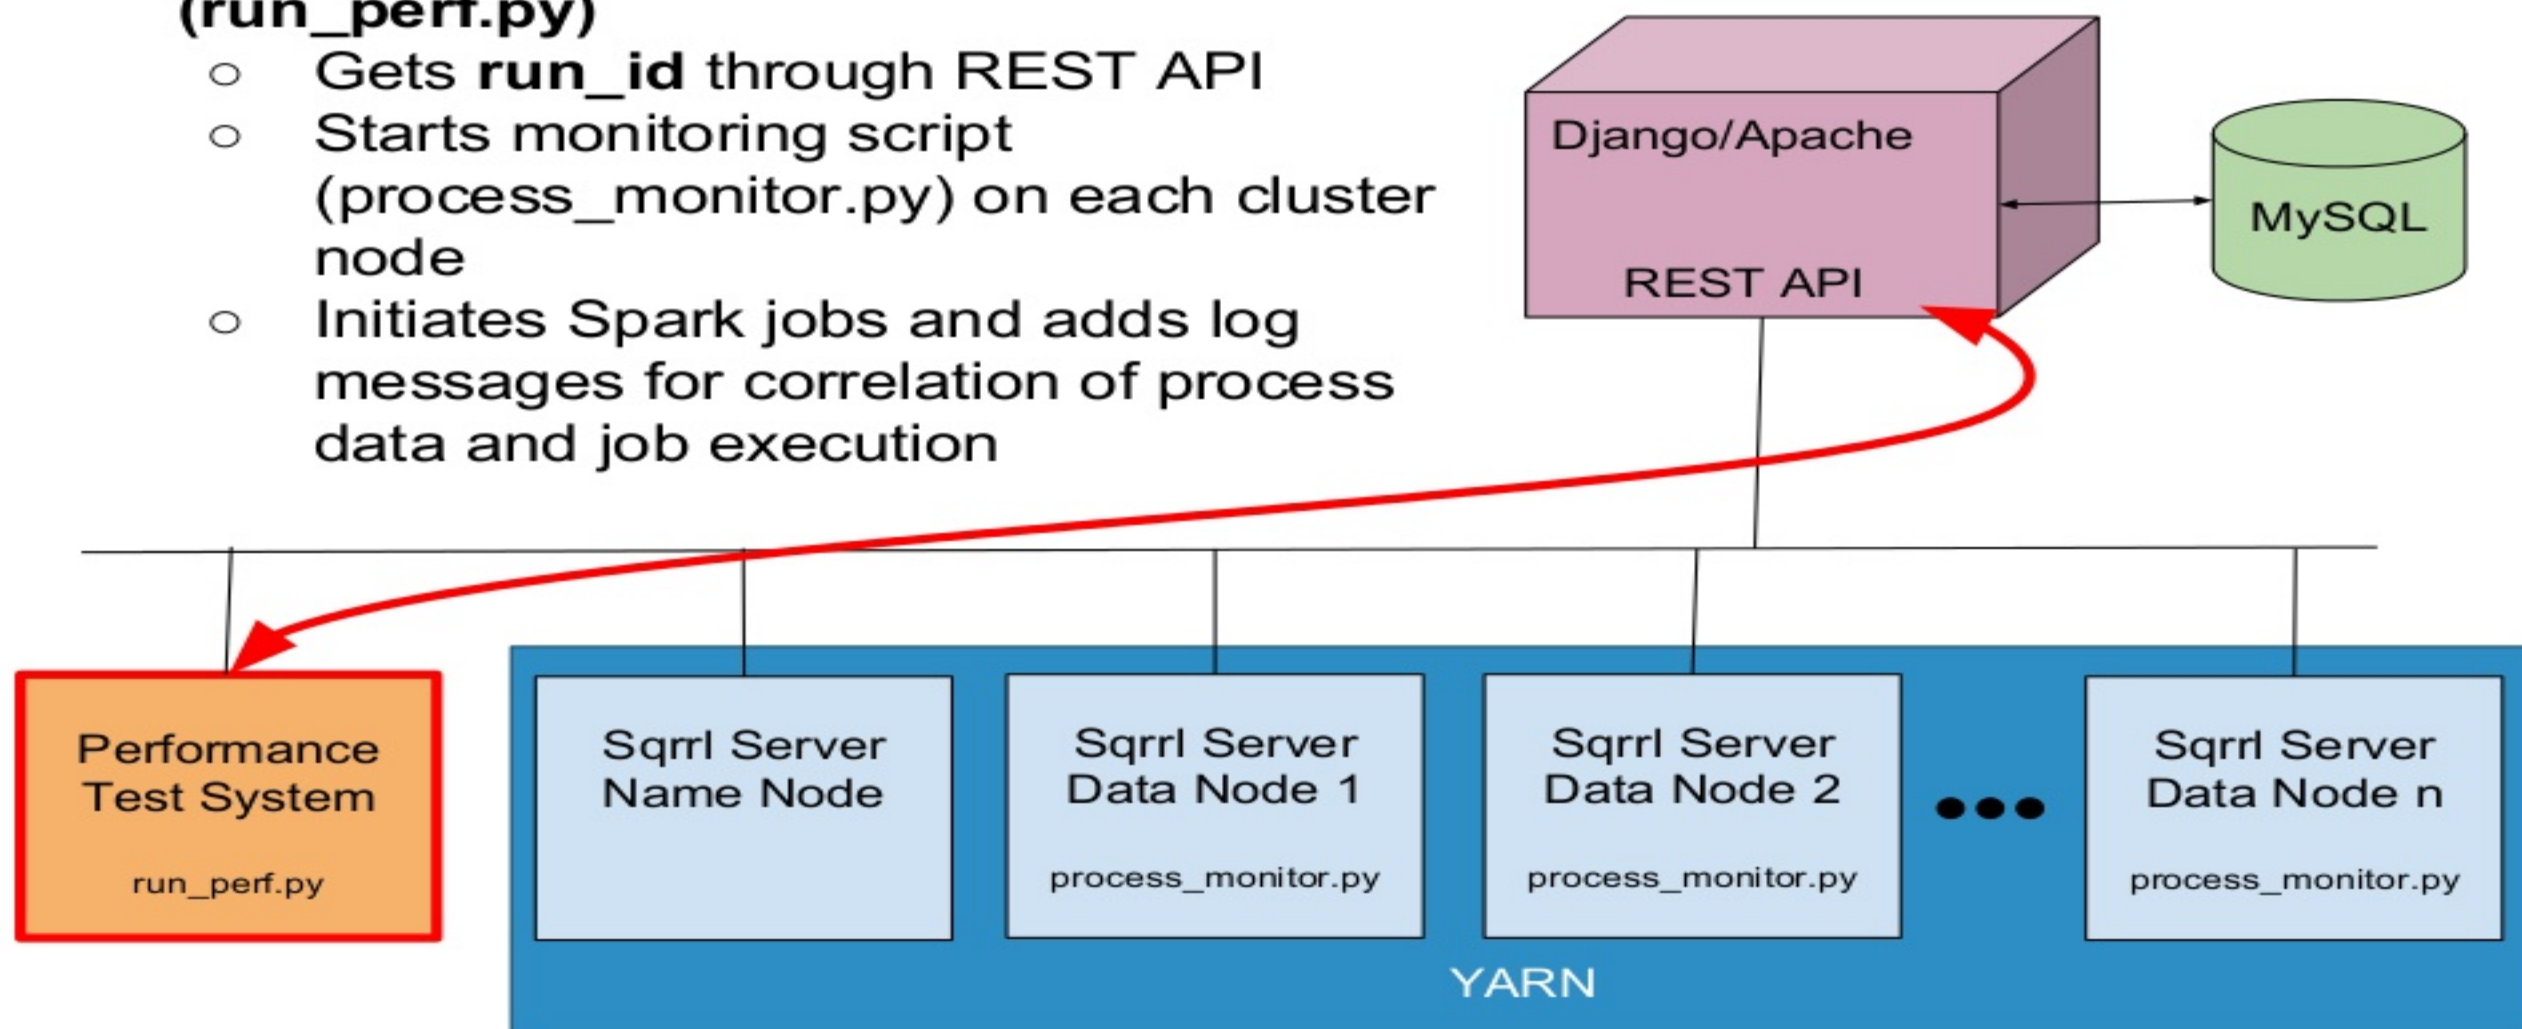
Overall System Design

- Centralized reporting/logging system
 - REST API for scripts
 - Browser-based reporting
 - Database for process data and logging
- Main performance script (run_perf.py)
- Process monitoring script (process_monitor.py)



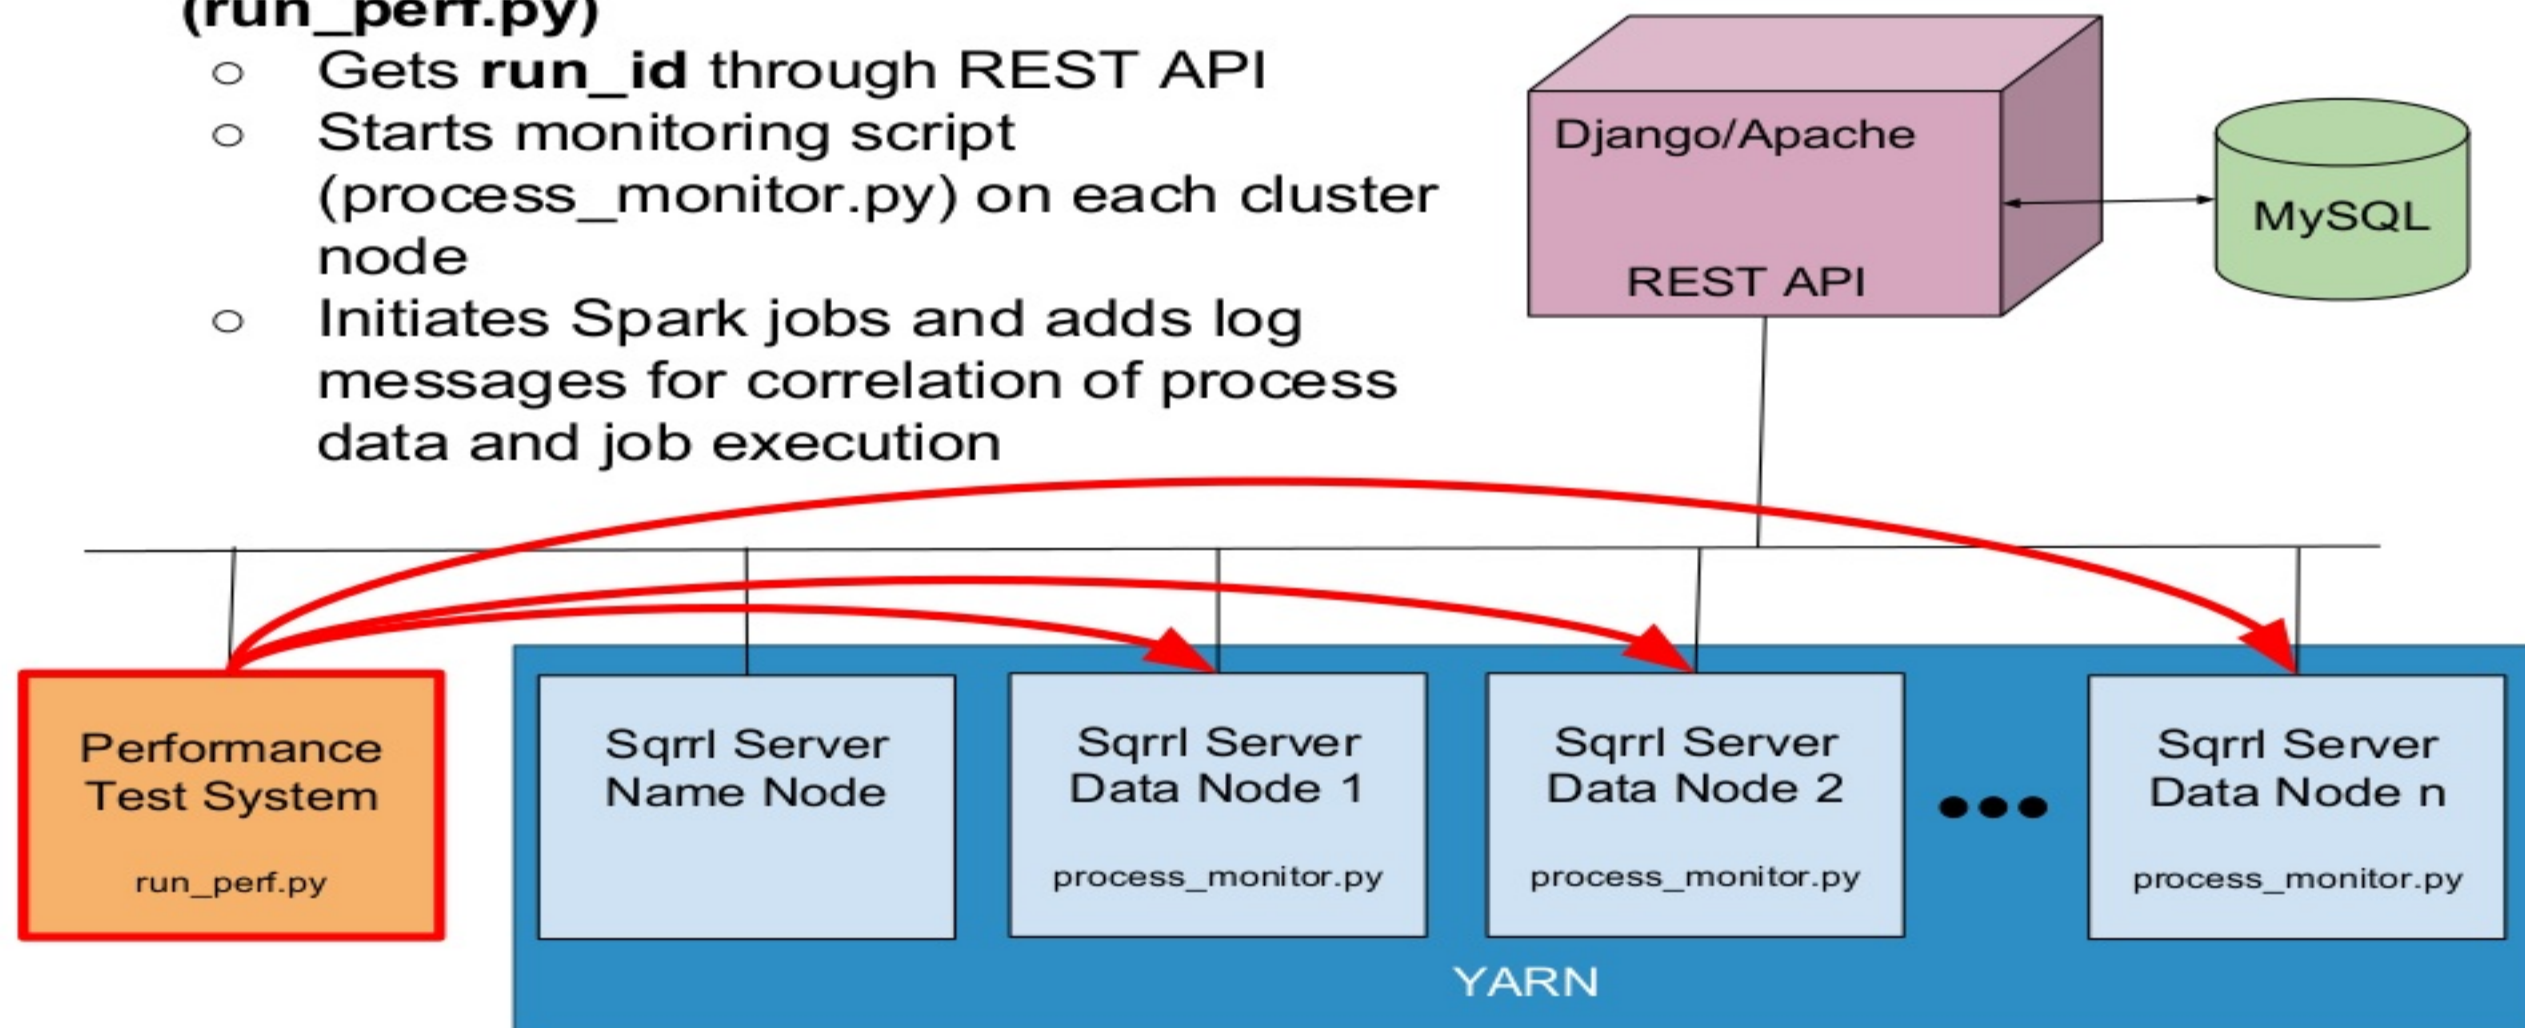
Scripts/Operation

- **Main performance harness (run_perf.py)**
 - Gets **run_id** through REST API
 - Starts monitoring script (process_monitor.py) on each cluster node
 - Initiates Spark jobs and adds log messages for correlation of process data and job execution



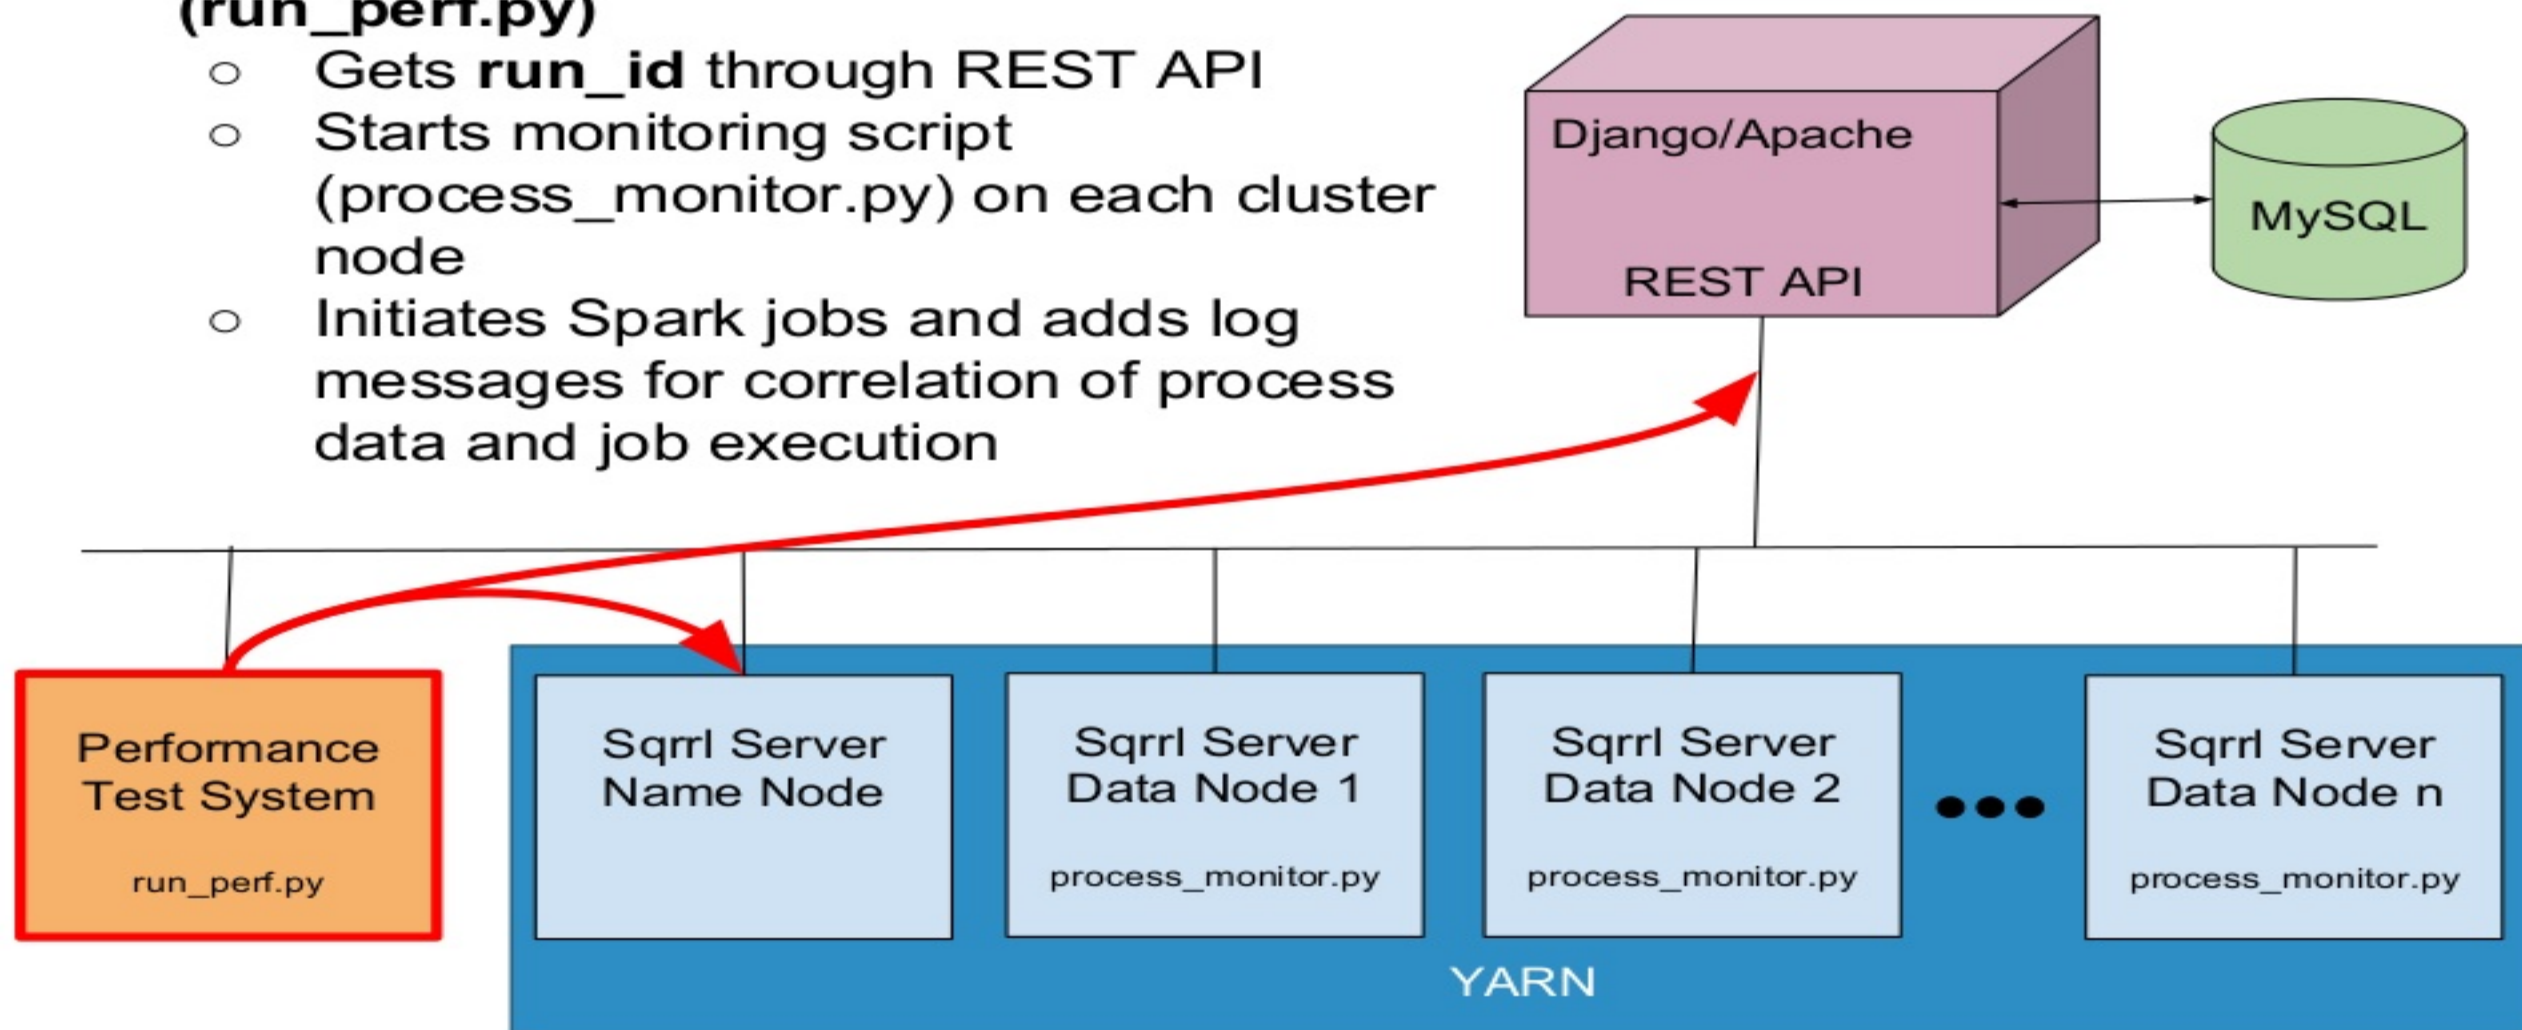
Scripts/Operation

- **Main performance harness (run_perf.py)**
 - Gets **run_id** through REST API
 - Starts monitoring script (process_monitor.py) on each cluster node
 - Initiates Spark jobs and adds log messages for correlation of process data and job execution



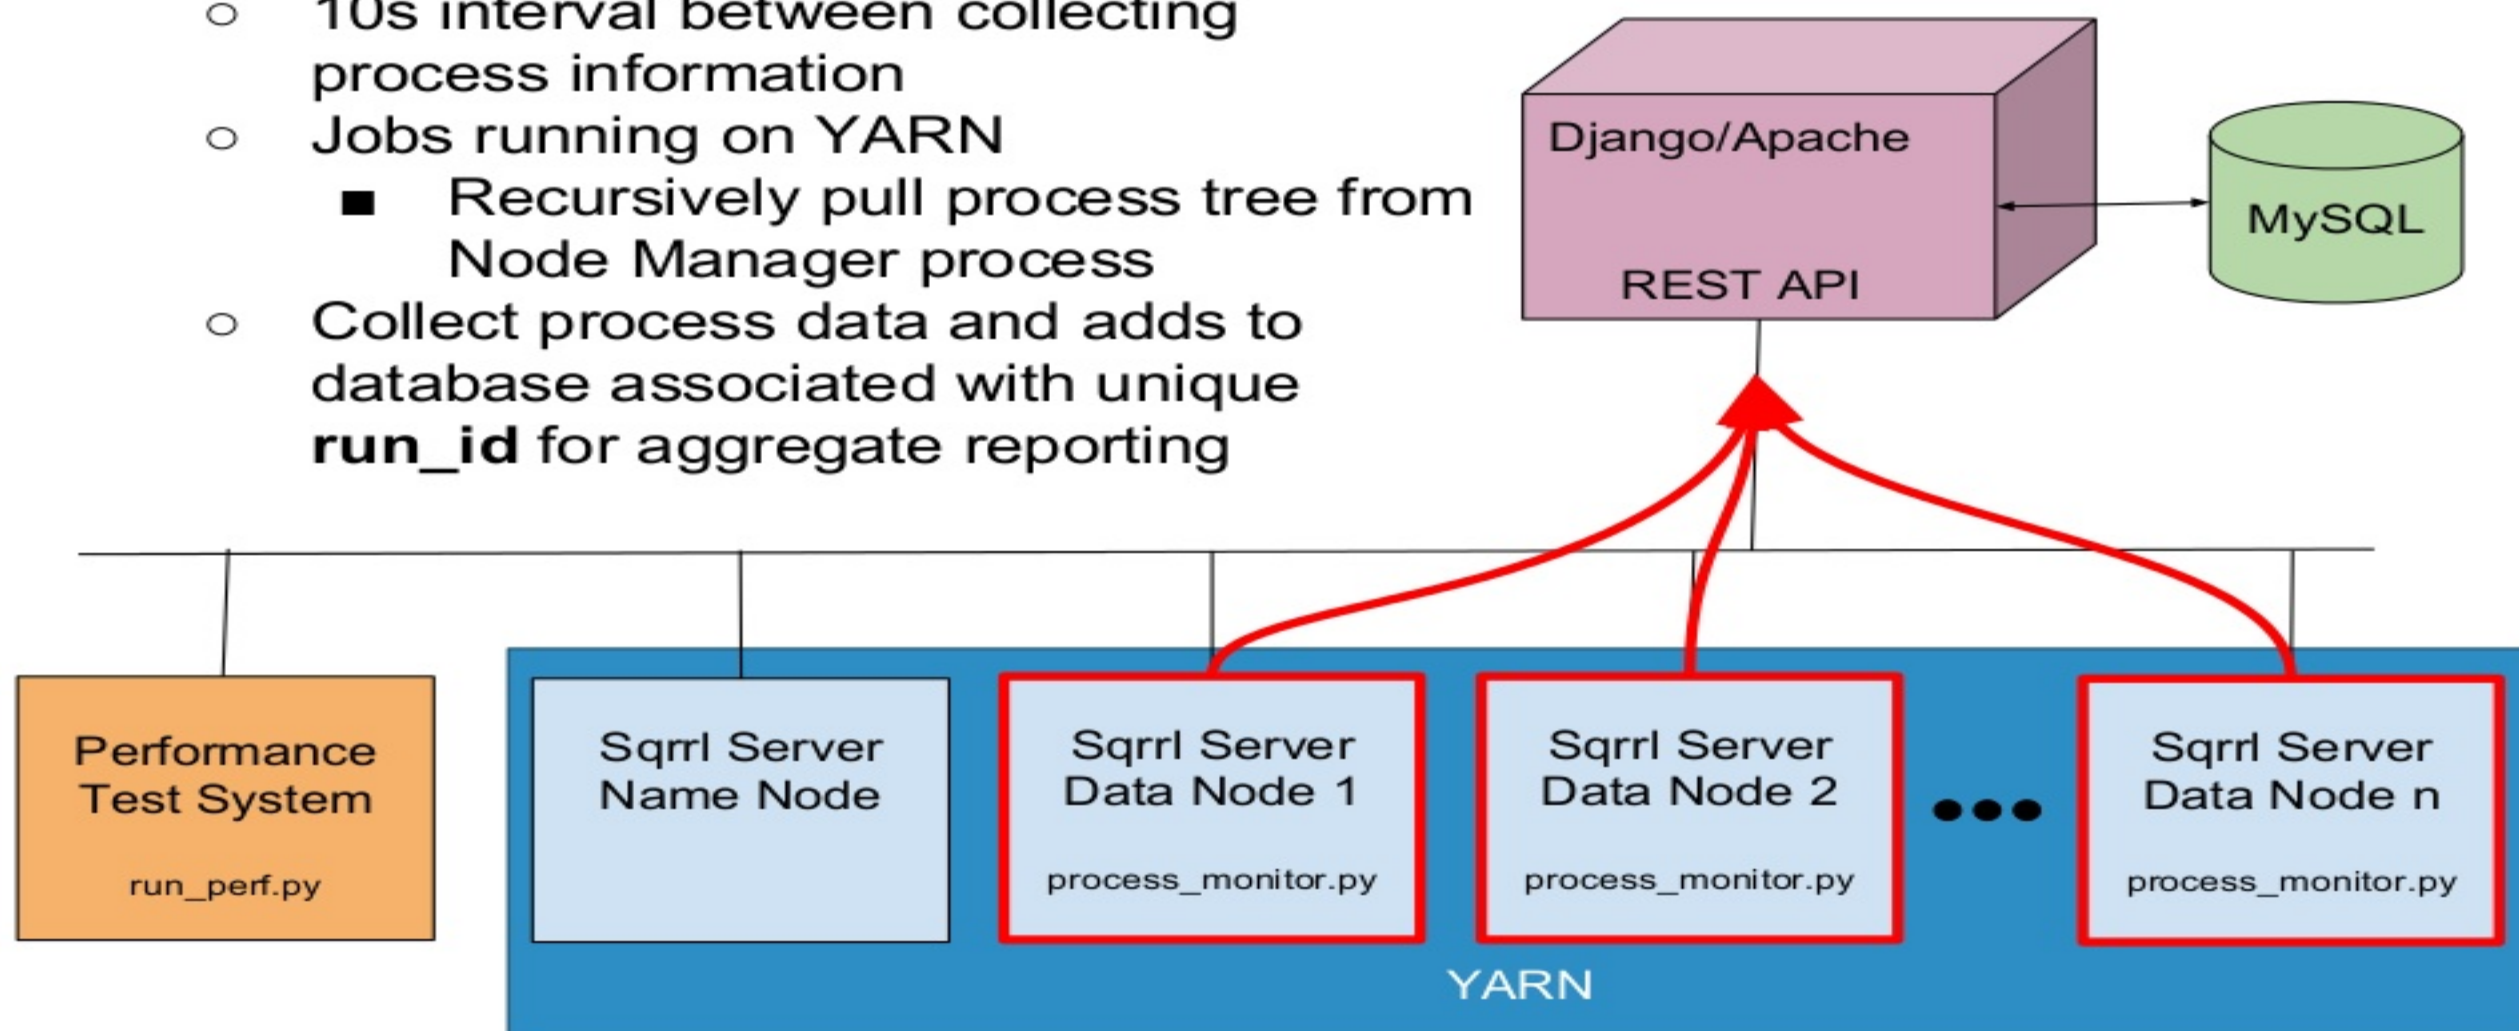
Scripts/Operation

- **Main performance harness (run_perf.py)**
 - Gets **run_id** through REST API
 - Starts monitoring script (process_monitor.py) on each cluster node
 - Initiates Spark jobs and adds log messages for correlation of process data and job execution



Scripts/Operation

- **Monitoring script (process_monitor.py)**
 - 10s interval between collecting process information
 - Jobs running on YARN
 - Recursively pull process tree from Node Manager process
 - Collect process data and adds to database associated with unique **run_id** for aggregate reporting



Database Design

Runs
id
start_time
end_time
...
suite_name
user
overall_status

Log
id
run_id
msg_time
msg_type
message

ProcessMonitor
id
time_stamp
monitor
run_id
pid
ppid
pcpu
pmem
cpu_time
elapsed_time
resident_size
mapped
private
shared
short_command
long_command

User Experience

Dashboard Results Longevity Dashboard Infrastructure Dashboard	Run ID	Start Time	End Time	Group	Suite
	2196 Perf Proc	2016-12-01 16:40			run_detector_perf_no_ootb.py
	2195 Perf Proc	2016-12-01 03:00			run_detector_perf.py

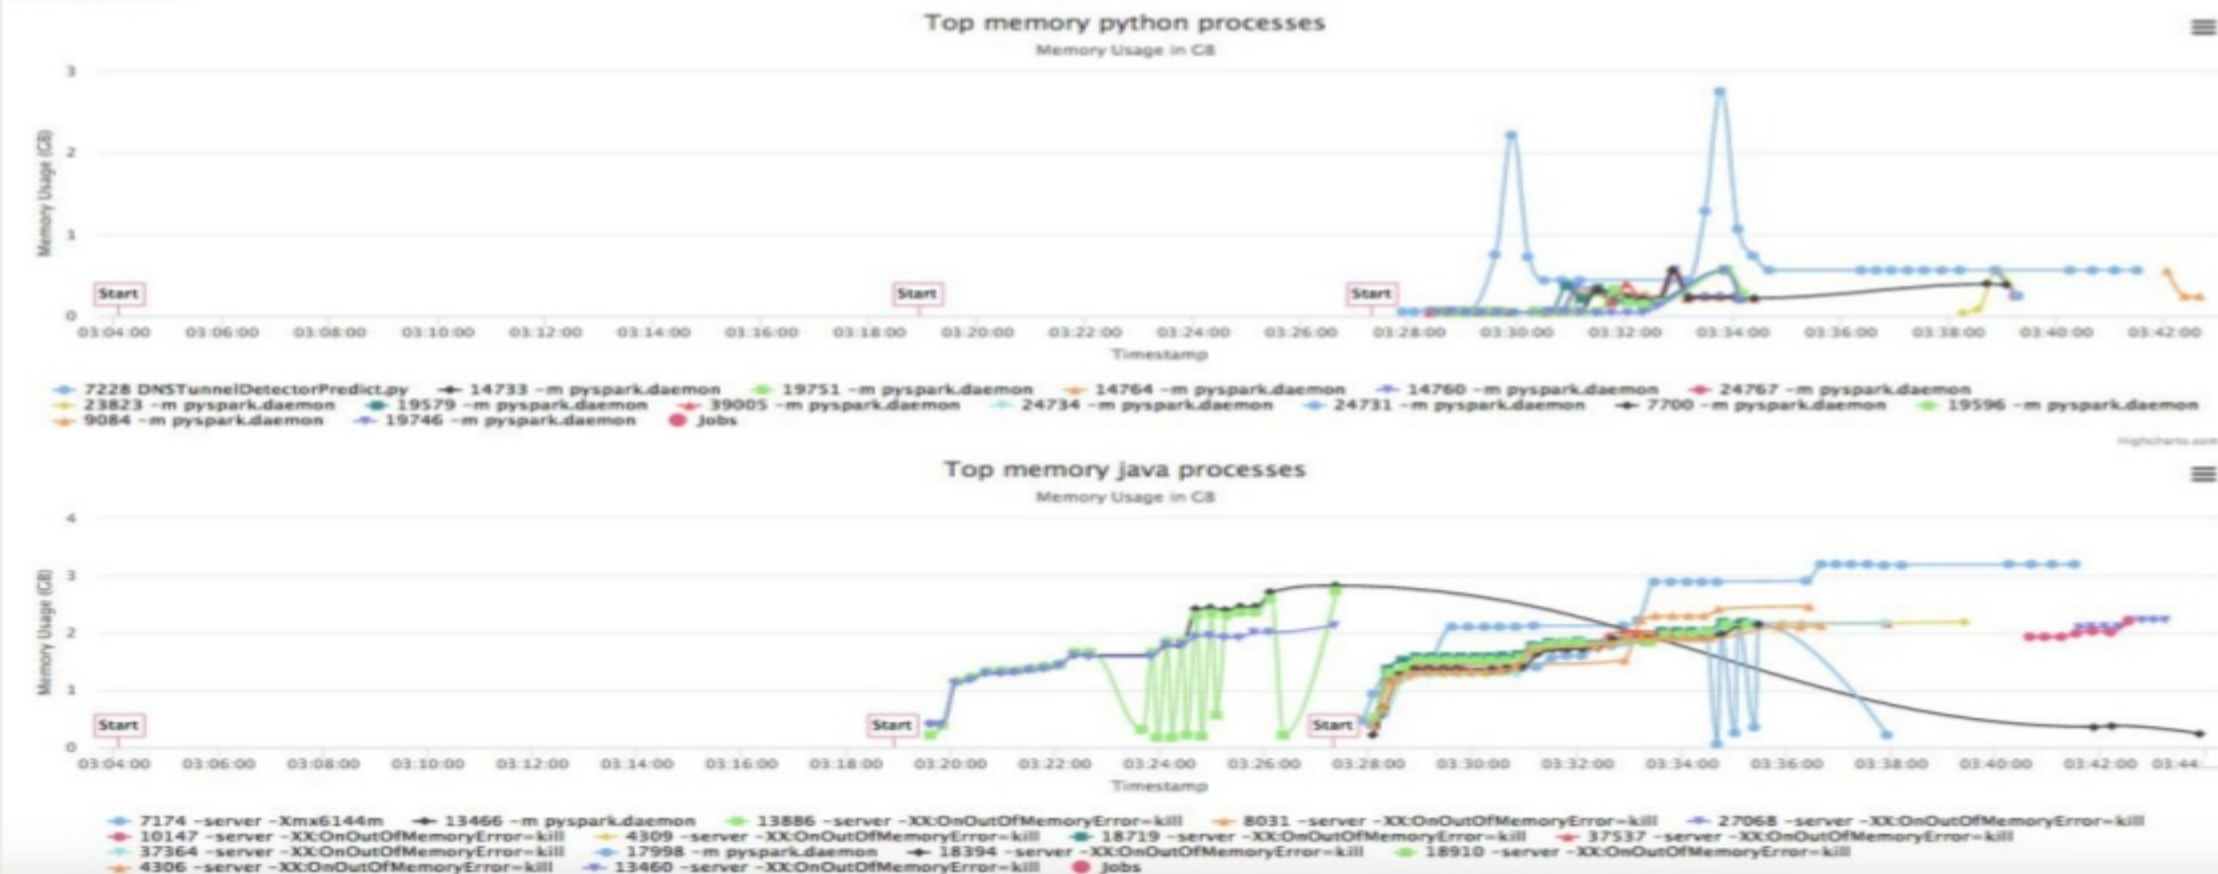
User Experience

Dashboard	Run ID	Start Time	End Time	Group	Suite
Results	2196 Perf Proc	2016-12-01 16:40			run_detector_perf_no_outb.py
Longevity Dashboard	2195 Perf Proc	2016-12-01 03:00			run_detector_perf.py
Infrastructure Dashboard					

Run: 2195

Measure Name	Type	PID	Max % CPU	Max Virtual (GB)	Max Resident Size (GB)
detectors dns-Microsoft-DNS-Debug-log-load-job-ms dns 70m.log	JAVA	529.00	5.841876	4.997276	
detectors dns-Microsoft-DNS-Debug-log-load-job-ms dns 70m.log	YARN	0.40	0.009244	0.001104	
detectors dns-DNS-Tunnel---training-for-Microsoft-DNS-Debug-Logs-ms dns 70m.log	JAVA	185.00	7.740376	2.8117	
detectors dns-DNS-Tunnel---training-for-Microsoft-DNS-Debug-Logs-ms dns 70m.log	PYTHON	251.00	3.474492	1.702632	
detectors dns-DNS-Tunnel---training-for-Microsoft-DNS-Debug-Logs-ms dns 70m.log	YARN	0.00	0.009412	0.00124	
detectors dns-DNS-Tunnel---prediction-for-Microsoft-DNS-Debug-Logs-ms dns 70m.log	JAVA	217.00	7.757036	3.193608	
detectors dns-DNS-Tunnel---prediction-for-Microsoft-DNS-Debug-Logs-ms dns 70m.log	PYTHON	434.00	4.556148	2.747848	
detectors dns-DNS-Tunnel---prediction-for-Microsoft-DNS-Debug-Logs-ms dns 70m.log	YARN	0.00	0.009412	0.00124	

The screenshot shows the AWS IAM console interface. On the left, the navigation menu has a red circle around the 'Users' link. The main content area shows the 'Users' page with a table of users. A red circle highlights the 'Users' tab in the top navigation bar, and a red arrow points to the 'Users' link in the left-hand navigation menu.



User Experience

Run: 2195	Process Name	Type	Path	Memory (MB)	Private (MB)	Shared (MB)	Open Files	Max Open Files	Max Connections	Max CPU	Max Memory	Max Private	Max Shared	Max Open Files	Max Connections	Max CPU	Max Memory	Max Private	Max Shared
7228	DNSTunnelDetectorPredict.py	Python	/usr/local/bin/DNSTunnelDetectorPredict.py	100	100	0	0	0	0	0	0	0	0	0	0	0	0	0	0
14733	-m pyspark.daemon	Java	/usr/lib/jvm/java-8-jdk/bin/java	100	100	0	0	0	0	0	0	0	0	0	0	0	0	0	0
19751	-m pyspark.daemon	Java	/usr/lib/jvm/java-8-jdk/bin/java	100	100	0	0	0	0	0	0	0	0	0	0	0	0	0	0
14764	-m pyspark.daemon	Java	/usr/lib/jvm/java-8-jdk/bin/java	100	100	0	0	0	0	0	0	0	0	0	0	0	0	0	0
14767	-m pyspark.daemon	Java	/usr/lib/jvm/java-8-jdk/bin/java	100	100	0	0	0	0	0	0	0	0	0	0	0	0	0	0
24767	-m pyspark.daemon	Java	/usr/lib/jvm/java-8-jdk/bin/java	100	100	0	0	0	0	0	0	0	0	0	0	0	0	0	0
23823	-m pyspark.daemon	Java	/usr/lib/jvm/java-8-jdk/bin/java	100	100	0	0	0	0	0	0	0	0	0	0	0	0	0	0
19579	-m pyspark.daemon	Java	/usr/lib/jvm/java-8-jdk/bin/java	100	100	0	0	0	0	0	0	0	0	0	0	0	0	0	0
39005	-m pyspark.daemon	Java	/usr/lib/jvm/java-8-jdk/bin/java	100	100	0	0	0	0	0	0	0	0	0	0	0	0	0	0
24734	-m pyspark.daemon	Java	/usr/lib/jvm/java-8-jdk/bin/java	100	100	0	0	0	0	0	0	0	0	0	0	0	0	0	0
24731	-m pyspark.daemon	Java	/usr/lib/jvm/java-8-jdk/bin/java	100	100	0	0	0	0	0	0	0	0	0	0	0	0	0	0
7700	-m pyspark.daemon	Java	/usr/lib/jvm/java-8-jdk/bin/java	100	100	0	0	0	0	0	0	0	0	0	0	0	0	0	0
19596	-m pyspark.daemon	Java	/usr/lib/jvm/java-8-jdk/bin/java	100	100	0	0	0	0	0	0	0	0	0	0	0	0	0	0
9084	-m pyspark.daemon	Java	/usr/lib/jvm/java-8-jdk/bin/java	100	100	0	0	0	0	0	0	0	0	0	0	0	0	0	0
19746	-m pyspark.daemon	Java	/usr/lib/jvm/java-8-jdk/bin/java	100	100	0	0	0	0	0	0	0	0	0	0	0	0	0	0
Jobs	Jobs	Jobs																	

Run: 2195


Top memory python processes

Memory Usage in GB



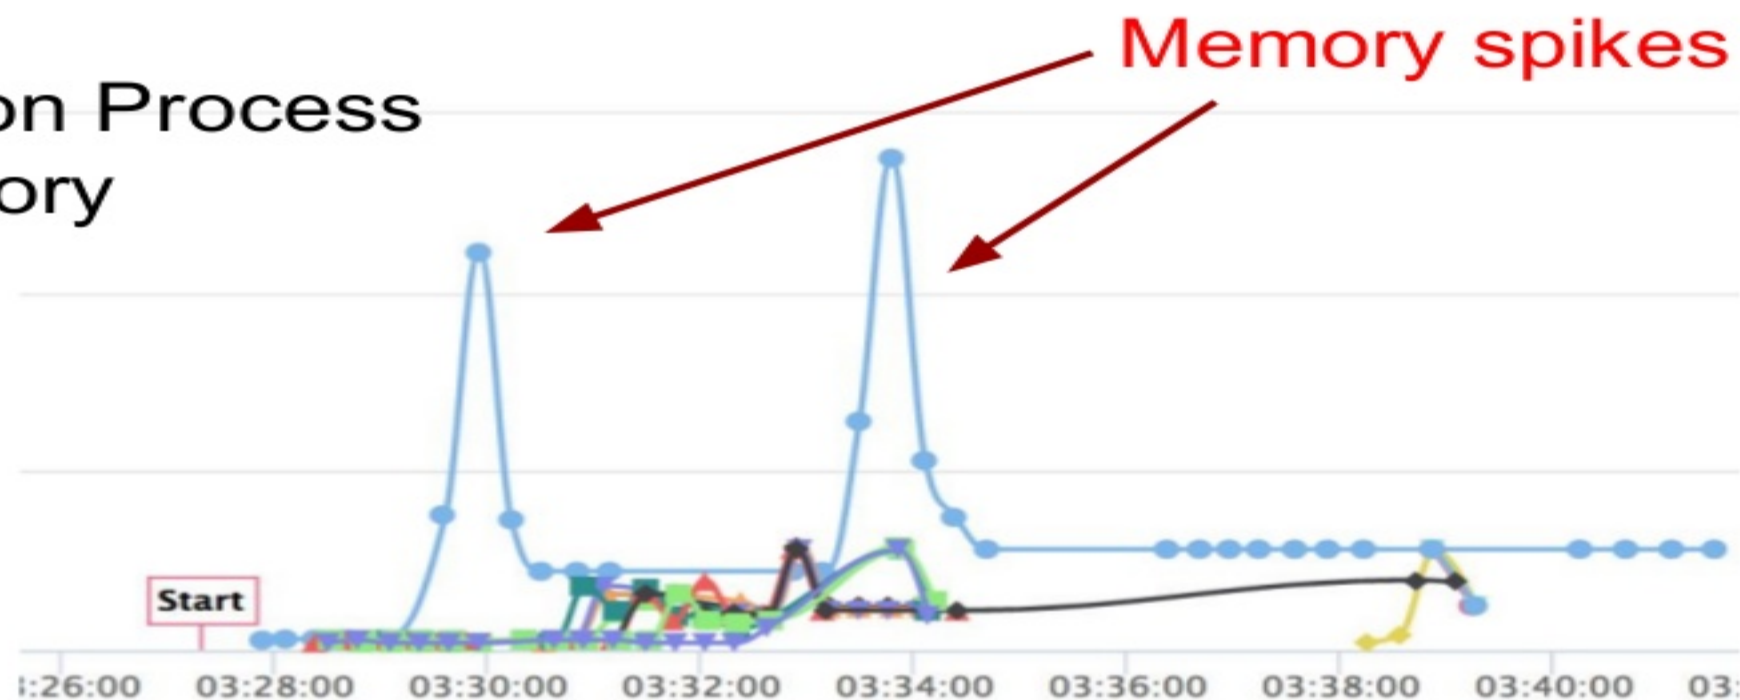
Top memory java processes

Memory Usage in GB



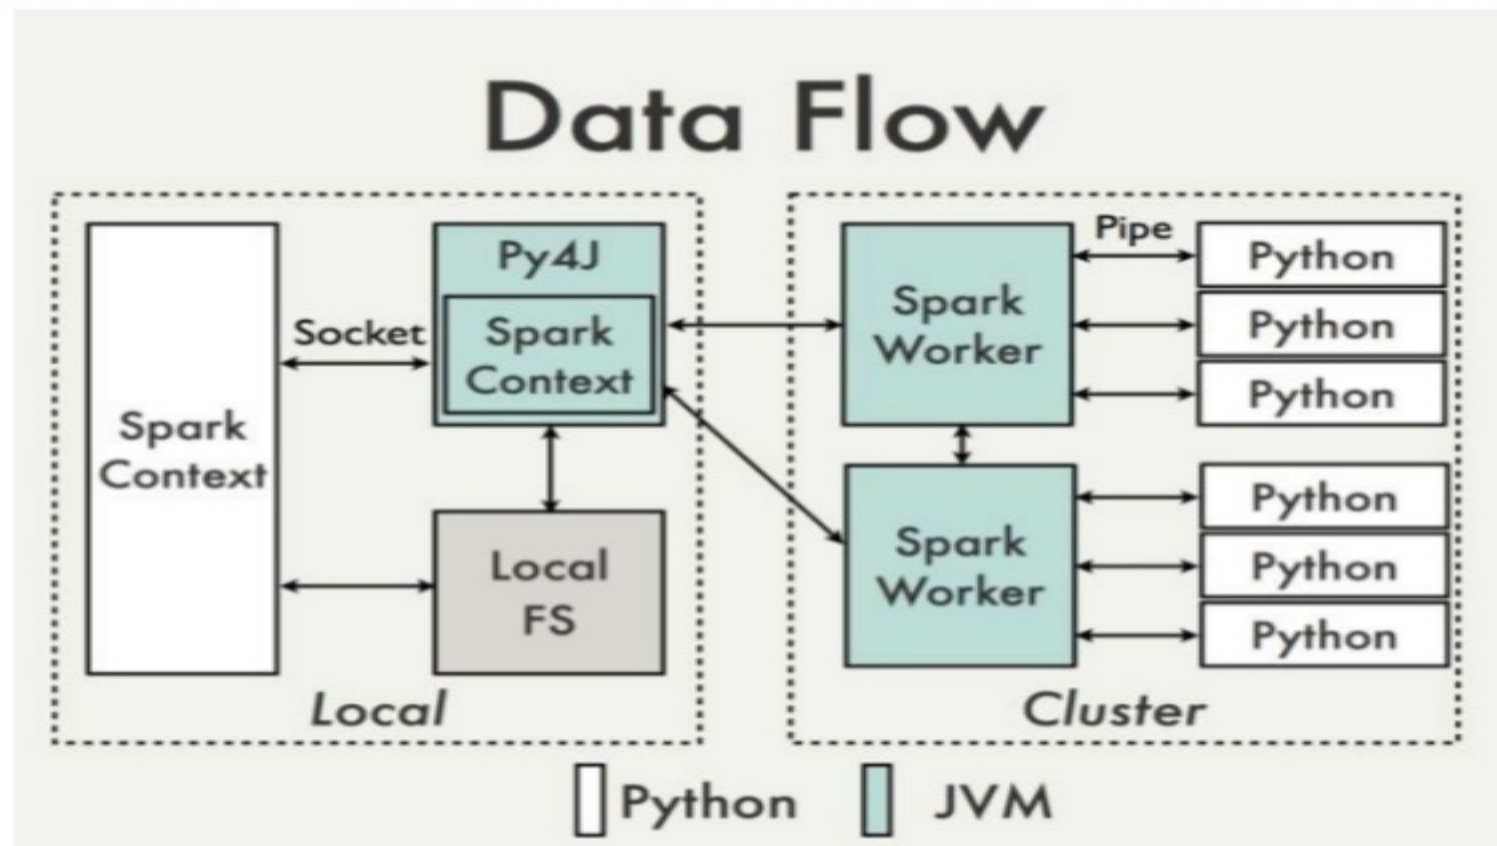
Py4J Issue: Evidence

Python Process
Memory




- Memory spiked during loading of trained ML models into python process.
- Only on driver!

Diagnosis and Conclusions



- Loading trained model from HDFS required decryption with java libraries and Py4J for loading into python.
- Py4J protocol inflated size by almost 100x!
- Solution: implement custom protocol for loading data from jvm to python via socket.

Py4J issue: After Fix



- Memory spikes are gone!
- Ready for production!

Recommendations & Lessons Learned

- Do not take scalability for granted!
- Understand Spark's architecture
 - Python/JVM interaction
- Follow best practices
 - Iterators not Lists
 - Careful with joins
- Understand your computing demands
- Test at scale
- Invest in tools
- Think distributed and your code will shine!

Thank You.



sqrri.com
@SqrriData

ebarnes@sqrri.com
ruslan@sqrri.com
chris@sqrri.com

