

# Netflix's Recommendation ML Pipeline using Apache Spark


DB Tsai  
Spark Summit East - Feb 8, 2017

**NETFLIX**

**At Netflix, we use ML everywhere**

**NETFLIX**

# Everything is a Recommendation



**Over 80%** of what members watch comes from our recommendations

Recommendations are driven by **Machine Learning Algorithms**

**NETFLIX**

A person in a dark suit and light shirt stands on a stage, gesturing with both hands. Behind them is a large screen displaying the text '#netflixeverywhere' in a bold, sans-serif font. The word 'netflix' is red, and the rest is white. The background of the screen is a dark world map composed of small, light-colored squares. The stage is dimly lit, with some red lights visible in the background.

**#netflixeverywhere**

**Jan 6th, 2016**



# #NetflixEverywhere

- 93+ Million Members
- 190+ Countries
- 125+ Million streaming hours / day
- 1000 hours of Original content in 2017
- 1/3 of US internet traffic during evenings



**Constantly Innovating  
through A/B tests**



## **Data Driven**

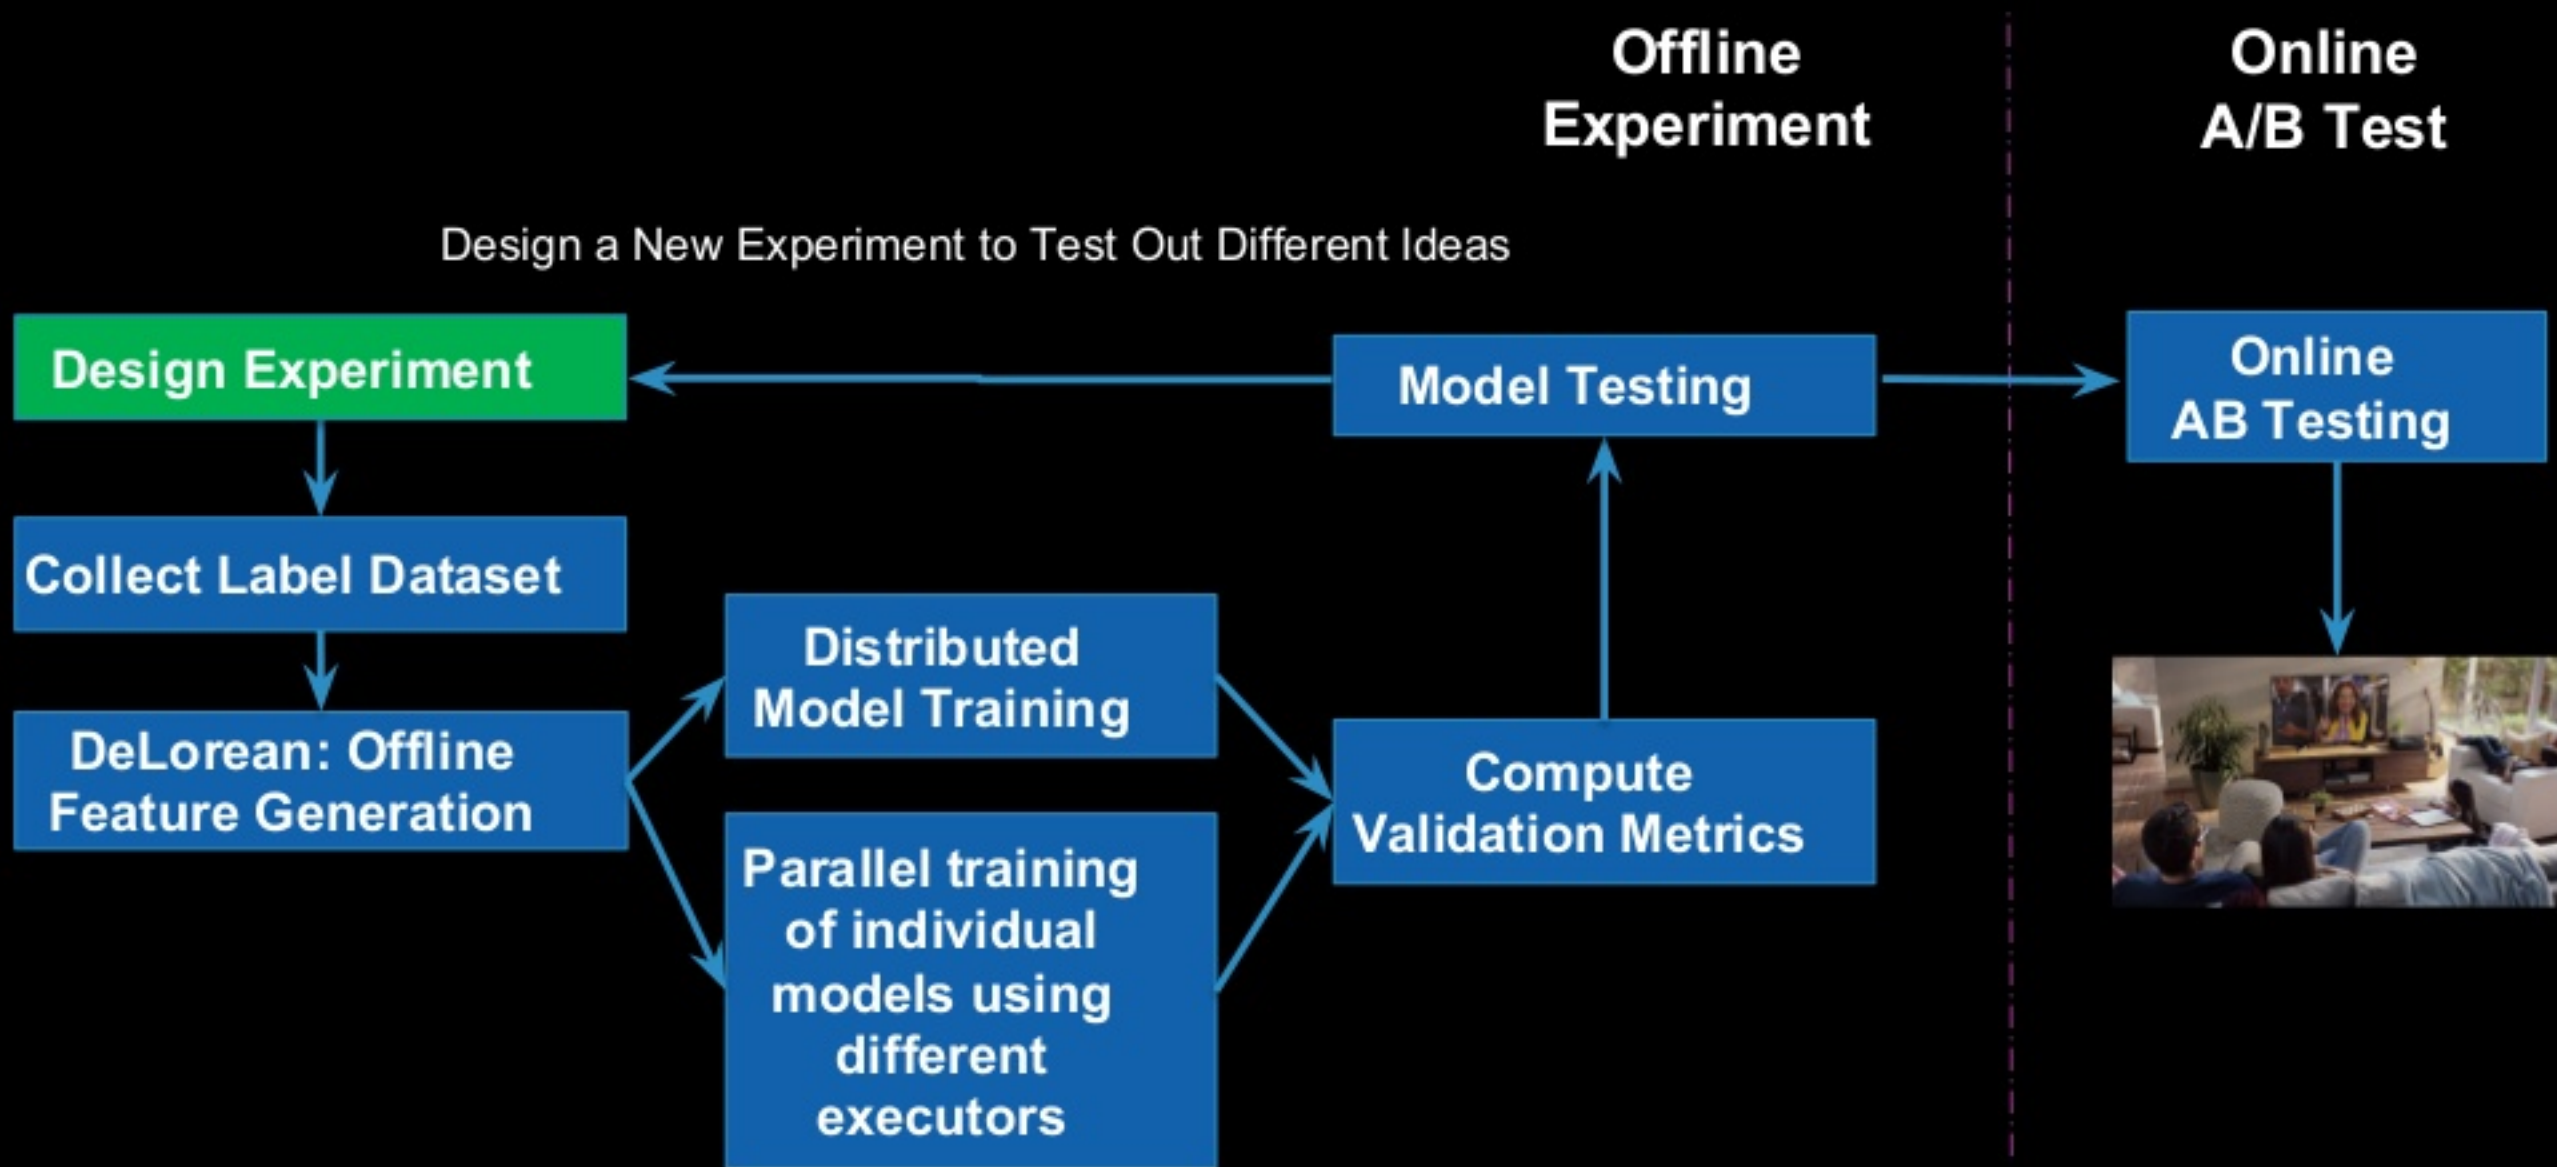
**Try an idea offline using historical data to see if they would have made better recommendations**

**If it does, deploy a live A/B test to see if it performs well in Production**

**NETFLIX**



# Running an Experiment



**We use a standardized data format across multiple ranking pipelines**

**This standardized data format is used by common tooling, libraries, and algorithms**

**NETFLIX**

## Ranking problems

**Contexts:** The setting for evaluating a set of items (e.g. tuples of member profiles, country, time, device, etc.)

**Items:** The elements to be trained on, scored, and/or ranked (e.g. videos, rows, search entities)

**Labels:** For supervised learning, this will be the label (target) for each item

**NETFLIX**

# DeLorean Data Format a.k.a DMC-12

root

|-- profile\_id: long (nullable = false)

|-- country\_iso\_code: string (nullable = false)

|-- items: array (nullable = true)

| |-- element: struct (containsNull = false)

| | |-- show\_title\_id: long (nullable = false)

| | |-- label: double (nullable = false)

| | |-- weight: double (nullable = false)

| | |-- features: struct (nullable = false)

| | | |-- feature1: double (nullable = false)

| | | |-- feature2: double (nullable = false)

| | | |-- feature3: double (nullable = false)

NETFLIX

**The nested data structure avoids an expensive shuffle when ranking**

**The features are derived from Netflix data or the output of other trained models**

**The features are persisted in HIVE using Parquet**

**Ensemble methods are used to build rankers**

**NETFLIX**



# Transformer



[https://en.wikipedia.org/wiki/Spark\\_\(Transformers\)](https://en.wikipedia.org/wiki/Spark_(Transformers)) **NETFLIX**

**Transformer takes an input DataFrame and “lazily” returns an output DataFrame**

## **Item Transformer**

- **Extends Spark ML's Transformer**
- **Accepts DMC-12 DataFrame with contextual information**
- **Transforms DataFrame at the item level**

**NETFLIX**

# Why DataFrame?

Catalyst Optimizations

Up-front Schema Verification

We found a 4x speedup during feature generation by migrating from RDD-based implementation to DataFrame implementation

**NETFLIX**

# Negative Generator

Creating negatives from what  
member plays for supervised learning

## Facts

```
root
|-- profile_id: long (nullable = false)
|-- country_iso_code: string (nullable = false)
|-- items: array (nullable = true)
|   |-- element: struct (containsNull = false)
|   |   |-- show_title_id: long (nullable = false)
|   |   |-- label: double (nullable = false)
|   |   |-- weight: double (nullable = false)
```



## Facts with synthetic negatives

```
root
|-- profile_id: long (nullable = false)
|-- country_iso_code: string (nullable = false)
|-- items: array (nullable = true)
|   |-- element: struct (containsNull = false)
|   |   |-- show_title_id: long (nullable = false)
|   |   |-- label: double (nullable = false)
|   |   |-- weight: double (nullable = false)
```

**NETFLIX**



# DeLorean Feature Generator

Creating features based on common code base  
in offline and online system

```
root
|-- profile_id: long (nullable = false)
|-- country_iso_code: string (nullable = false)
|-- items: array (nullable = true)
|   |-- element: struct (containsNull = false)
|   |   |-- show_title_id: long (nullable = false)
|   |   |-- label: double (nullable = false)
|   |   |-- weight: double (nullable = false)
```



```
root
|-- profile_id: long (nullable = false)
|-- country_iso_code: string (nullable = false)
|-- items: array (nullable = true)
|   |-- element: struct (containsNull = false)
|   |   |-- show_title_id: long (nullable = false)
|   |   |-- label: double (nullable = false)
|   |   |-- weight: double (nullable = false)
|   |   |-- features: struct (nullable = false)
|   |   |   |-- feature1: double (nullable = false)
|   |   |   |-- feature2: double (nullable = false)
|   |   |   |-- feature3: double (nullable = false)
```



# Creating the Dataset for Algorithms

```
import org.apache.spark.ml.PipelineModel
val featurePipeline = new PipelineModel(Array(
  taggerTransformer,
  countryTenureStratifiedSampler,
  negativeGenTransformer,
  featureGenTransformer
))
val featureDF = featurePipeline.transform(playFeedDF)
```

## **Multithreading Model Training**

**For single machine multi-threading algorithms, we allocate one task to one machine. Multiple tasks are running in Spark for different parameters**

**Broadcast in Spark has datasize limitation, we write data into HDFS, and stream the data into the trainers in executors which run single-machine multi-threading algorithms**

## **Distributed Model Training**

**We use both Spark ML's algorithms and in-house ML implementations**

**We keep the interface similar for both multi-threading and distributed algorithms, so experimenters can try different ideas easily**

# Scoring and Ranking

**Scorer is also a Transformer  
returned from the Trainer**

**Multiple models can be scored at  
the same time in parallel**

**The ranks are derived from sorted  
scores**

**Together with labels, we can  
compute metrics, NMRR, NDCG,  
and Recall, etc**

```
root
|-- profile_id: long (nullable = false)
|-- country_iso_code: string (nullable = false)
|-- items: array (nullable = true)
|   |-- element: struct (containsNull = false)
|   |   |-- show_title_id: long (nullable = false)
|   |   |-- label: double (nullable = false)
|   |   |-- weight: double (nullable = false)
|   |   |-- features: struct (nullable = false)
|   |   |   |-- feature1: double (nullable = false)
|   |   |   |-- feature2: double (nullable = false)
|   |   |   |-- feature3: double (nullable = false)
|   |   |-- scores: struct (nullable = false)
|   |   |   |-- model1: double false
|   |   |   |-- model2: double false
```



# Lessons Learned - Pipeline Abstraction

## Pros

- **Modularity + Tests**
- **Plug-N-Play**
- **Notebook Prototyping**
- **Customizability**
- **Schema Verification**
- **Serializability (AC)**

## Cons

- **Dependent on Spark platform**
- **Not easy to bring to production**
- **Default Metric Evaluator doesn't support ranking multiple models and type of metrics in one pass**





— ANY —  
QUESTIONS