

Processing Terabyte Scale Genomics Datasets with ADAM

Frank Austin Nothaft
University of California, Berkeley
@fnothaft



Genome Resequencing

- When we sequence a human genome, we obtain several hundred GB of raw sequence data
- With a reference genome, we can use this sequence to compute diffs between individuals
- Two problems:
 - How do we compute this diff?
 - How do we make sense of the differences?

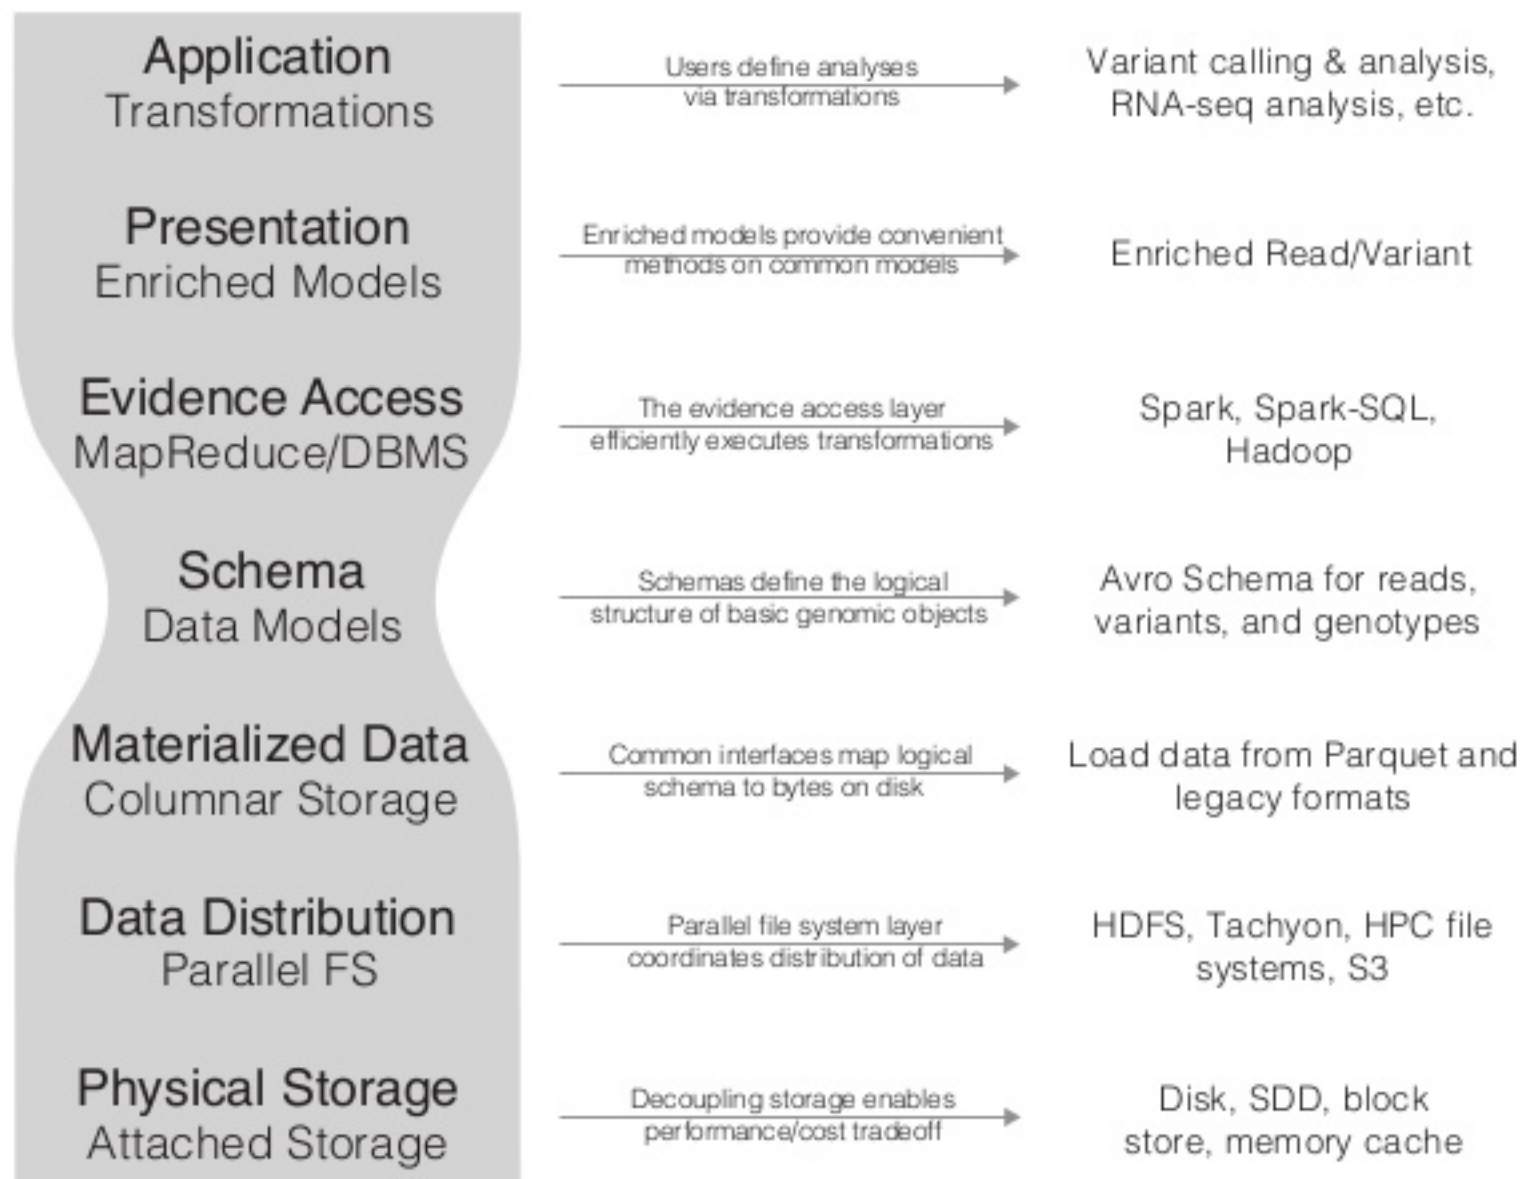
Building Scalable Genomics Tools on ADAM

- ADAM is an open source, high performance, distributed library for genomic analysis
- ADAM defines a:
 - Data schema and layout on disk
 - Programming interface for distributed processing of genomic data using Spark + Scala
- Goal is to enable both batch and exploratory analysis of all types of genomic data

Genomics is built around flattened, single node tools

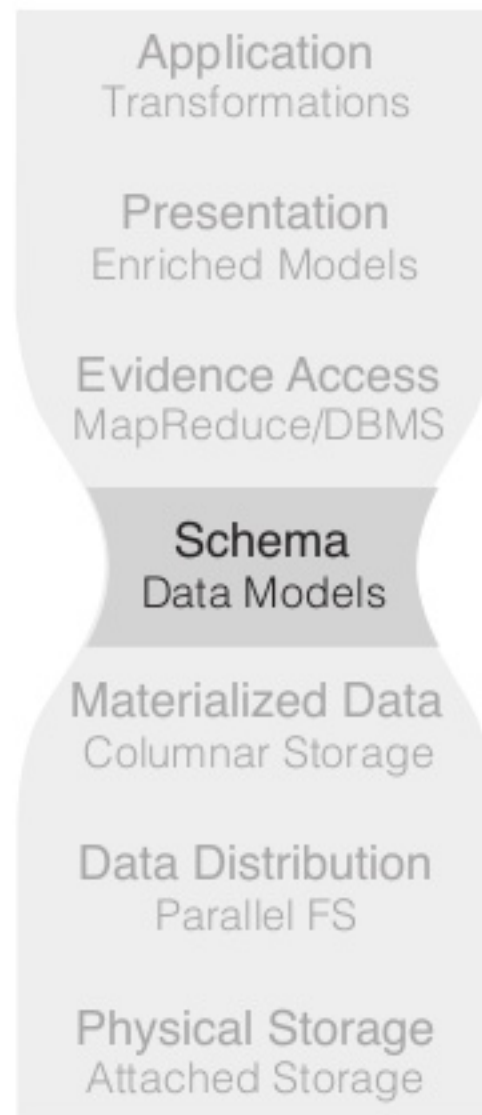
- Legacy flat-file formats:
 - Manually curated text/binary flat files
 - E.g., SAM/BAM → alignment, VCF → variants, BED/GTF/etc → features
- These formats scale poorly beyond single computer storage/compute capacity
- These legacy formats are functionally limiting and bug-prone:
 - What accesses can be optimized (read a full row)
 - What predicates can be evaluated (small number of genomic loci)
 - How we write genomic algorithms (sorted iterator over genome)
 - How we avoid technical lock-in (extend metadata)

ADAM uses a schema as a narrow waist



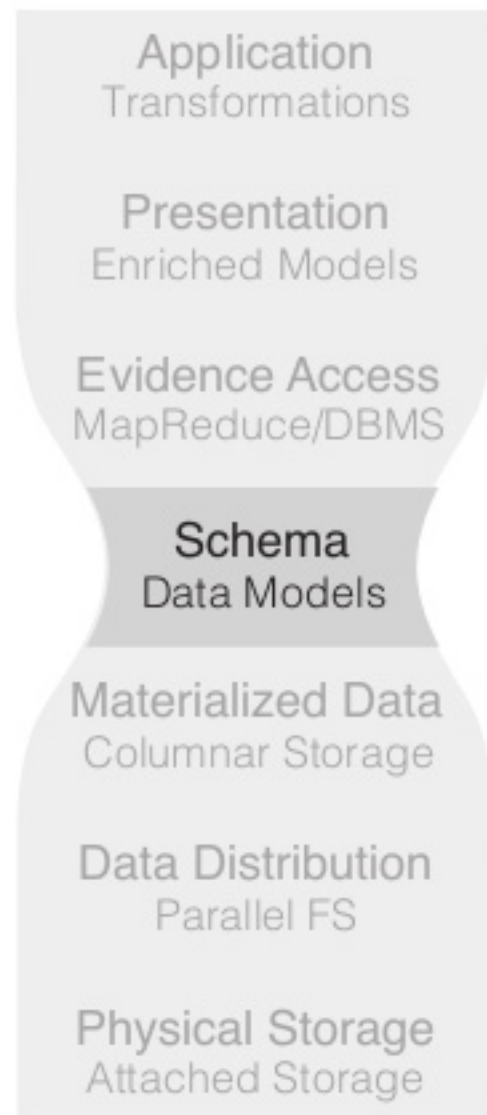
ADAM uses a schema as a narrow waist

```
record AlignmentRecord {  
  union null, Contig } contig = null;  
  union null, long } start = null;  
  union null, long } end = null;  
  union null, int } mapq = null;  
  union null, string } readName = null;  
  union null, string } sequence = null;  
  union null, string } mateReference = null;  
  union null, long } mateAlignmentStart = null;  
  union null, string } cigar = null;  
  union null, string } qual = null;  
  union null, string } recordGroupName = null;  
  union int, null } basesTrimmedFromStart = 0;  
  union int, null } basesTrimmedFromEnd = 0;  
  union boolean, null } readPaired = false;  
  union boolean, null } properPair = false;  
  union boolean, null } readMapped = false;  
  union boolean, null } mateMapped = false;  
  union boolean, null } firstOfPair = false;  
  union boolean, null } secondOfPair = false;  
  union boolean, null } failedVendorQualityChecks = false;  
  union boolean, null } duplicateRead = false;  
  union boolean, null } readNegativeStrand = false;  
  union boolean, null } mateNegativeStrand = false;  
  union boolean, null } primaryAlignment = false;  
  union boolean, null } secondaryAlignment = false;  
  union boolean, null } supplementaryAlignment = false;  
  union null, string } mismatchingPositions = null;  
  union null, string } origQual = null;  
  union null, string } attributes = null;  
  union null, string } recordGroupSequencingCenter = null;  
  union null, string } recordGroupDescription = null;  
  union null, long } recordGroupRunDateEpoch = null;  
  union null, string } recordGroupFlowOrder = null;  
  union null, string } recordGroupKeySequence = null;  
  union null, string } recordGroupLibrary = null;  
  union null, int } recordGroupPredictedMedianInsertSize = null;  
  union null, string } recordGroupPlatform = null;  
  union null, string } recordGroupPlatformUnit = null;  
  union null, string } recordGroupSample = null;  
  union null, Contig } mateContig = null;  
}
```

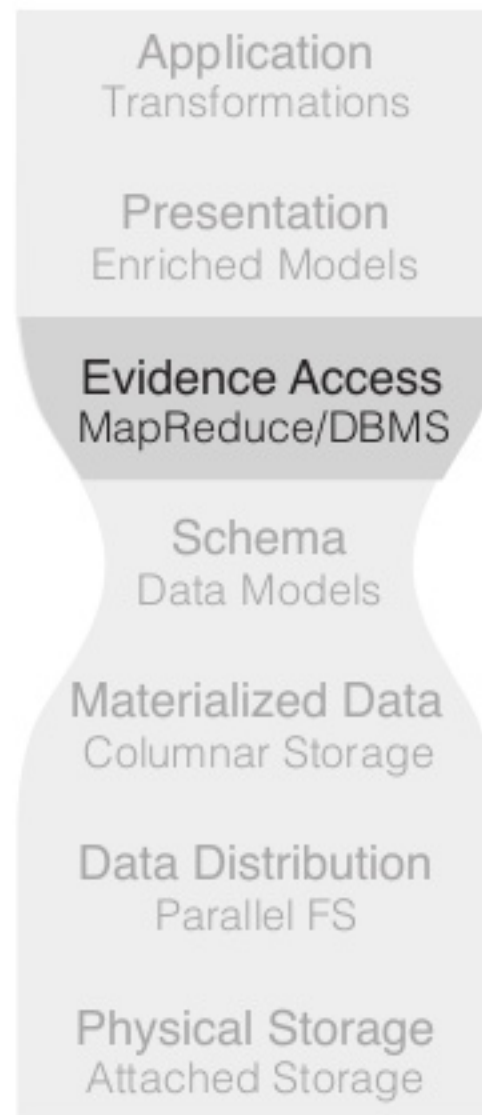


ADAM uses a schema as a narrow waist

- ADAM has schemas for:
 - Reads: SAM/BAM/CRAM, FASTQ
 - Features: BED/GTF/GFF2,3/NarrowPeak/IntervalList
 - Variants/Genotypes: (g)VCF/BCF1
 - Sequence: FASTA

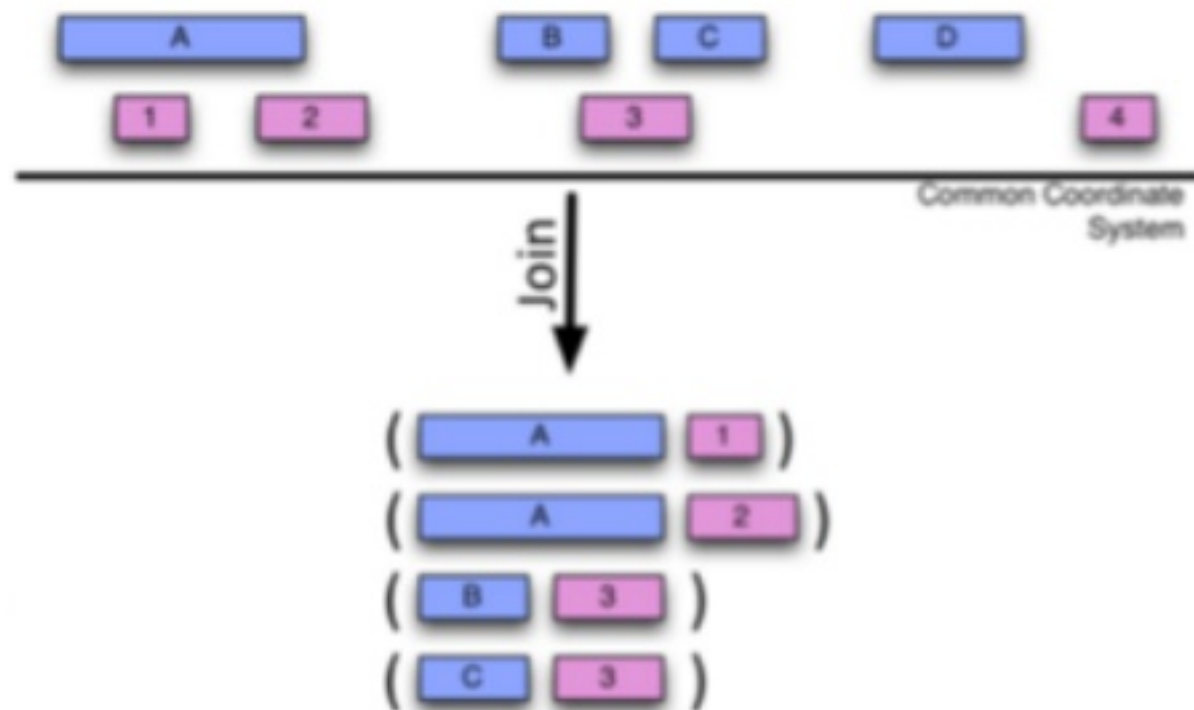


**Having a stack makes it easy to
accelerate genomic queries**



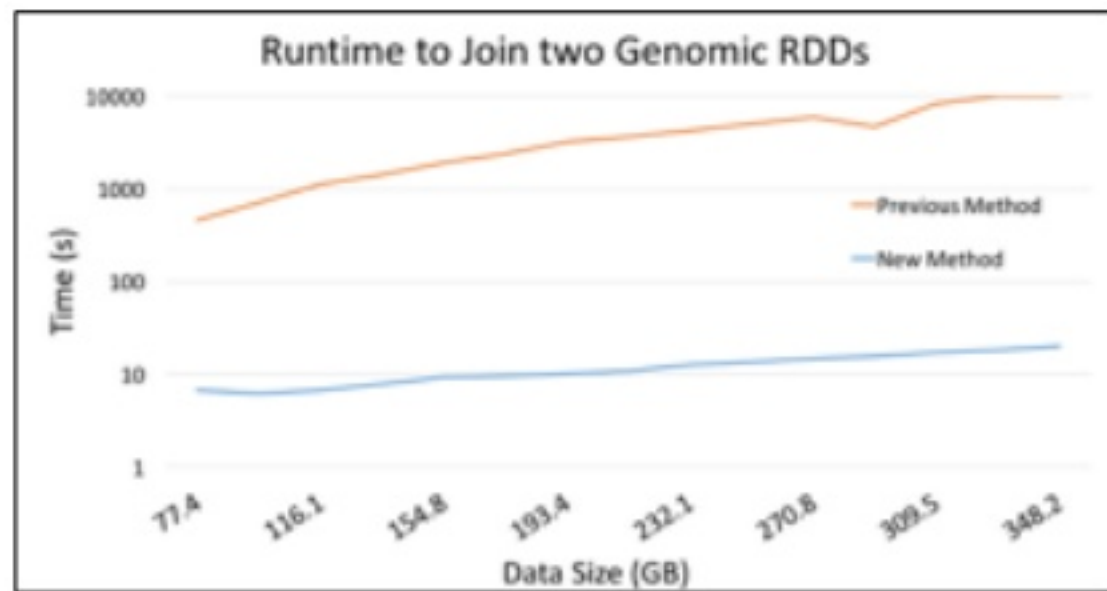
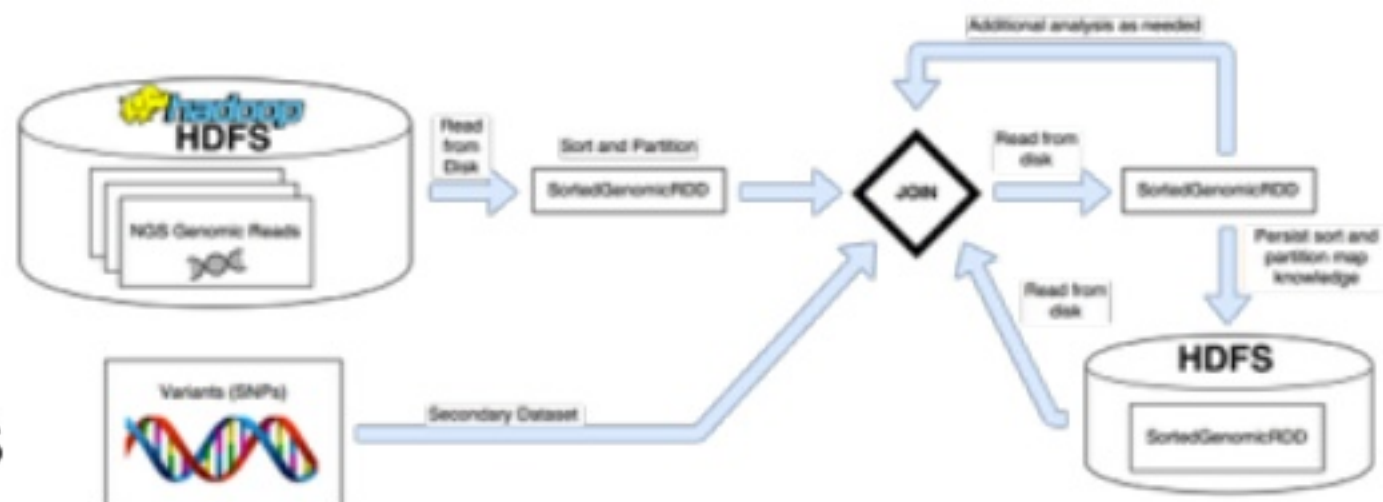
...while also providing higher level abstractions

- ADAM eliminates need to use “genome walker”:
 - Use region join for overlap computation
 - Use group/reduceByKey functions from Spark to process features aligned at a genomic coordinate point
- Can reduce targeted regions across a genome via sort + fold
- Higher level primitives enable optimizations:
 - Can leverage indices/sort orders
 - Can push down join/filter queries into storage

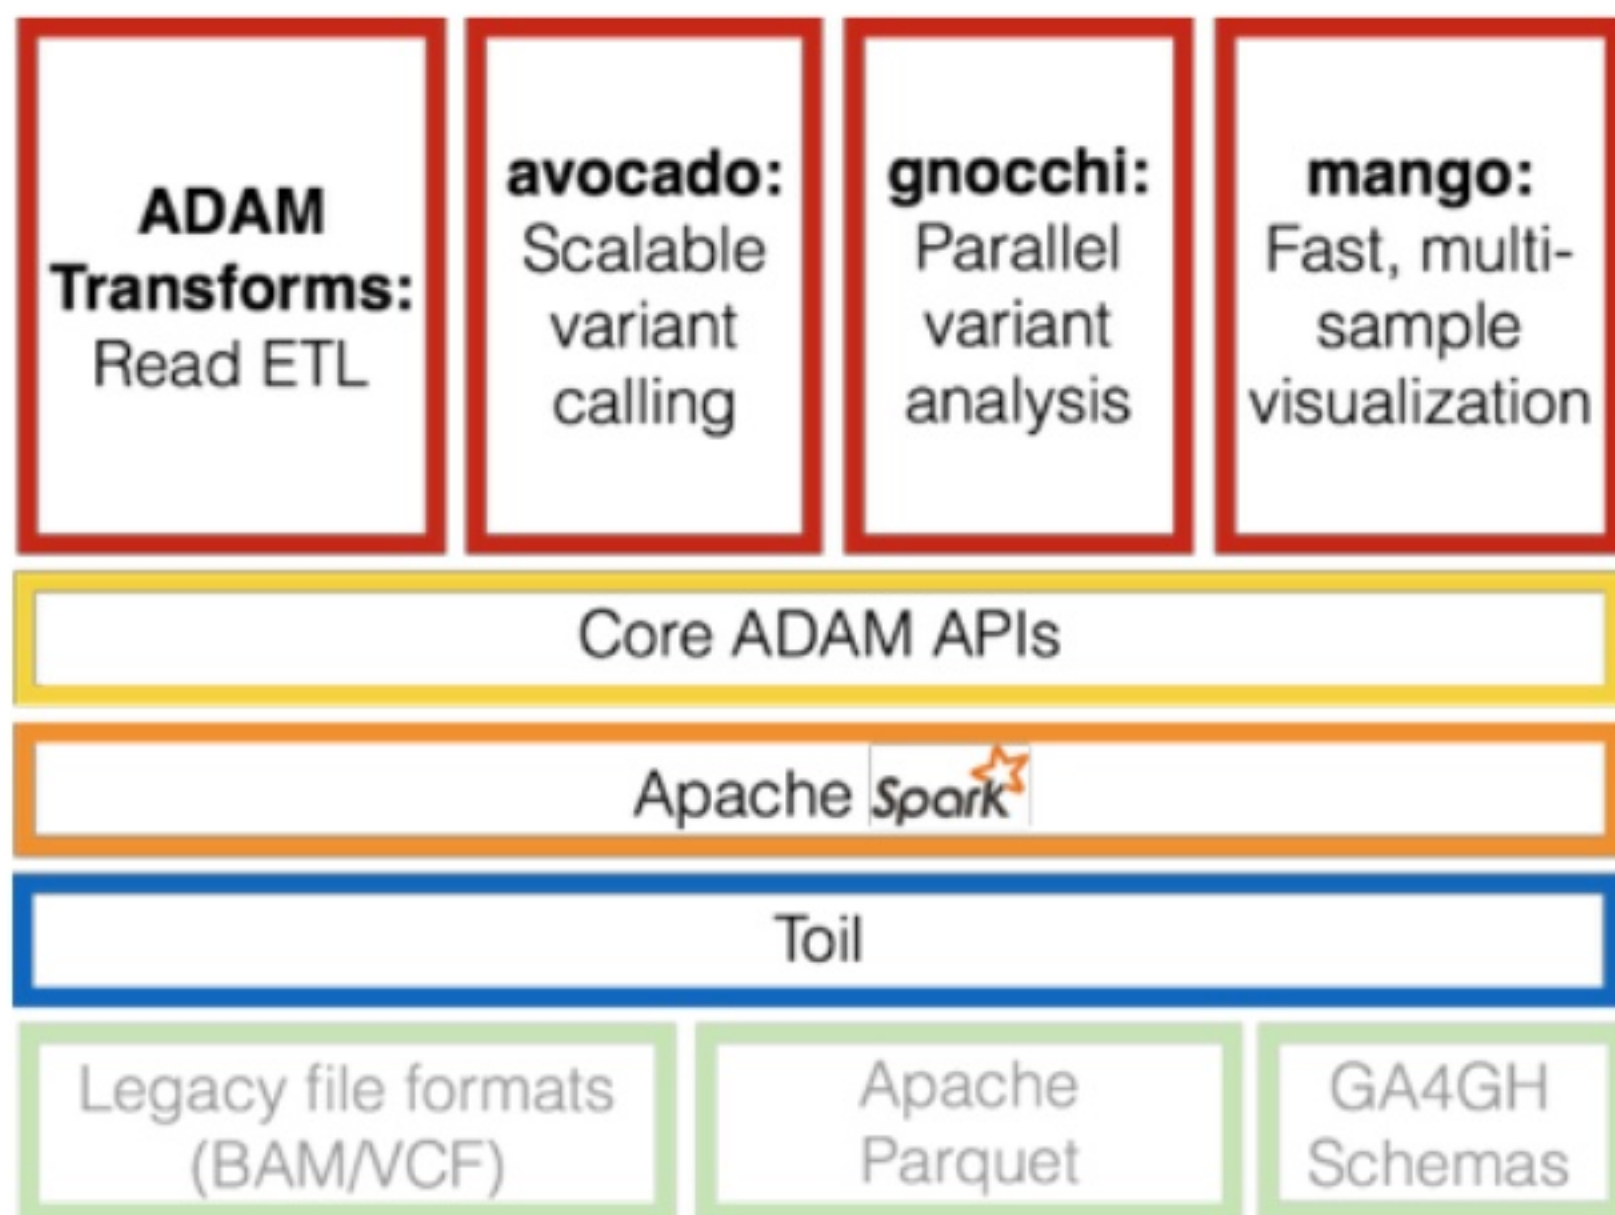


Higher level primitives enable optimizations

- Maintain sort order across runs and optimize to reduce data skew
- Leverage indices/sort orders
- Push down join/filter queries into storage
- Use join optimizations to develop BEDtools equivalent

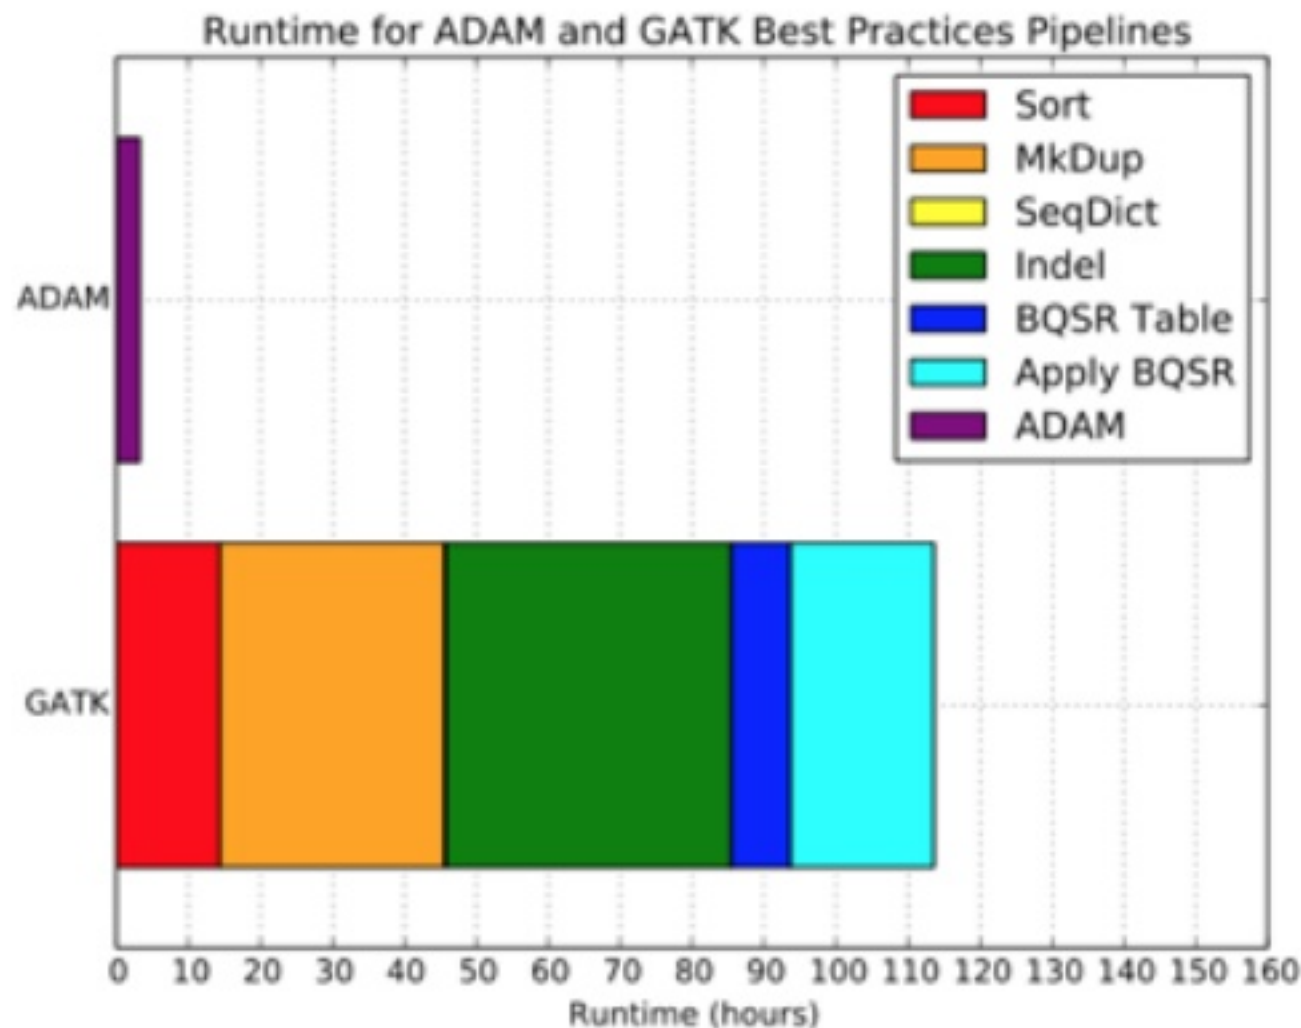


Big Data Genomics Stack



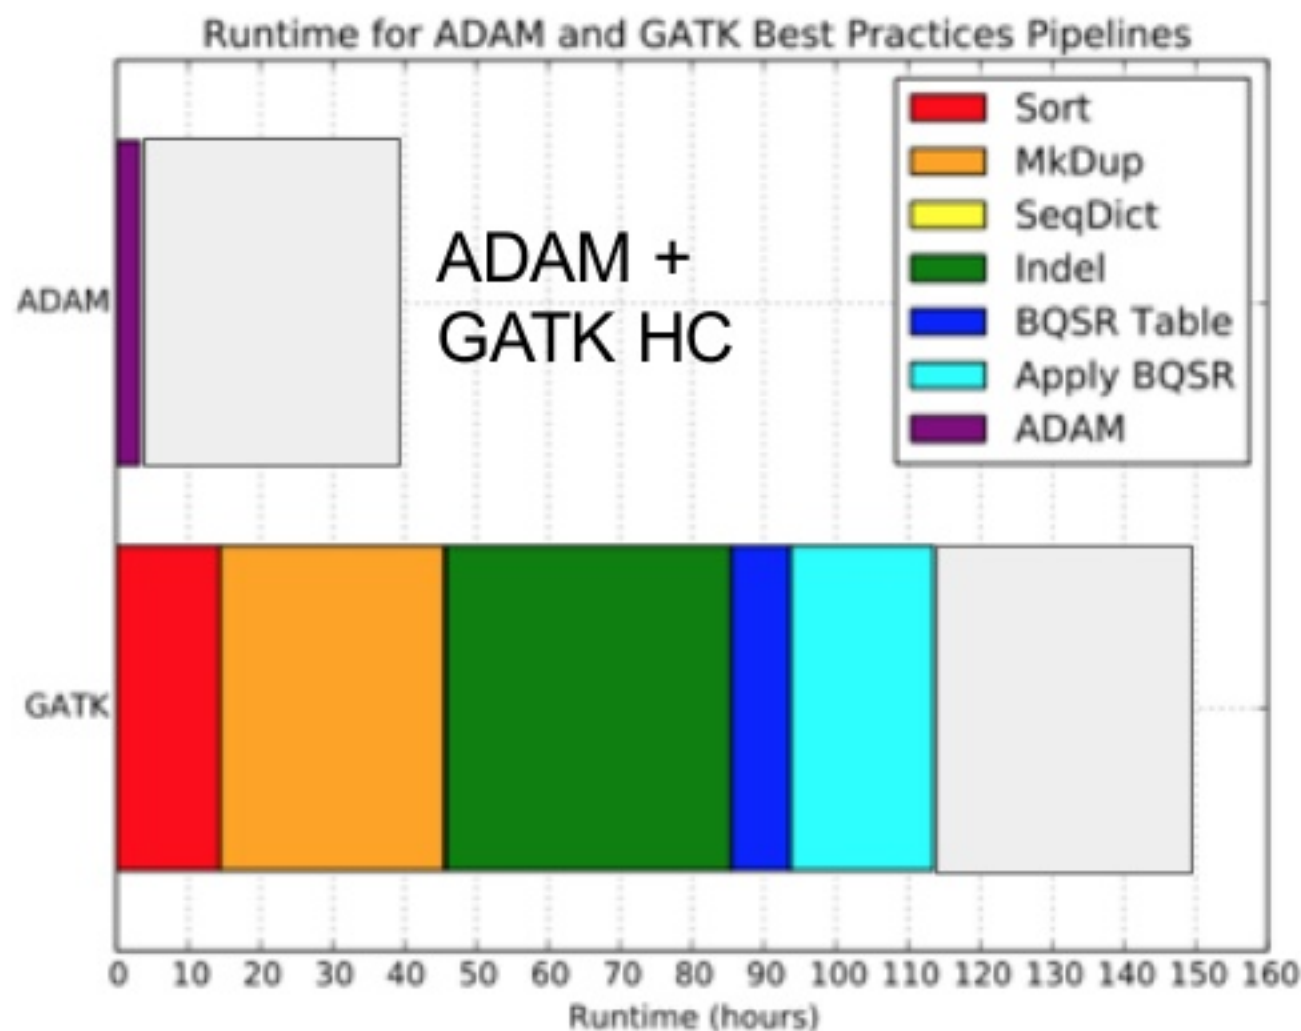
Benchmarking ADAM

- ADAM produces statistically equivalent results to the GATK best practices pipeline
- Read preprocessing is >30x faster and 3x cheaper



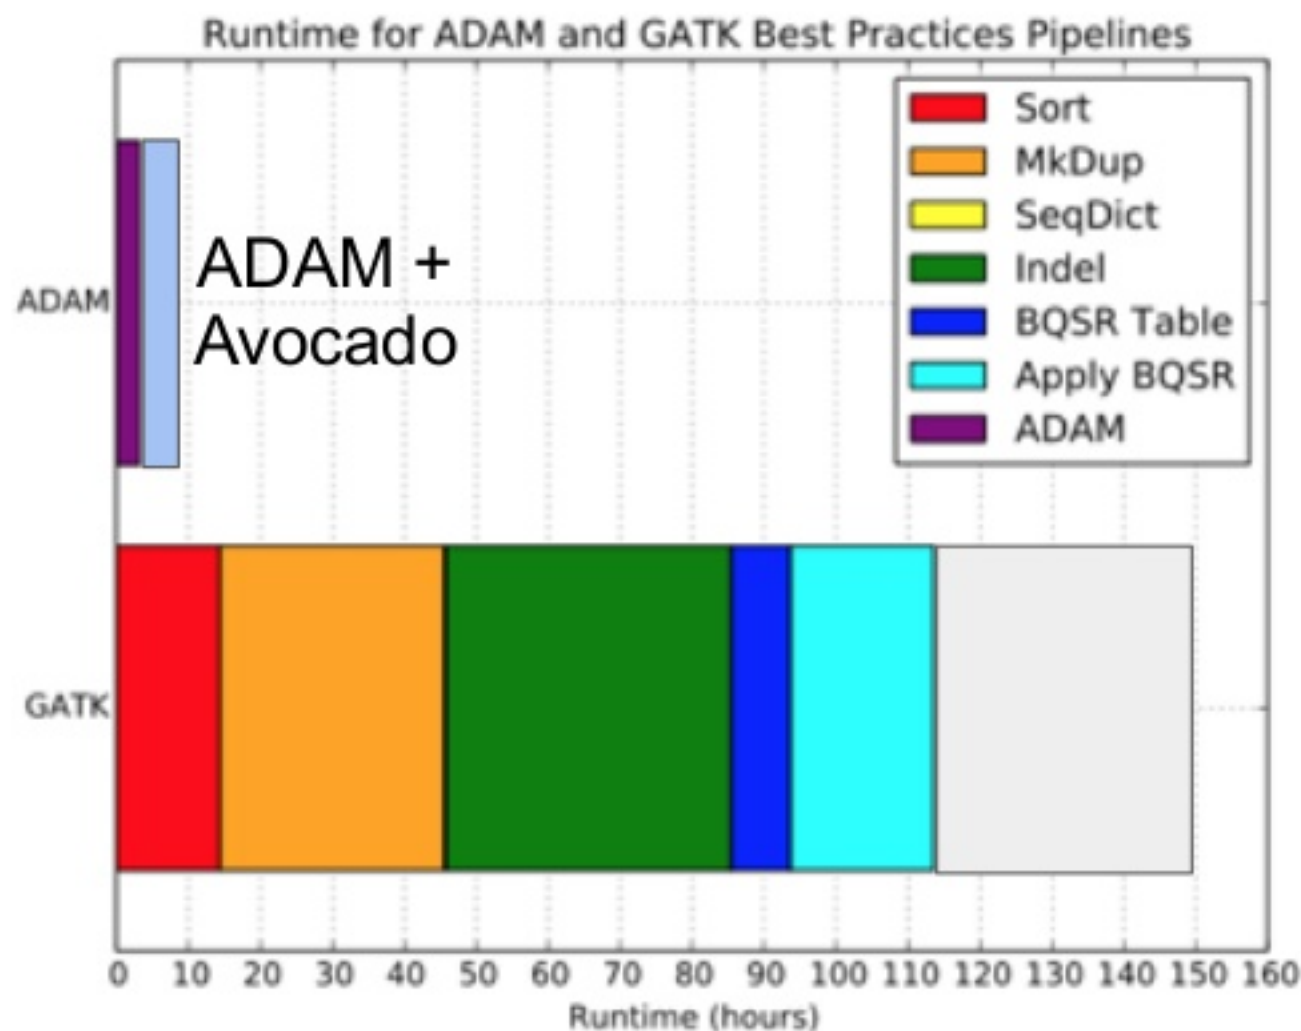
Benchmarking ADAM

- ADAM produces statistically equivalent results to the GATK best practices pipeline
- Read preprocessing is >30x faster and 3x cheaper, end-to-end pipeline is 4x faster, 3.5x cheaper

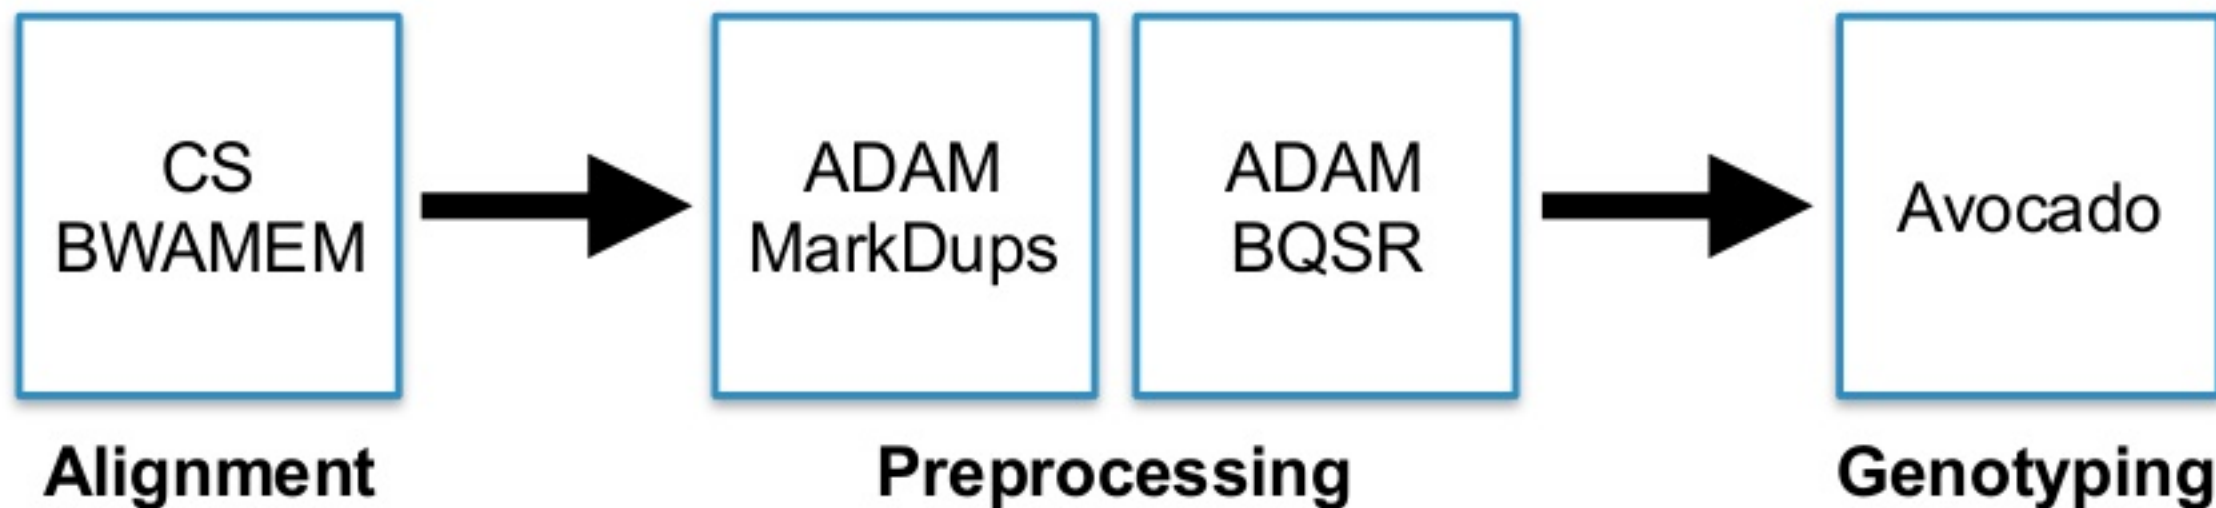


Benchmarking ADAM + Avocado

- Avocado outperforms GATK at SNP calling, slightly behind on INDELs
- Overall pipeline is >17x faster and 2x cheaper
- Avocado relies on novel, efficient INDEL canonicalization engine, drops INDEL discovery cost by 5x



End-to-end variant analysis in Spark



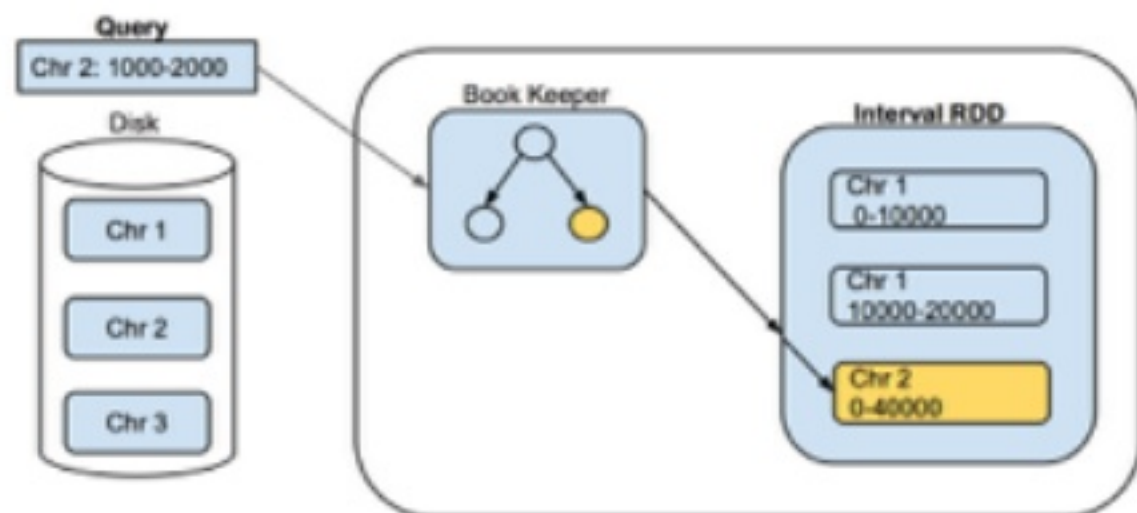
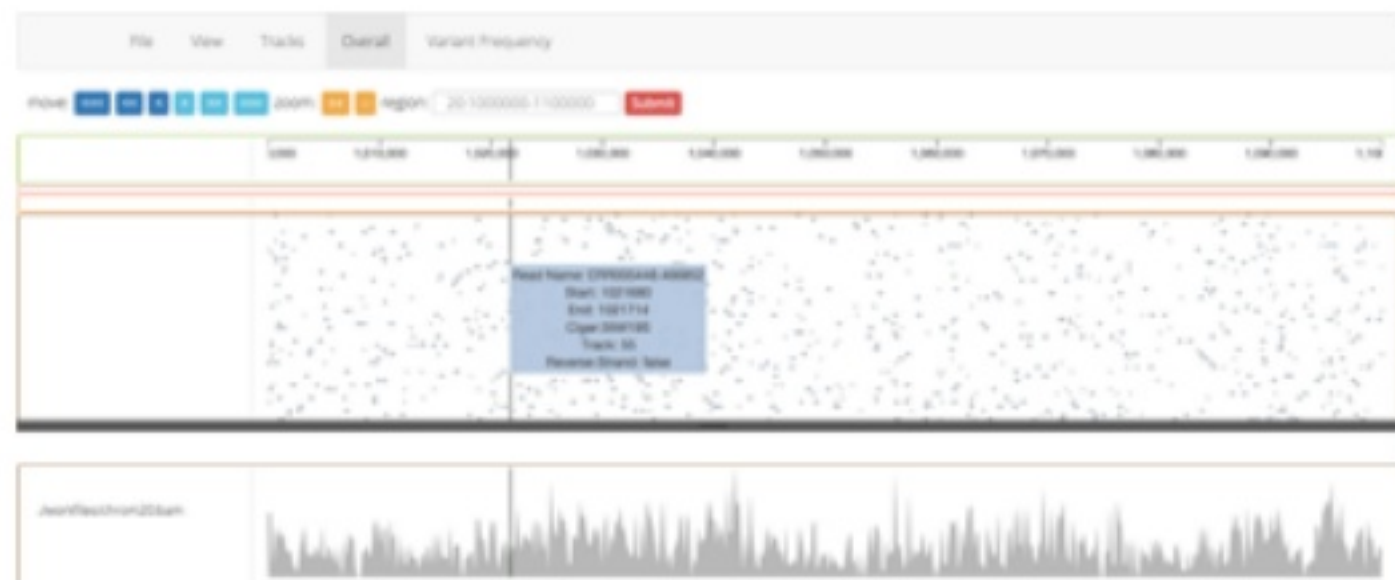
- Can process a 65x whole genome in <2hrs on 1,024 cores
- CS-BWAMEM: <https://github.com/ytchen0323/cloud-scale-bwamem>

End-to-end variant analysis in Spark



- Can process a 65x whole genome in <2hrs on 1,024 cores
- CS-BWAMEM: <https://github.com/ytchen0323/cloud-scale-bwamem>

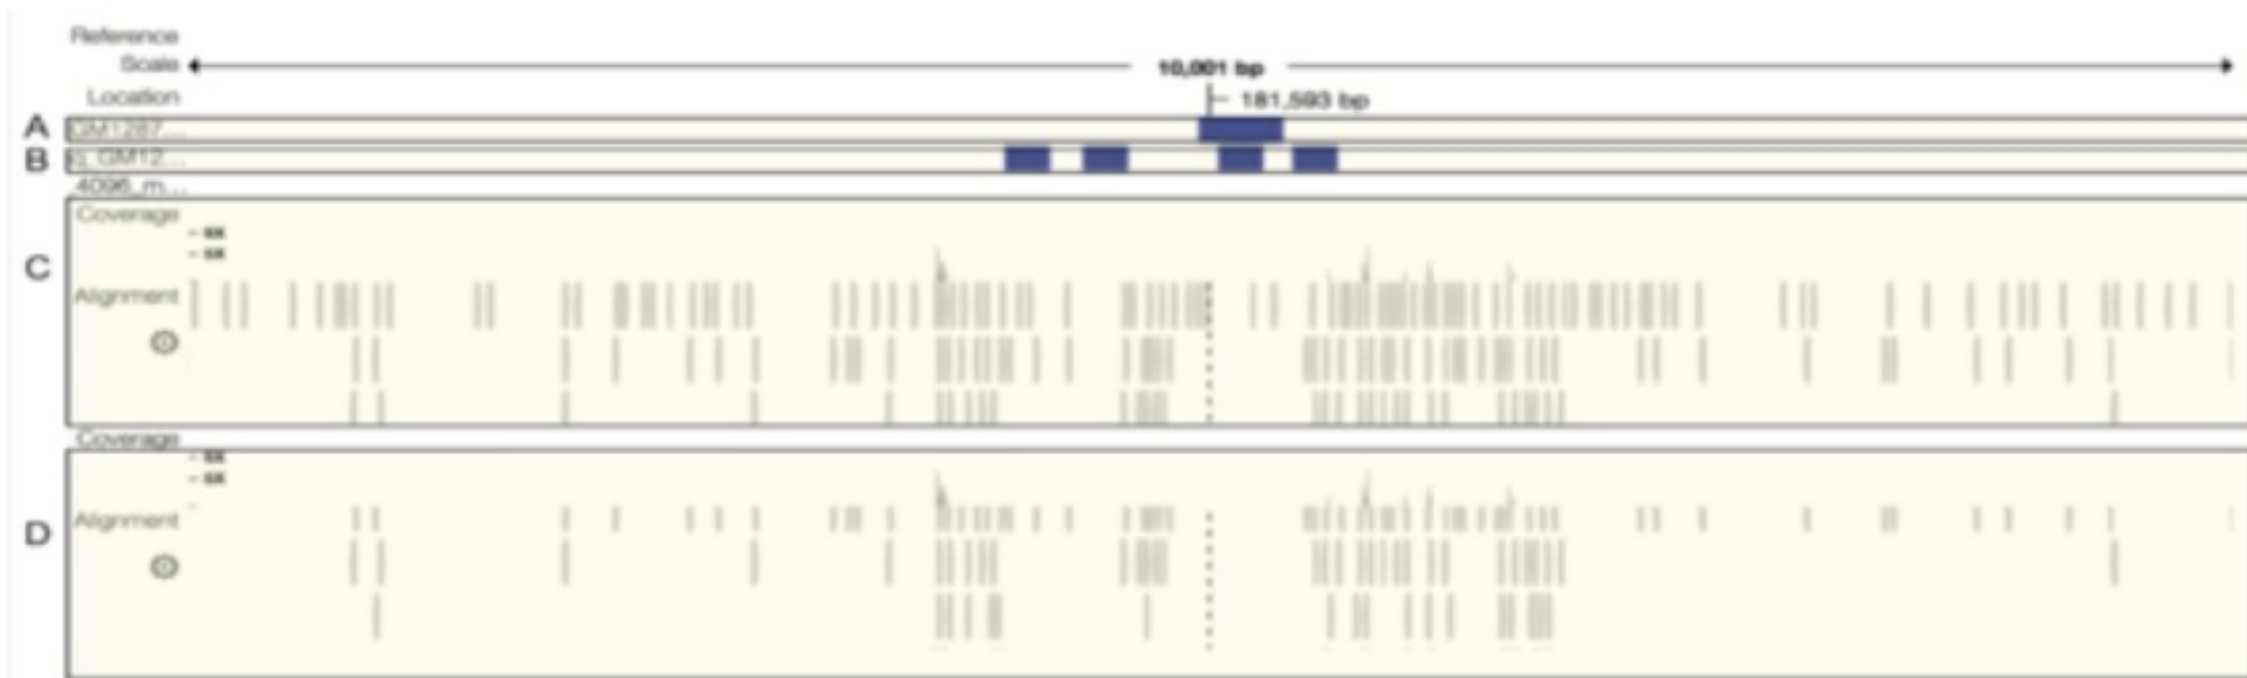
Optimizing for genomic EDA



Narrow waist in stack → can swap in/out levels of stack

- For interactive queries against genomic loci, swap in RDD implementation optimized for point/range queries
- Persistent store optimizations minimize initial overhead for fetching raw data
- Memory optimizations minimize latency for genomic range queries

Benefit of stack: intersection of technologies



- Apply the model interactively to a new dataset using Mango, use join query to overlap “ground truth” data against predictions
- 10kbp query+apply latency: ~400ms

Ongoing work: variant warehousing

- Reads yield genotypes, but we're often interested in statistical aggregates across genotypes:
 - Probability of seeing a genotype in a population
 - Probability of a genotype associating with a phenotype
- Data typically is arriving (near) continuously

Demonstrating Incremental Update in Gnocchi

- Problem: want to compute associations between genotypes and phenotypes (linear/logistic regression)
- Solution: Incremental update of many small GnocchiModels
 - Train each distributed model using standard methods
 - When new data added, build locally optimized model on new data and merge resulting model with old model (do not need old data!)
- Work in progress:
 - Requires periodic recomputes over entire cohort to remain close to full recompute solution
 - Can limit the number of recomputes by being smart about haplotype blocks

Acknowledgements

- **UC Berkeley:** Matt Massie, Timothy Danford, André Schumacher, Jey Kottalam, Karen Feng, Eric Tu, Alyssa Morrow, Niranjana Kumar, Ananth Pallaseni, Michael Heuer, Justin Paschall, Taner Dagdelen, Devin Petersohn, Anthony D. Joseph, Dave Patterson
- **Mt. Sinai:** Arun Ahuja, Neal Sidhwaney, Ryan Williams, Michael Linderman, Jeff Hammerbacher, Uri Laserson
- **GenomeBridge:** Carl Yeksigian
- **Cloudera:** Uri Laserson, Tom White
- **Microsoft Research:** Ravi Pandya, Bill Bolosky
- **UC Santa Cruz:** Benedict Paten, David Haussler, Hannes Schmidt, Beau Norgeot, Audrey Musselman-Brown, John Vivian
- And many other open source contributors, especially Neil Ferguson, Andy Petrella, Xavier Tordior, Deborah Siegel, Denny Lee
- Over 60 contributors to ADAM/BDG from >12 institutions

Thank You.

Check out the code: <https://github.com/bigdatagenomics>

Check out a demo: <https://databricks.com/blog/2016/05/24/genome-sequencing-in-a-nutshell.html>

Run ADAM in Databricks CE: <http://goo.gl/xK8x7s>

