

Real-time analytical query processing and predictive model building on high dimensional document datasets with timestamps

Debasish Das

Data & Artificial Intelligence, Verizon

Contributors

Algorithm: Santanu Das, Zhengming Xing

Platform: Ponrama Jegan

Frontend: AltaffShaik, Jon Leonhardt



© Verizon 2017 All Rights Reserved

Information contained herein is provided AS IS and subject to change without notice. All trademarks used herein are property of their respective owners.



Data Overview

- Location data
 - Each srcip defined as unique row key
 - Provides approximate location of each srcip
 - Timeseries containing latitude, longitude, error bound, duration, timezone for each srcip
- Clickstream data
 - Contains clickstream data of each row key
 - Contains startTime, duration, httpHost, httpURI, upload/download bytes, httpMethod
 - Compatible with IPFIX/Netflow formats



Marketing Analytics

- Aggregate Anonymous analysis for insights
- Spark Summit Europe 2016



- Spark Summit East 2017



Data Model

- Schema: srcip, timestmap, tld, zip, tldvisits, zipvisits
- Dense dimension, dense measure
 - Data: 10.1.13.120, d1H2, [company1.com](#), 94555, 2, 4
- Sparse dimension, dense measure
 - Data: 10.1.13.120, d1, {[company1.com](#), [company2.com](#)}, {94555, 94301}, 10, 15
- Sparse dimension, sparse measure
 - Data: 10.1.13.120, d1, {[company1.com](#), [company2.com](#)}, {94555, 94301}, {[company1.com](#):4, [company2.com](#):6}, {94555:8, 94301:7}
- Timestamp optional
- Competing technologies: PowerDrill, Druid, LinkedIn Pinot, Essbase



Lucene Document Mapping

- Example

Schema: srcip, timestamp, tld, zip, tldvisits, zipvisits

Data: 10.1.13.120, d1, {[company1.com](#), [company2.com](#)}, 94555, 10, 15

Data: 10.1.13.120, d4, {[company1.com](#), [company3.com](#)}, 94301, 12, 8

- DataFrame Row to Lucene Document mapping

schema	Row	Document	OLAP
srcip	StringType	Stored	Measure
timestamp	TimestampType	Stored	Dimension
tld	ArrayType[StringType]	Indexed + Stored	Dimension
zip	StringType	Indexed + Stored	Dimension
tld/zipvisits	IntegerType	Stored	Measure

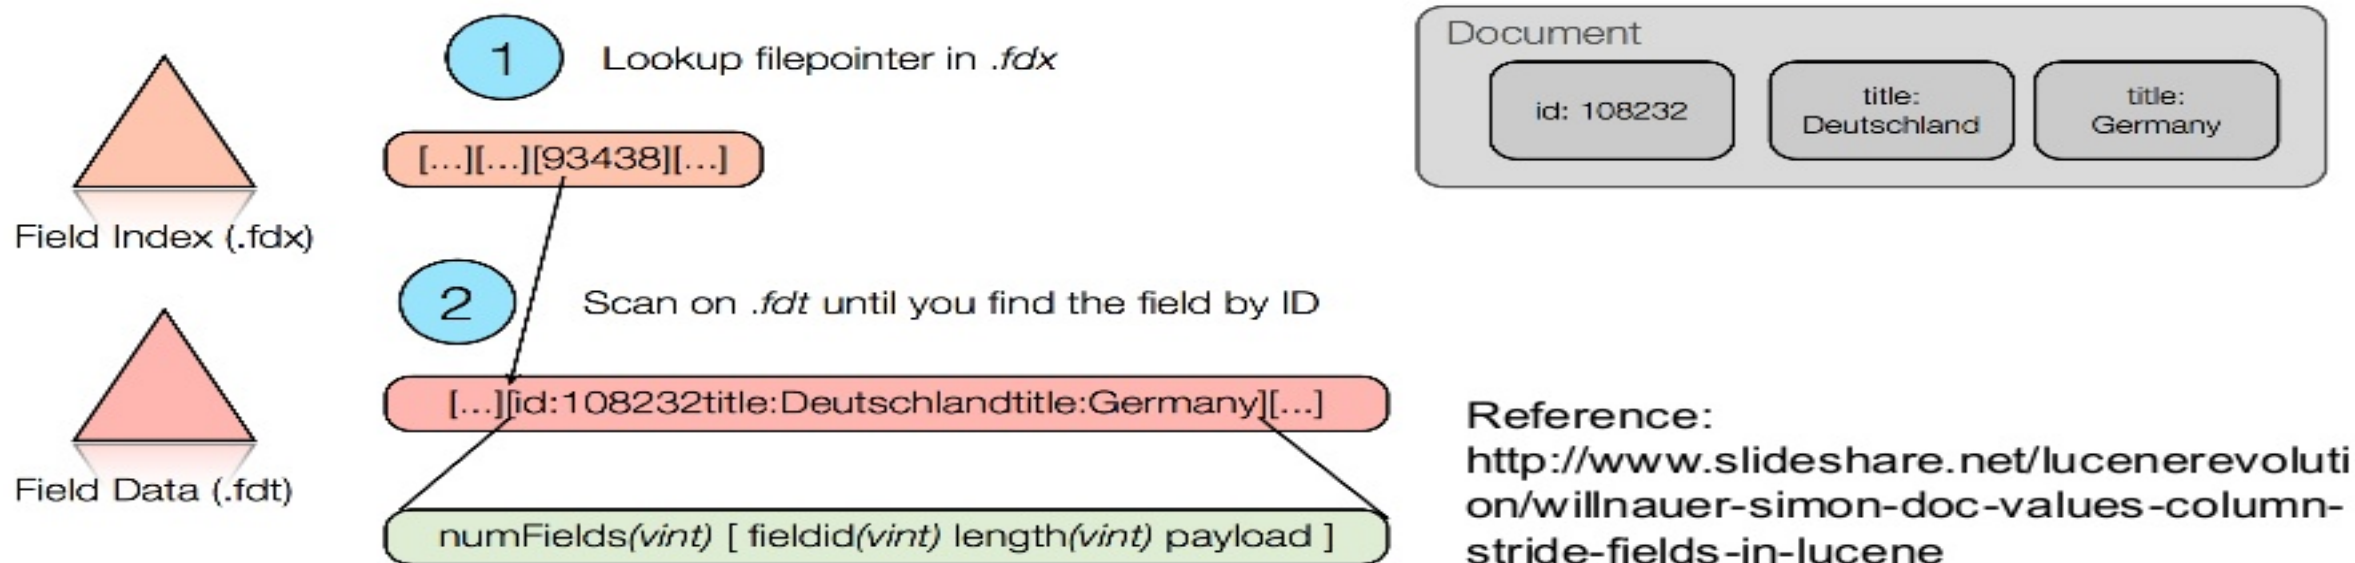


© Verizon 2017 All Rights Reserved

Information contained herein is provided AS IS and subject to change without notice. All trademarks used herein are property of their respective owners.

Lucene Storage

- Row storage: Spark Summit Europe 2016
 - 2 indirect disk seeks for retrieval




Lucene Column Store

- Column storage: Spark Summit East 2017
 - References: LUCENE-3108, LUCENE-2935, LUCENE-2168, LUCENE-1231
 - Cache friendly column retrieval: 1 direct disk seek
 - Integer column: Min-Max encoding
 - Numeric column: Uncompressed
 - Binary column: Referenced
 - Complex Type: Binary + Kryo

field: id
1
5
3
4
6
9
8
7
12
14
22
32
100
33
34
35
36
37
38

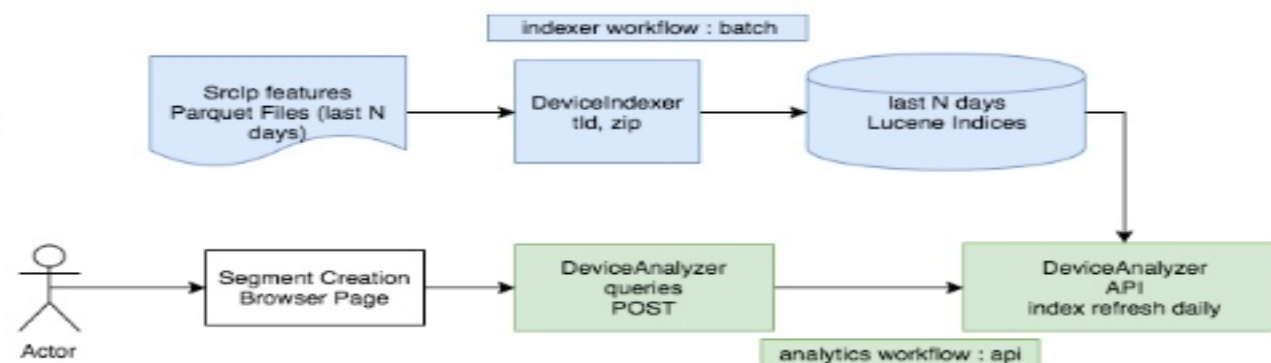
Integer



Binary

DeviceAnalyzer

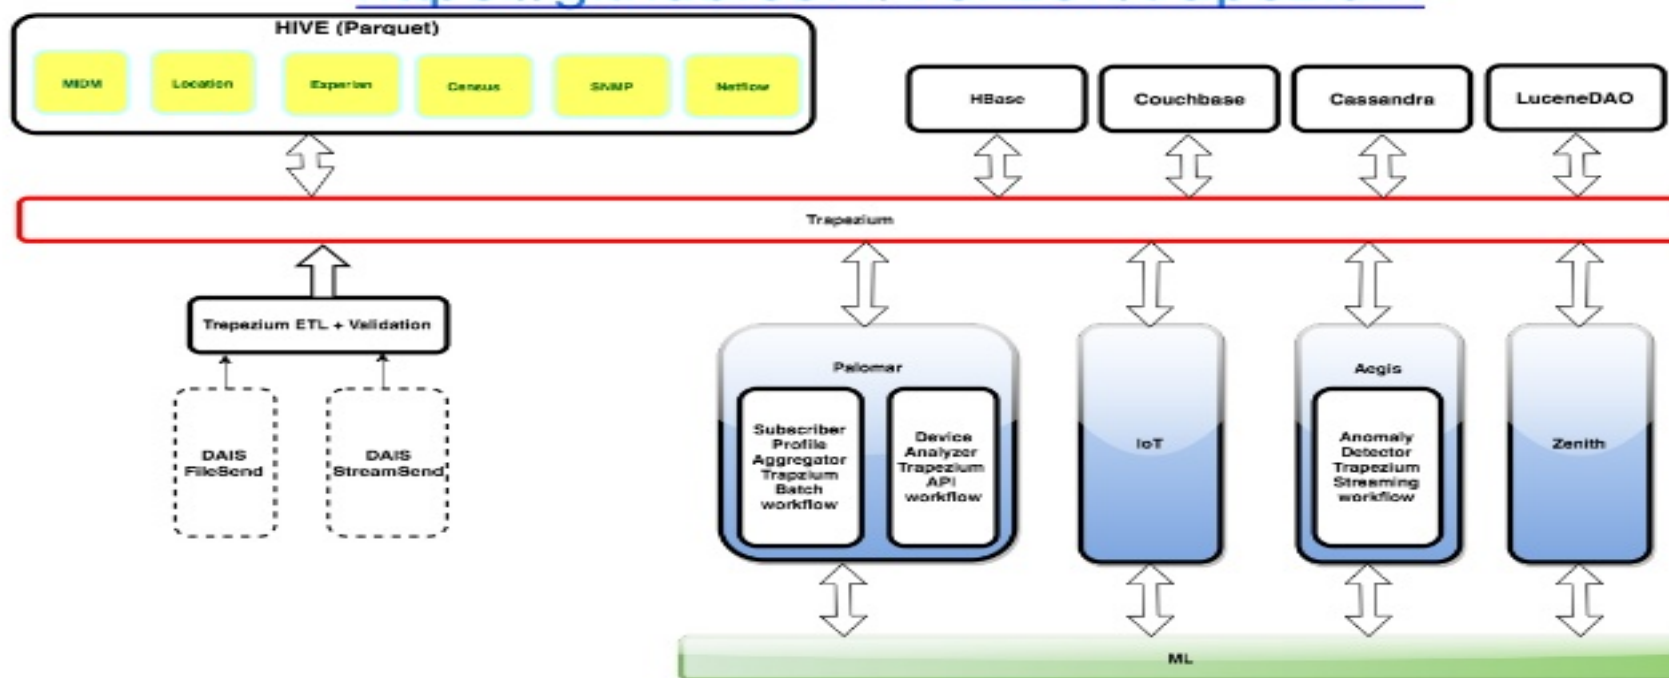
- Goals
 - srcip/visits as dense measure
 - Real-Time queries
 - Aggregate
 - Group
 - Timeseries
 - Real-Time Timeseries forecast



Trapezium

DAIS Open Source framework to build batch, streaming and API services

<https://github.com/Verizon/trapezium>



© Verizon 2017 All Rights Reserved

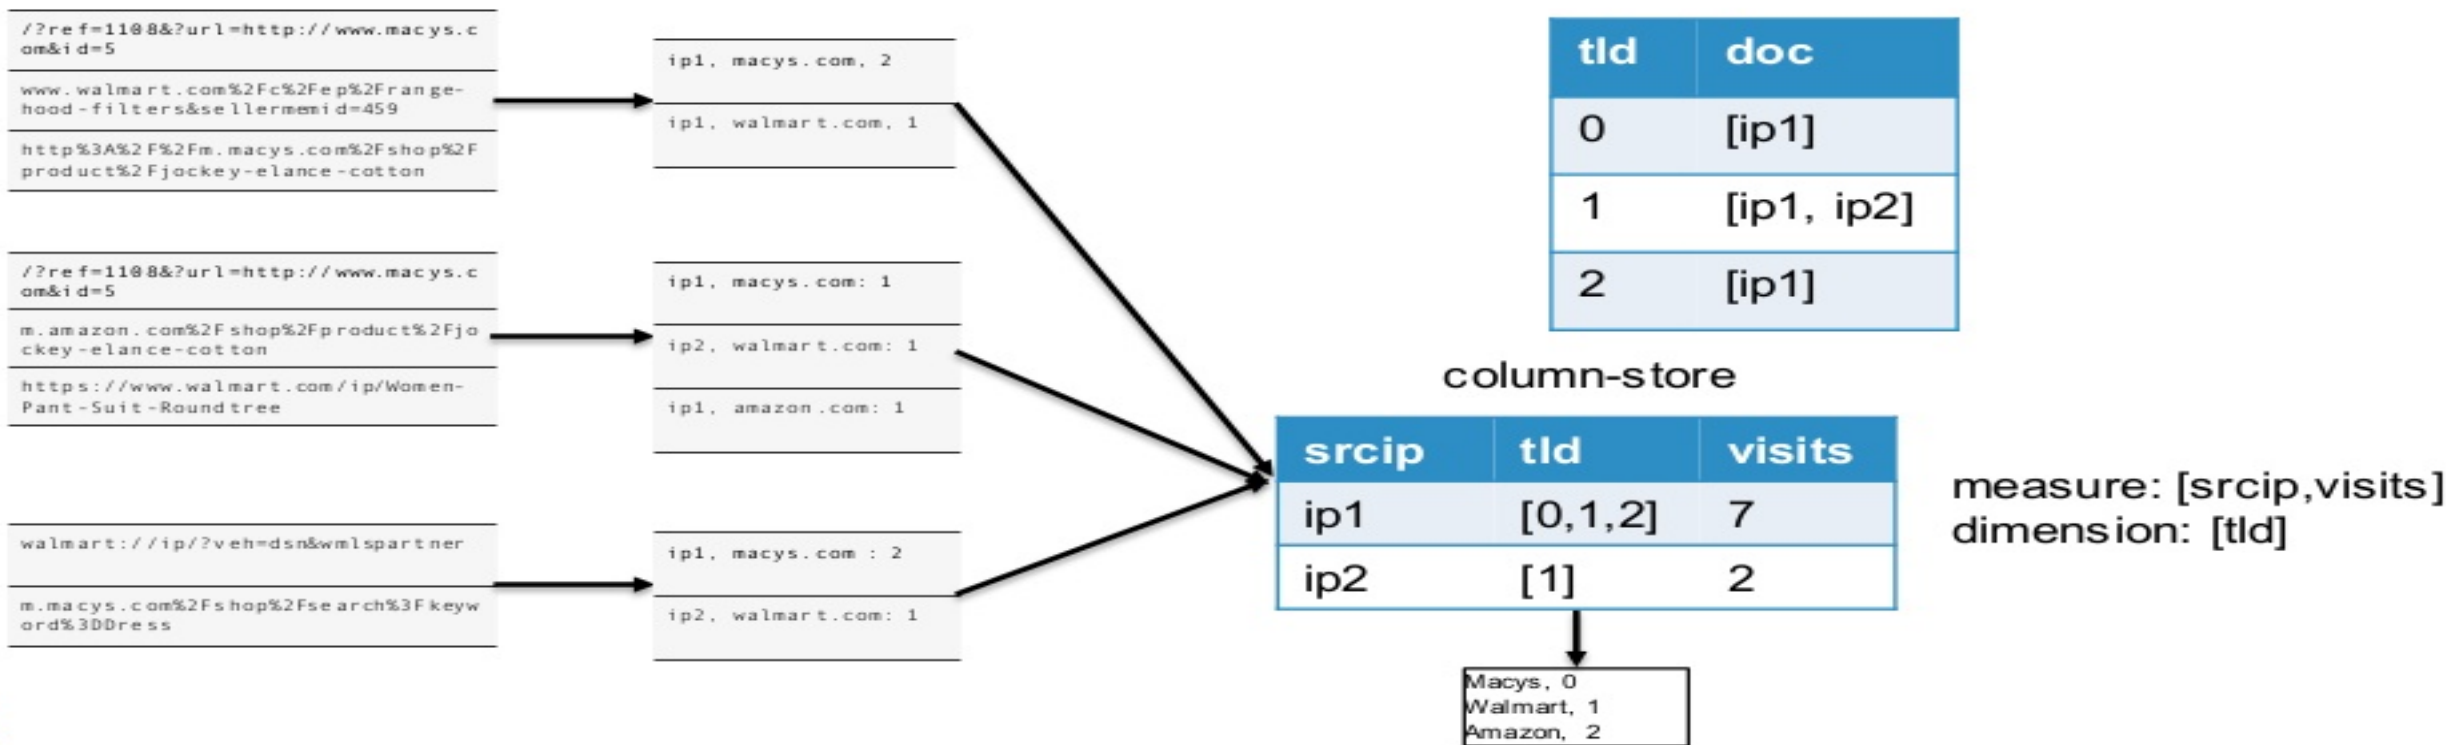
Information contained herein is provided AS IS and subject to change without notice. All trademarks used herein are property of their respective owners.

Trapezium LuceneDAO

- SparkSQL optimized for full scan
 - Column indexing not supported
- Fulfills Real-Time requirements for OLAP queries
- Lucene for indexing + storage per executor
- Spark operators for distributed aggregation
 - treeAggregate
 - mapPartition + treeReduce
- Features
 - Build Distributed Lucene Shards from Dataframe
 - Access saved shards through LuceneDAO for Analytics + ML pipelines
 - Save shards to HDFS for QueryProcessor like SolrCloud



LuceneDAO Indexing



© Verizon 2017 All Rights Reserved

Information contained herein is provided AS IS and subject to change without notice. All trademarks used herein are property of their respective owners.

LuceneDAO API

Index Creation

```
import trapezium.dal.lucene._
import org.apache.spark.sql.types._

object DeviceIndexer extend BatchTransaction {
  process(dfs: Map[String, DataFrame], batchTime: Time): {
    df = dfs("DeviceStore")
    olapDf = rollup(df)
  }

  persist(df: DataFrame, batchTime: Time): {
    val dimensions = Set("tld", "zip")
    val types = Map("tld" -> LuceneType(true, StringType),
                    "srcip" -> LuceneType(false, StringType),
                    "visits" -> LuceneType(false, IntegerType))
    val dao = new LuceneDAO("path", dimension, types)
    dao.index(df, new Time(batchTime))
  }
}
```

Query Processing

```
import trapezium.dal.lucene._
import org.apache.spark.sql.types._
```

Load:

```
val dimensions = Set("tld", "zip")
val types = Map("tld" -> LuceneType(true, StringType),
                "srcip" -> LuceneType(false, StringType),
                "visits" -> LuceneType(false, IntegerType))
val dao = new LuceneDAO("path", dimension, types)
dao.load(sc)
```

Queries:

```
dao.aggregate(query: String, measure: String, aggregator: String)
dao.group(query: String, dimension: String, measure: String,
          aggregator: String)
dao.timeseries(query: String, minTime: Long, maxTime: Long,
               rollup: Long, measure: String, aggregator:
String)
dao.search(query: String, columns: Seq[String]): DataFrame
```

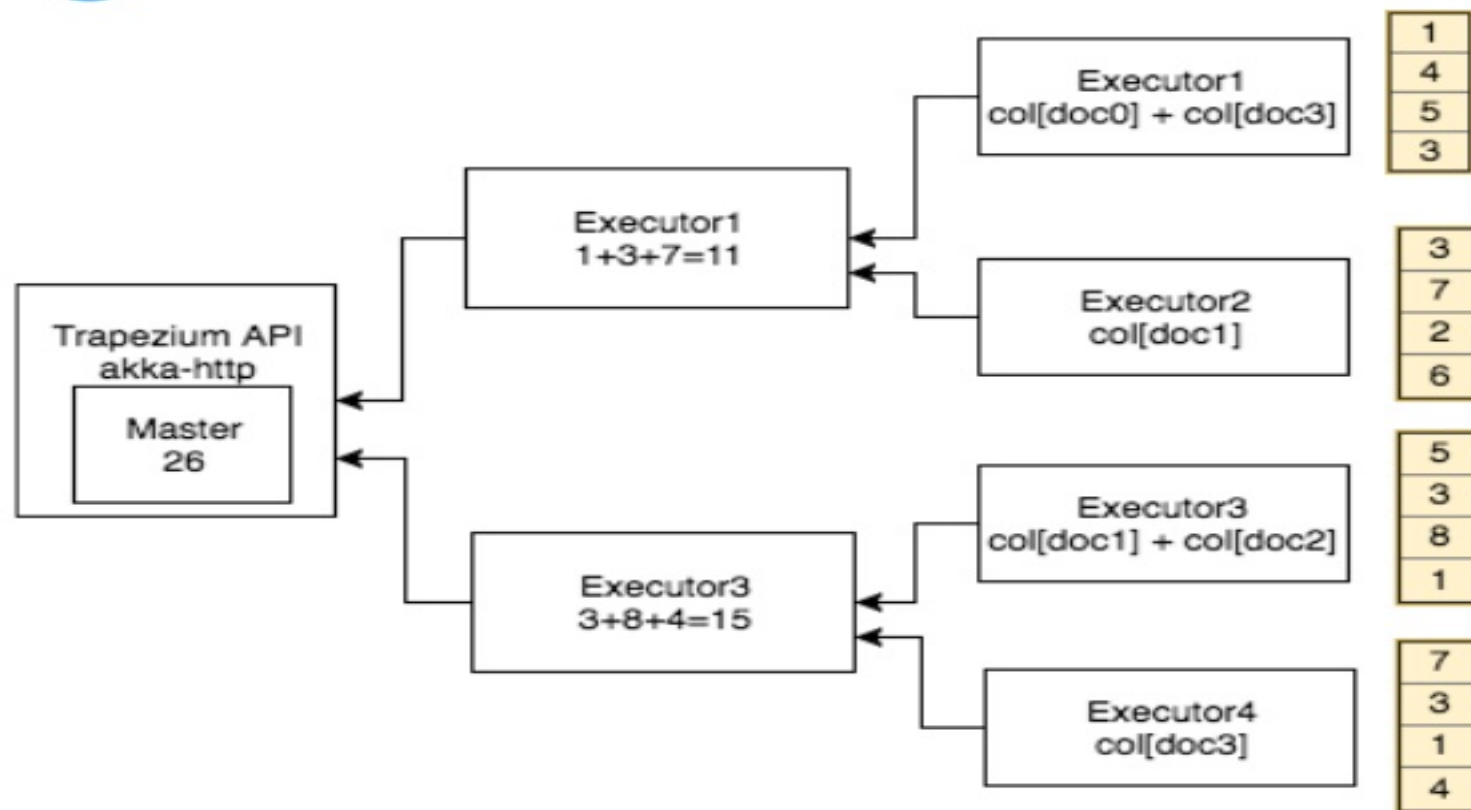


LuceneDAO Internals

- Retrieve documents with/without relevance
- Column Accessor over dimension + measures
- Disk / In-Memory Column Accessor
- C-store style while loops over dimension
- Spark ML style aggregators
- treeAggregate for distributed aggregation



Aggregation Architecture



Index Generation

- Dataset details:
57M devices, 4.2B docs
- Parquet: 79 GB
- Lucene Reverse Index: 16 GB
- Lucene DocValues: 59.6 GB
- Global Dictionary Size: 5.5 MB
- Executors: 20 Cores: 8
- RAM Driver: 16g Executor: 16g
- Runtime
 - Parquet:
 - 1831.87 s
 - Dictionary:
 - 213.7 s
 - Index + Stored:
 - 360 s



Aggregate Queries


- HashSet aggregation

- SparkSQL

```
df.select("srcip","tld")  
.where(array_contains(df("tld"),  
"company1.com"))  
.agg(countDistinct("srcip") as "visits")  
.collect()
```

- LuceneDAO

```
dao.aggregate("tld:company1.com",  
"srcip", "count")
```



Group Queries


- HLL aggregation

- SparkSQL

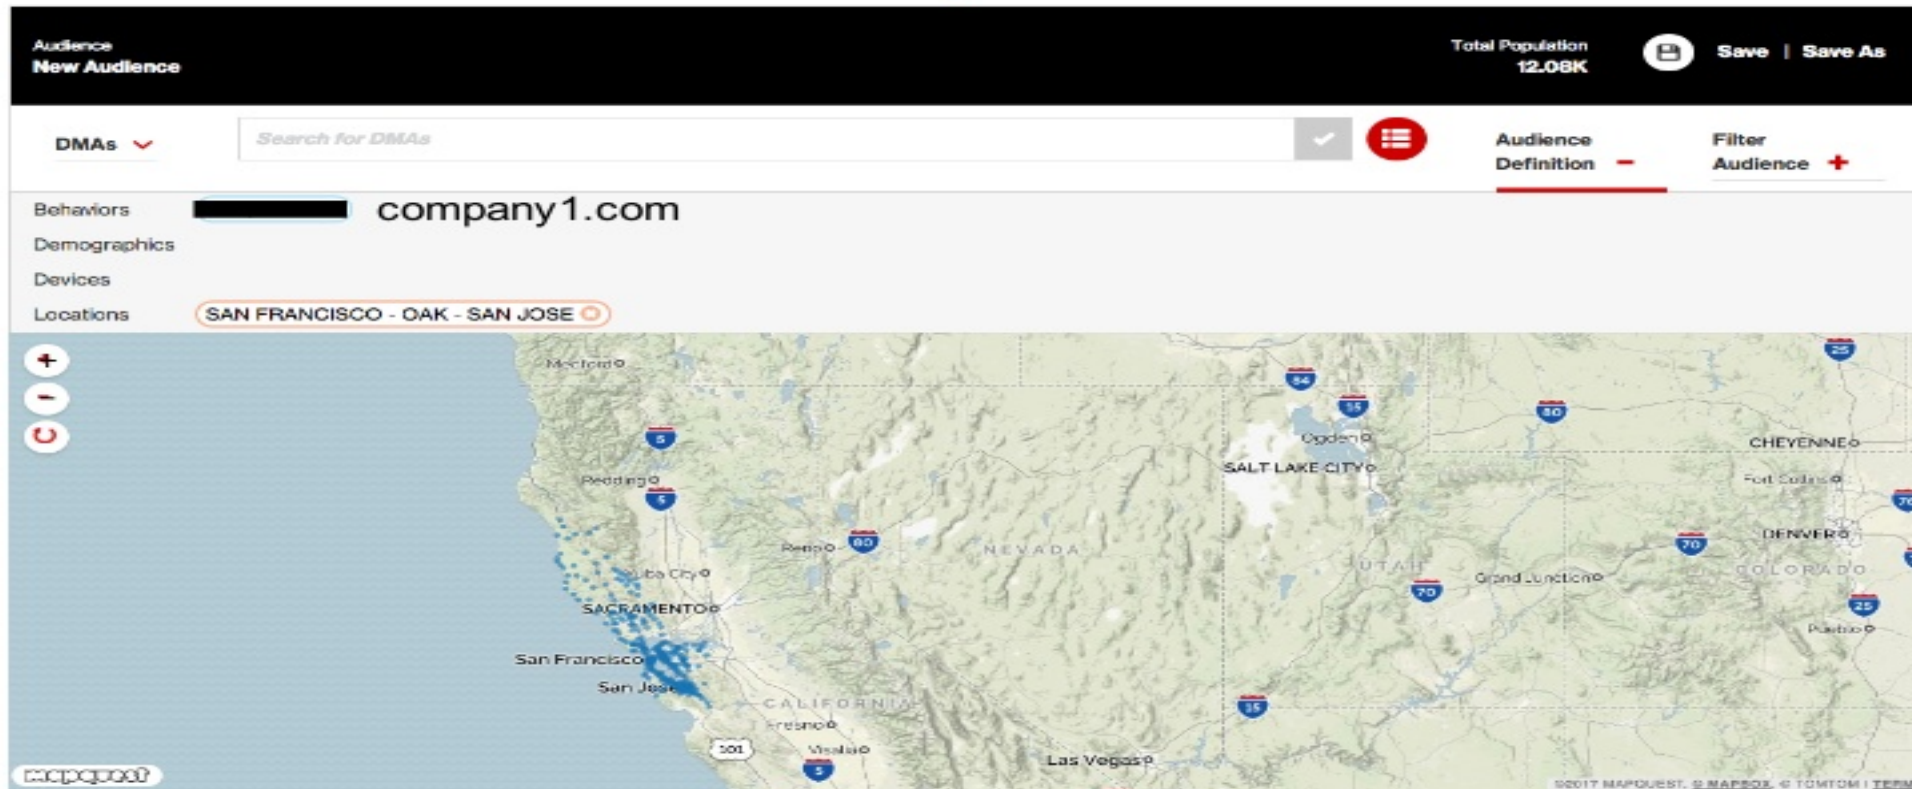
```
df.select("srcip","tld", "zip")  
.where(array_contains(df("tld"),  
"company1.com"))  
.select("zip", "srcip").groupBy("zip")  
.agg(approxCountDistinct("srcip") as  
"visits")  
.collect()
```

- LuceneDAO

```
dao.aggregate("tld:company1.com", "srcip",  
"count")
```



Device Heat-Map



Timeseries Queries


- HLL aggregation

- SparkSQL

```
df.select("time","srcip","tld")  
.where(array_contains(df("tld"),  
"company1.com"))  
.select("time", "srcip").groupBy("time")  
.agg(approxCountDistinct("srcip") as "visits")  
.collect()
```

- LuceneDAO

```
dao.aggregate("tld:company1.com", "srcip",  
"count")
```



TimeSeries Forecast

Trapezium ML

- Given a query:
select
timestamp, (srcip) as deviceCount
where
tld='company1.com' AND state='CA'
- Predict deviceCount for next
timestamp
- Forecast deviceCount for next N
timestamps

TimeSeriesKNNRegression.predict

Input:

timeseries: Array[Double]

topk: Int

featureDim: Int

normalize: Boolean

multiStep: Int

metric: KernelType=Euclidean

Output:

predicted values: Array[Double]



© Verizon 2017 All Rights Reserved

Information contained herein is provided AS IS and subject to change without notice. All trademarks used herein are property of their respective owners.

Forecast Service

Powered by Trapezium API

```
httpServer = {  
  provider = "akka"  
  hostname = "localhost"  
  port = 19999  
  contextPath = "/"  
  endPoints = [{  
    path = "analyzer-api"  
    className =  
    "TimeseriesEndPoint"  
  }]  
}
```


```
class TimeseriesEndPoint(sc: SparkContext)  
extends SparkServiceEndPoint(sc) {  
  override def route : timeseriesRoute  
  
  val types = Map("tld" -> LuceneType(true, StringType),  
                  "srcip" -> LuceneType(false, StringType),  
                  "visits" -> LuceneType(false, IntegerType))  
  val dao = new LuceneDAO("path", dimension, types)  
  dao.load(sc)  
  
  def timeseriesRoute : {  
    post { request => {  
      ts = dao.timeseries(request, minTime, maxTime, rollout,  
                           "srcip", "count_approx")  
      predicted = TimeseriesKNNRegression.predict(ts, topk=5,  
                                                    featureDim=3, normalize=false, multiStep=5,  
                                                    metric=Euclidean)  
      generateResponse(ts, predicted)  
    }  
  }  
}
```





© Verizon 2017 All Rights Reserved


Information contained herein is provided AS IS and subject to change without notice. All trademarks used herein are property of their respective owners.


Device-Count Forecast


DMA 

Search for DMAs




Audience Definition 

Filter Audience 

Behaviors  company1.com

Demographics

Devices

Locations SAN FRANCISCO - OAK - SAN JOSE 

Monthly Time Series



5 step prediction



SPARK
SUMMIT
EAST 2017

© Verizon 2017 All Rights Reserved

Information contained herein is provided AS IS and subject to change without notice. All trademarks used herein are property of their respective owners.

Thank You.

Q&A

Join us and make machines intelligent
Data & Artificial Intelligence Systems
499 Hamilton Ave, Palo Alto
California



© Verizon 2017 All Rights Reserved

Information contained herein is provided AS IS and subject to change without notice. All trademarks used herein are property of their respective owners.

