

TRANSFORMATIONS AND ACTIONS

<http://training.databricks.com/visualapi.pdf>



A Visual Guide of the API



[LinkedIn](#)

Blog: [data-frack](#)


Databricks would like to give a special thanks to Jeff Thomspon for contributing 67 visual diagrams depicting the Spark API under the MIT license to the Spark community.

Jeff's original, creative work can be found [here](#) and you can read more about Jeff's project in his [blog post](#).

After talking to Jeff, Databricks commissioned [Adam Breindel](#) to further evolve Jeff's work into the diagrams you see in this deck.



- Founded in late 2013
- by the creators of Apache Spark
- Original team from UC Berkeley AMPLab
- Raised \$47 Million in 2 rounds
- ~55 employees
- We're hiring! (<http://databricks.workable.com>)
- Level 2/3 support partnerships with
 - Hortonworks
 - MapR
 - DataStax



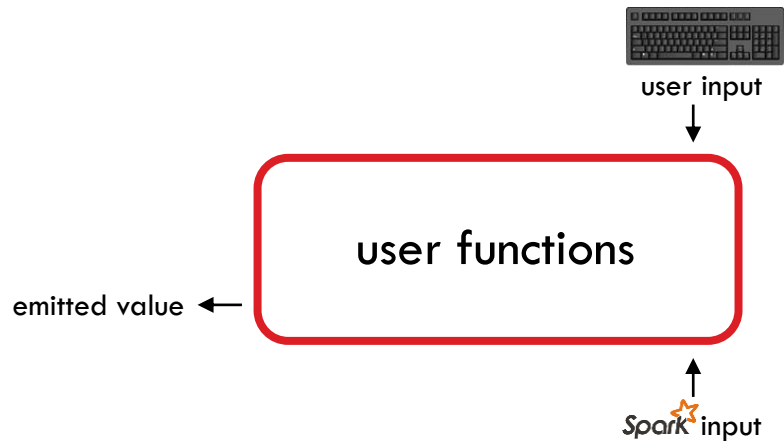
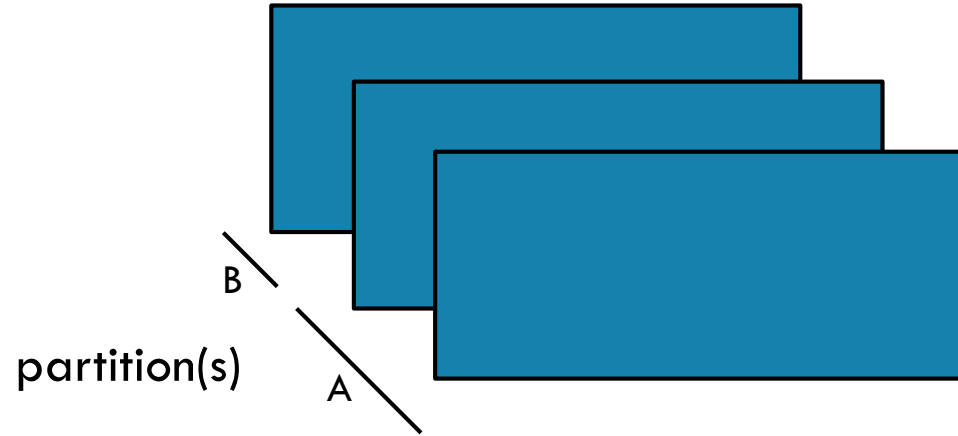
Databricks Cloud:

“A unified platform for building Big Data pipelines – from ETL to Exploration and Dashboards, to Advanced Analytics and Data Products.”

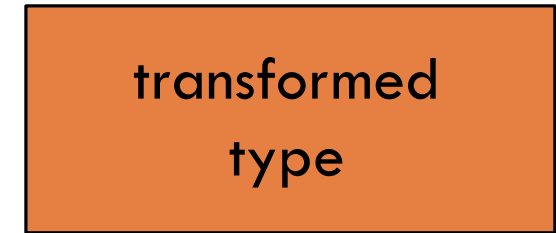
RDD



Legend




RDD Elements

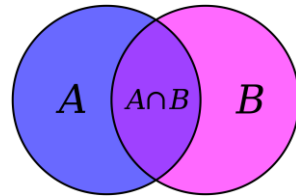




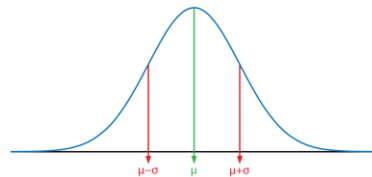
Legend



Randomized operation



Set Theory / Relational operation



Numeric calculation

Spark  Operations =



TRANSFORMATIONS

+



ACTIONS

 = easy

 = medium

Essential Core & Intermediate Spark Operations

TRANSFORMATIONS

General

- map
- filter
- flatMap
- mapPartitions
- mapPartitionsWithIndex
- groupBy
- sortBy

Math / Statistical

- sample
- randomSplit

Set Theory / Relational

- union
- intersection
- subtract
- distinct
- cartesian
- zip

Data Structure / I/O

- keyBy
- zipWithIndex
- zipWithUniqueId
- zipPartitions
- coalesce
- repartition
- repartitionAndSortWithinPartitions
- pipe

ACTIONS

- reduce
- collect
- aggregate
- fold
- first
- take
- foreach
- top
- treeAggregate
- treeReduce
- foreachPartition
- collectAsMap


- count
- takeSample
- max
- min
- sum
- histogram
- mean
- variance
- stdev
- sampleVariance
- countApprox
- countApproxDistinct

- takeOrdered

- saveAsTextFile
- saveAsSequenceFile
- saveAsObjectFile
- saveAsHadoopDataset
- saveAsHadoopFile
- saveAsNewAPIHadoopDataset
- saveAsNewAPIHadoopFile



 = easy

 = medium

Essential Core & Intermediate PairRDD Operations

TRANSFORMATIONS

General

- flatMapValues
- groupByKey
- reduceByKey
- reduceByKeyLocally
- foldByKey
- aggregateByKey
- sortByKey
- combineByKey

Math / Statistical

- sampleByKey

Set Theory / Relational

- cogroup (=groupWith)
- join
- subtractByKey
- fullOuterJoin
- leftOuterJoin
- rightOuterJoin

Data Structure

- partitionBy

ACTIONS

- keys
- values

- countByKey
- countByValue
- countByValueApprox
- countApproxDistinctByKey
- countApproxDistinctByKey
- countByKeyApprox
- sampleByKeyExact





vs

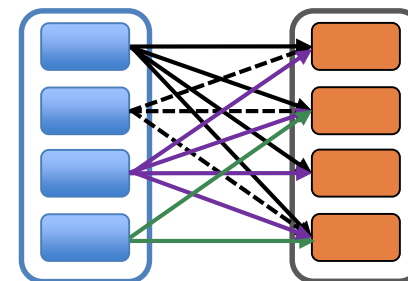
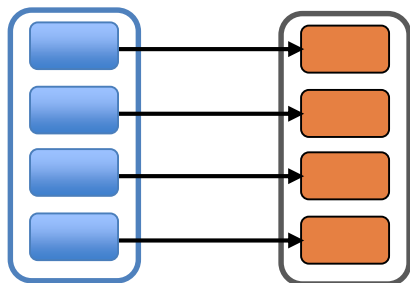


narrow

wide

*each partition of the parent RDD is used by
at most one partition of the child RDD*

*multiple child RDD partitions may depend
on a single parent RDD partition*



LINEAGE

“One of the challenges in providing RDDs as an abstraction is choosing a representation for them that can track lineage across a wide range of transformations.”

“The most interesting question in designing this interface is how to represent dependencies between RDDs.”


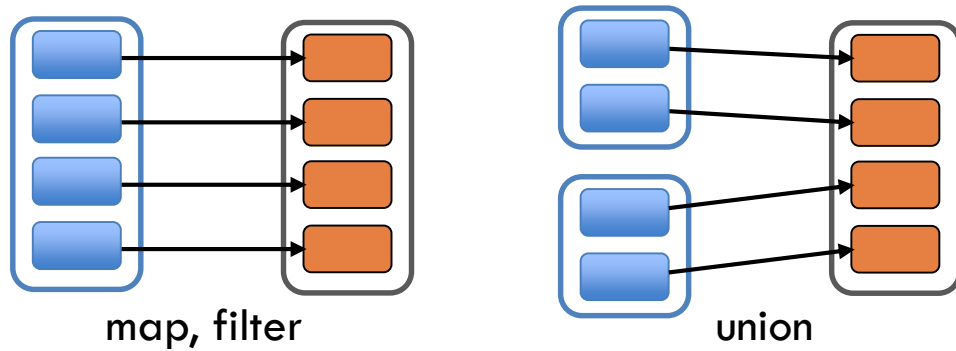
“We found it both sufficient and useful to classify dependencies into two types:

- **narrow dependencies**, where each partition of the parent RDD is used by at most one partition of the child RDD
- **wide dependencies**, where multiple child partitions may depend on it.”



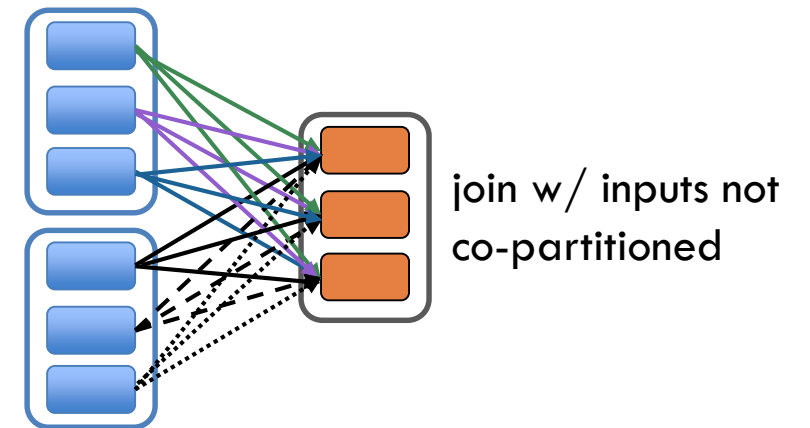
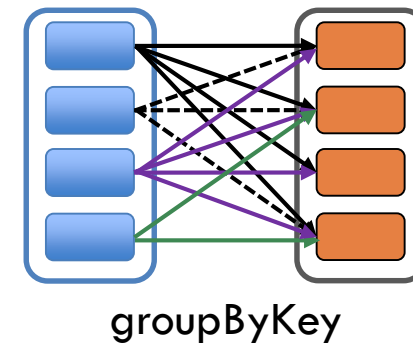
narrow

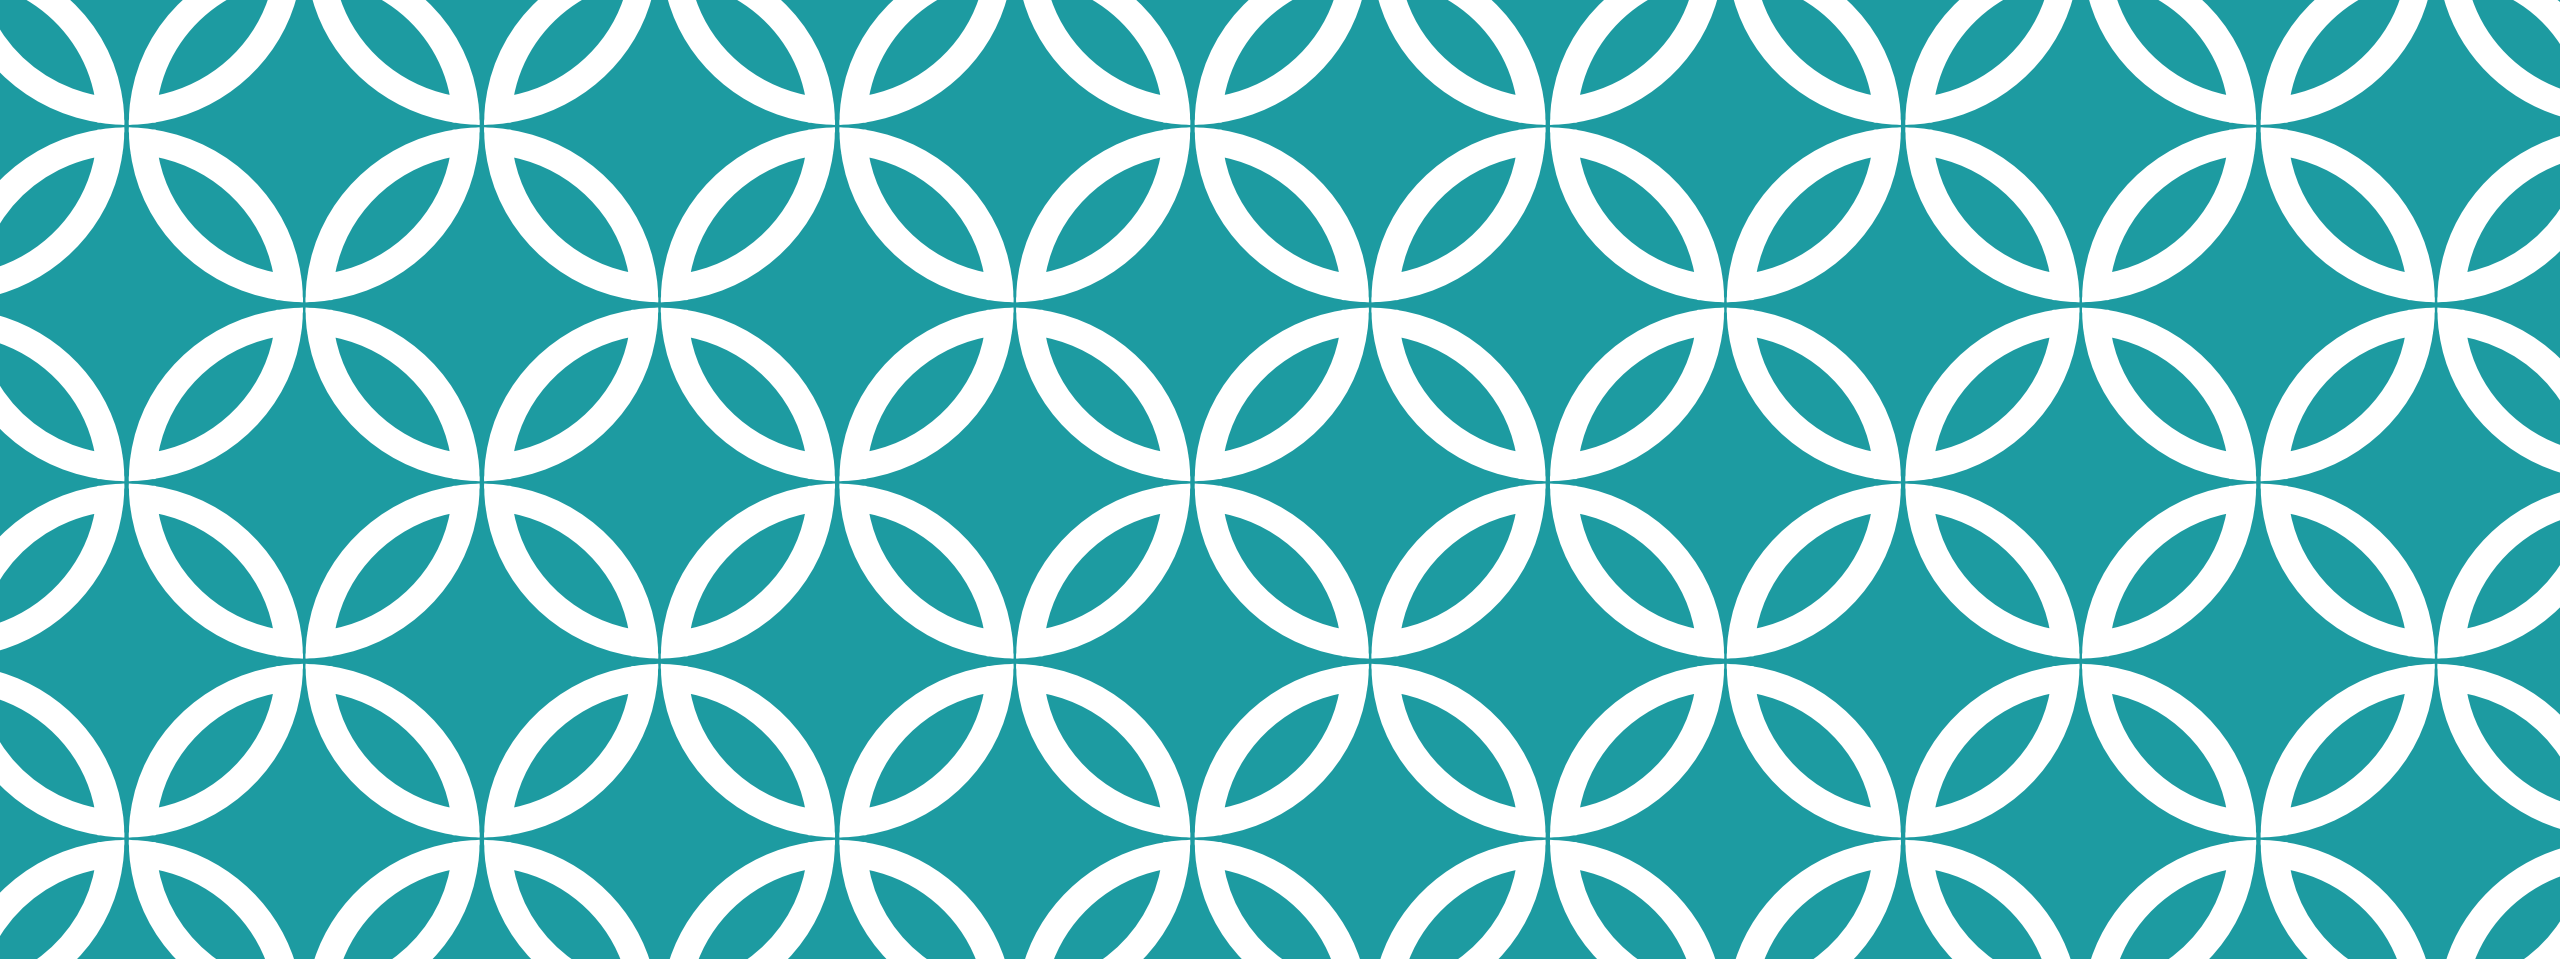
each partition of the parent RDD is used by at most one partition of the child RDD




wide

multiple child RDD partitions may depend on a single parent RDD partition





TRANSFORMATIONS




Core Operations



MAP

RDD: x

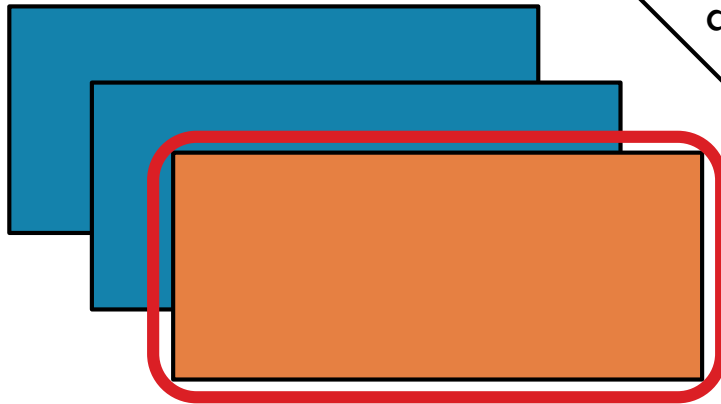
3 items in RDD





MAP

RDD: **x**



User function
applied item by item

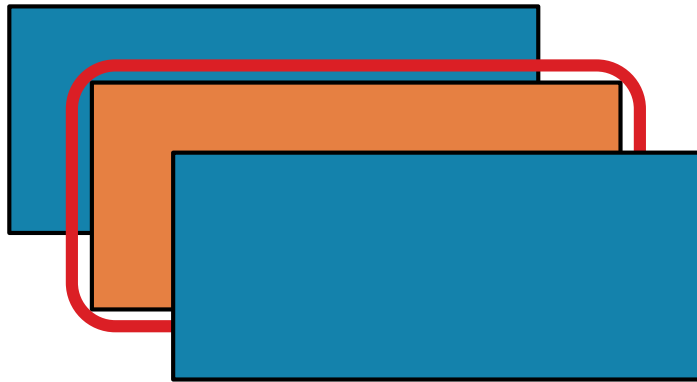
RDD: **y**





MAP

RDD: **x**



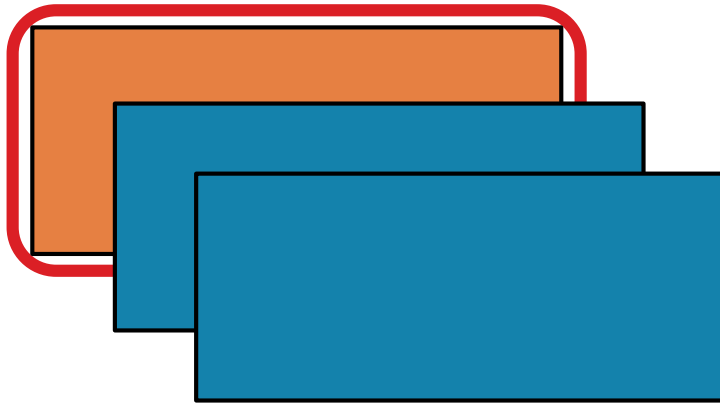
RDD: **y**



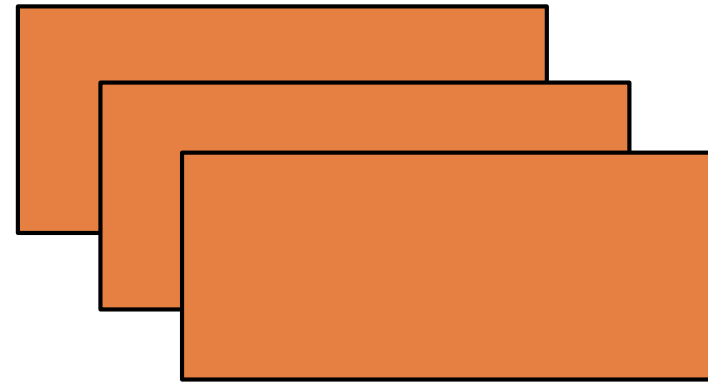


MAP

RDD: **x**



RDD: **y**

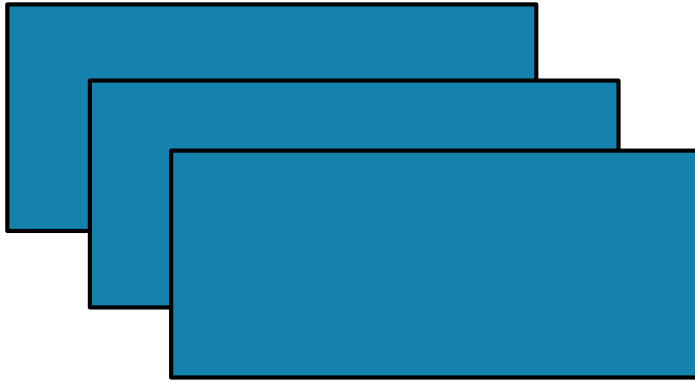




After map() has been applied...


MAP

RDD: **x**



before

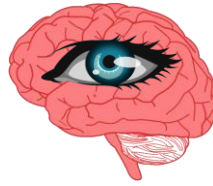
RDD: **y**




after




MAP



RDD: **x**



RDD: **y**




Return a new RDD by applying a function to each element of this RDD.




MAP

RDD: **x**



RDD: **y**



`map(f, preservesPartitioning=False)`

Return a new RDD by applying a function to each element of this RDD



```
x = sc.parallelize(["b", "a", "c"])
y = x.map(lambda z: (z, 1))
print(x.collect())
print(y.collect())
```



x: ['b', 'a', 'c']

y: [('b', 1), ('a', 1), ('c', 1)]



```
val x = sc.parallelize(Array("b", "a", "c"))
val y = x.map(z => (z,1))
println(x.collect().mkString(", "))
println(y.collect().mkString(", "))
```

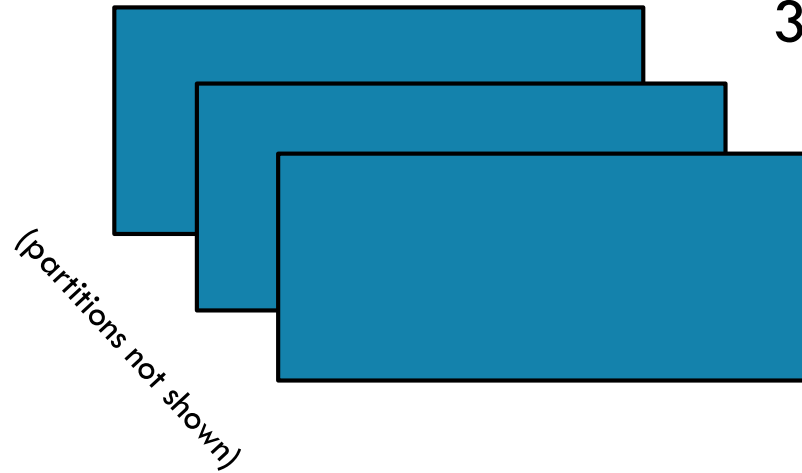




FILTER

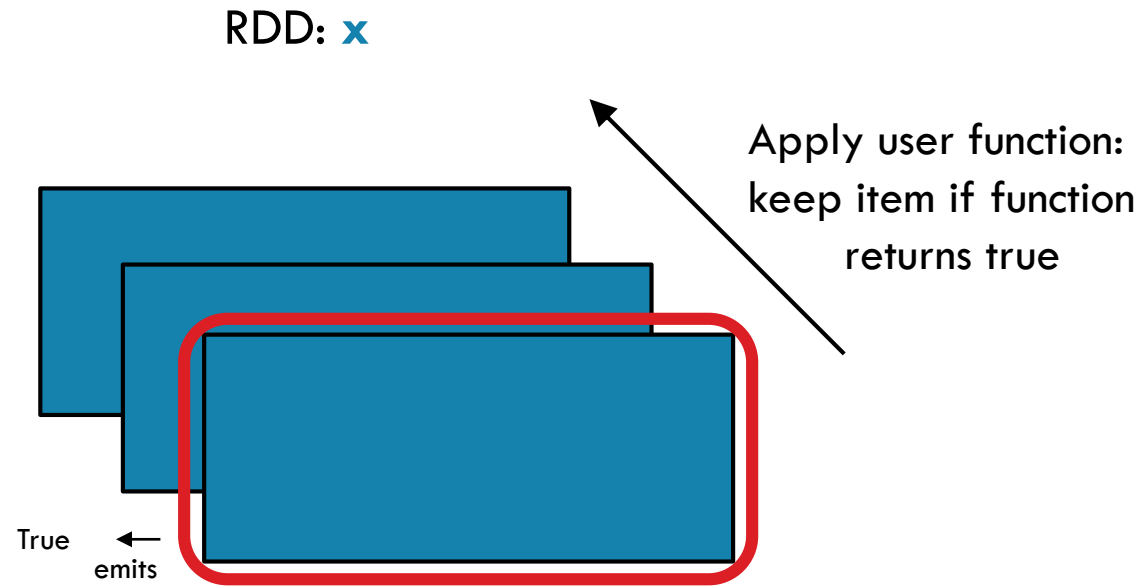
RDD: x

3 items in RDD





FILTER



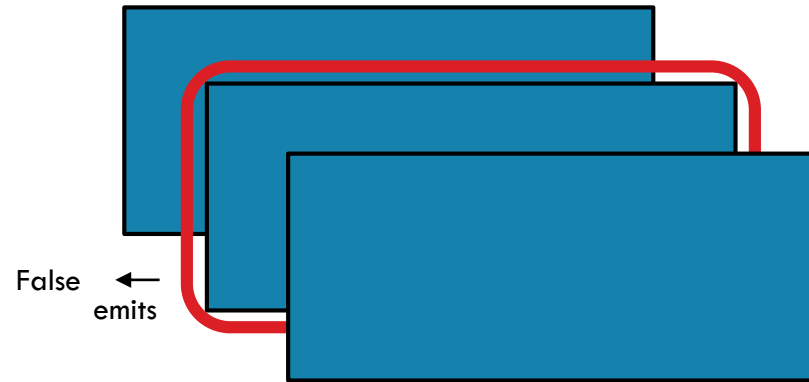
RDD: **y**





FILTER

RDD: **x**



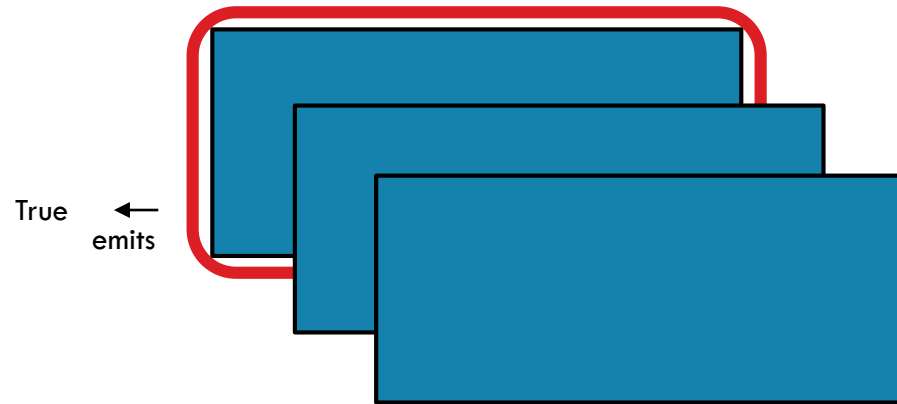
RDD: **y**



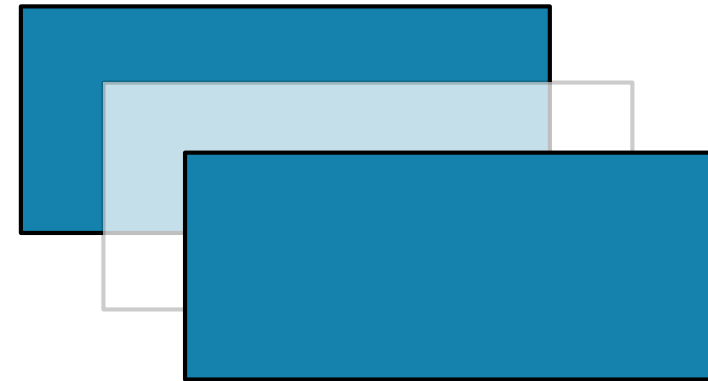


FILTER

RDD: **x**



RDD: **y**





After `filter()` has been applied...

FILTER

RDD: **x**



before

RDD: **y**

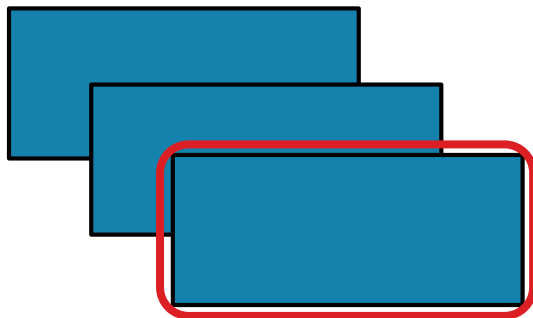


after



FILTER

RDD: **x**



RDD: **y**



`filter(f)`

Return a new RDD containing only the elements that satisfy a predicate



```
x = sc.parallelize([1,2,3])
y = x.filter(lambda x: x%2 == 1) #keep odd values
print(x.collect())
print(y.collect())
```



x: [1, 2, 3]

y: [1, 3]



```
val x = sc.parallelize(Array(1,2,3))
val y = x.filter(n => n%2 == 1)
println(x.collect().mkString(", "))
println(y.collect().mkString(", "))
```

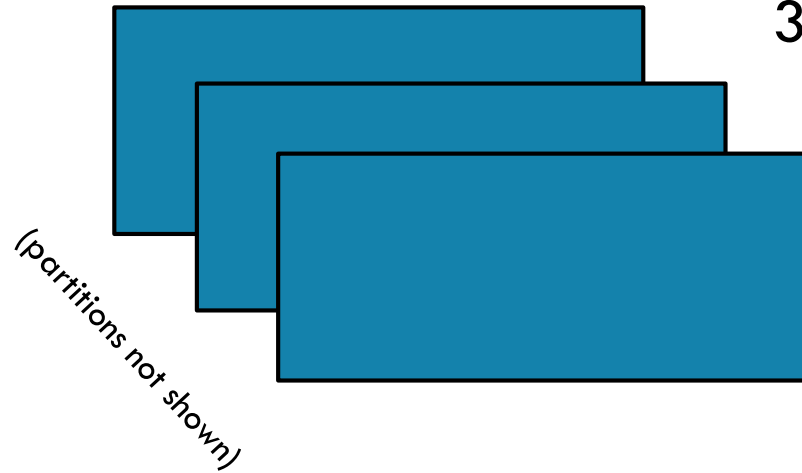




FLATMAP

RDD: **x**

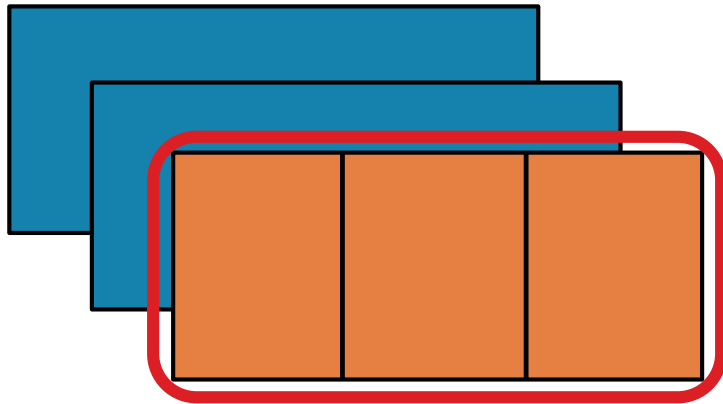
3 items in RDD



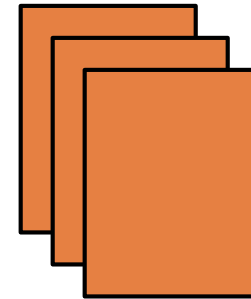


FLATMAP

RDD: **x**



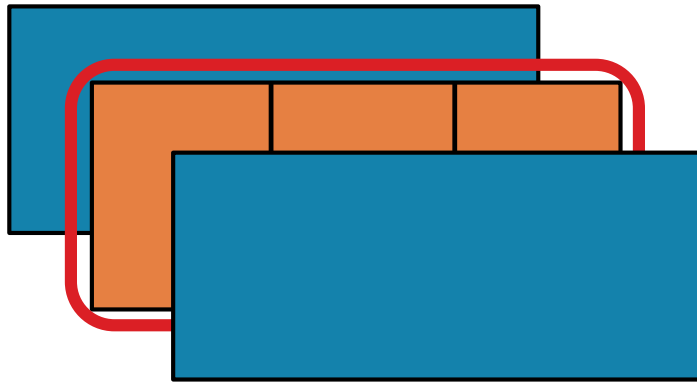
RDD: **y**



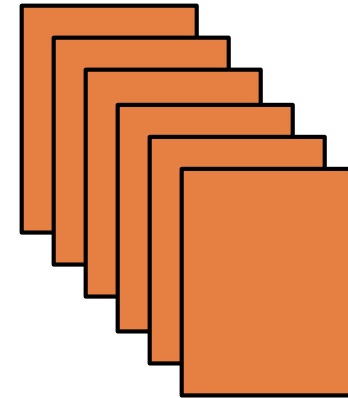


FLATMAP

RDD: **x**



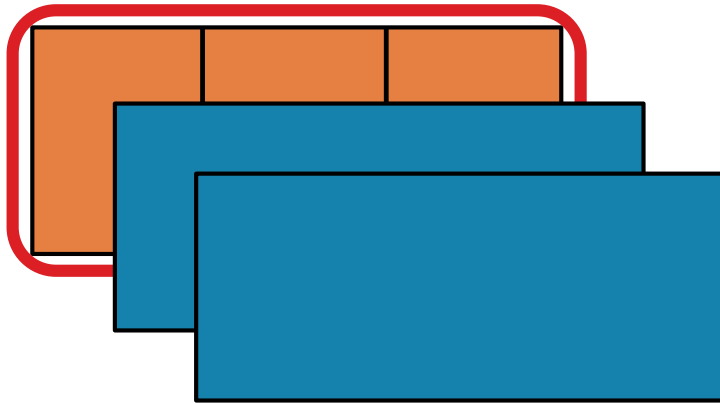
RDD: **y**



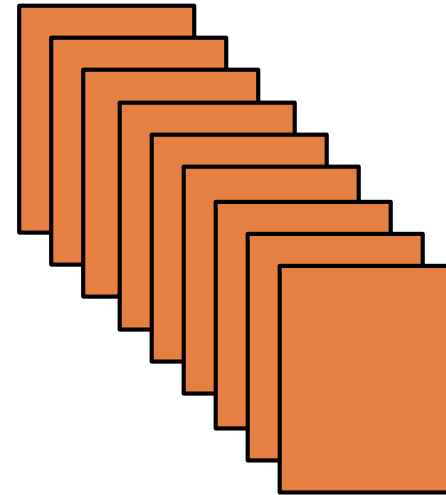


FLATMAP

RDD: **x**



RDD: **y**

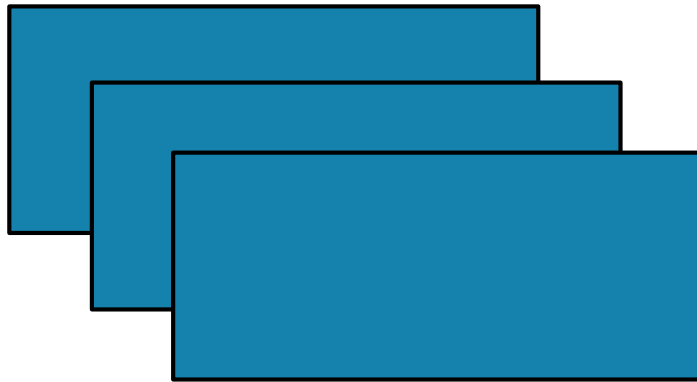




After flatmap() has been applied...

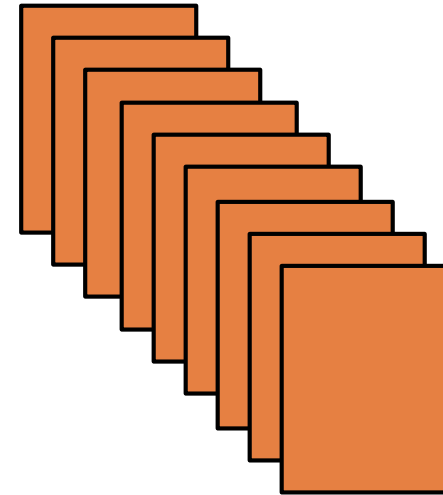
FLATMAP

RDD: **x**



before

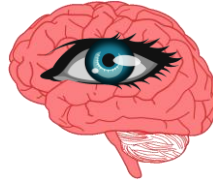
RDD: **y**



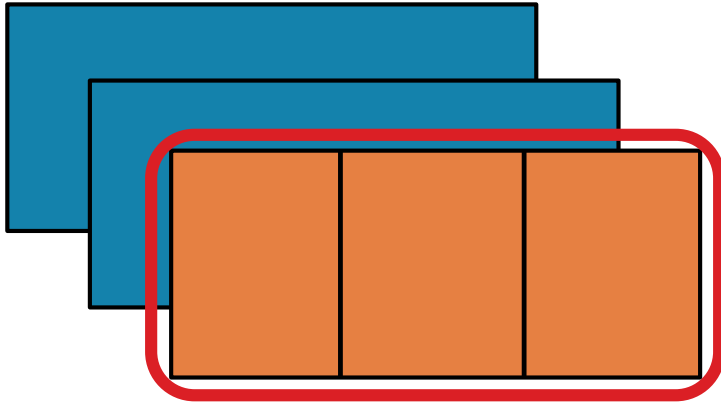
after



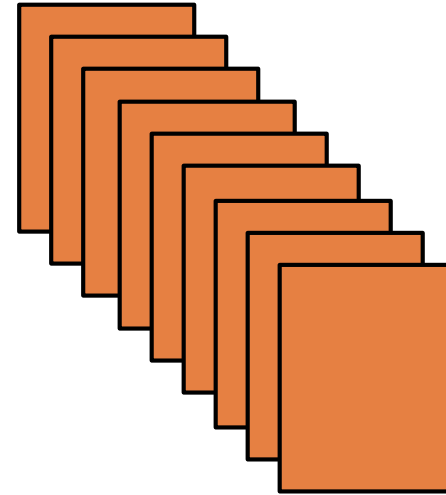
FLATMAP



RDD: **x**



RDD: **y**




Return a new RDD by first applying a function to all elements of this RDD, and then flattening the results



FLATMAP

RDD: **x**



RDD: **y**



`flatMap(f, preservesPartitioning=False)`

Return a new RDD by first applying a function to all elements of this RDD, and then flattening the results



```
x = sc.parallelize([1,2,3])
y = x.flatMap(lambda x: (x, x*100, 42))
print(x.collect())
print(y.collect())
```



x: [1, 2, 3]

y: [1, 100, 42, 2, 200, 42, 3, 300, 42]



```
val x = sc.parallelize(Array(1,2,3))
val y = x.flatMap(n => Array(n, n*100, 42))
println(x.collect().mkString(", "))
println(y.collect().mkString(", "))
```

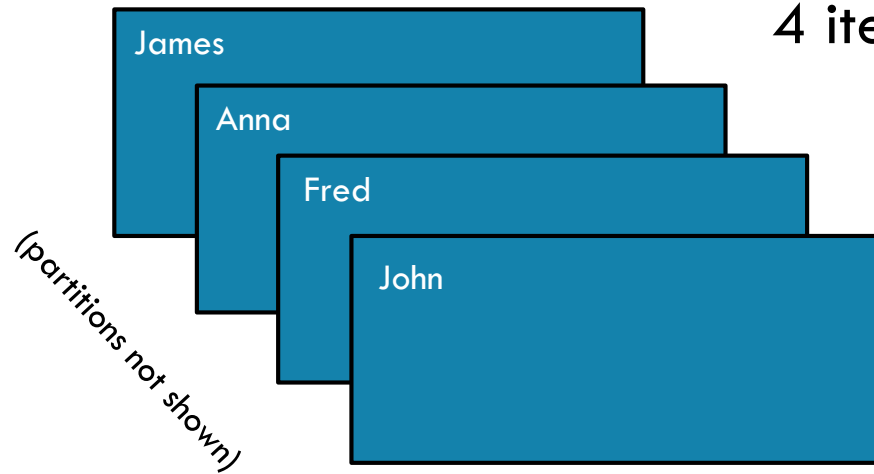




GROUPBY

RDD: x

4 items in RDD

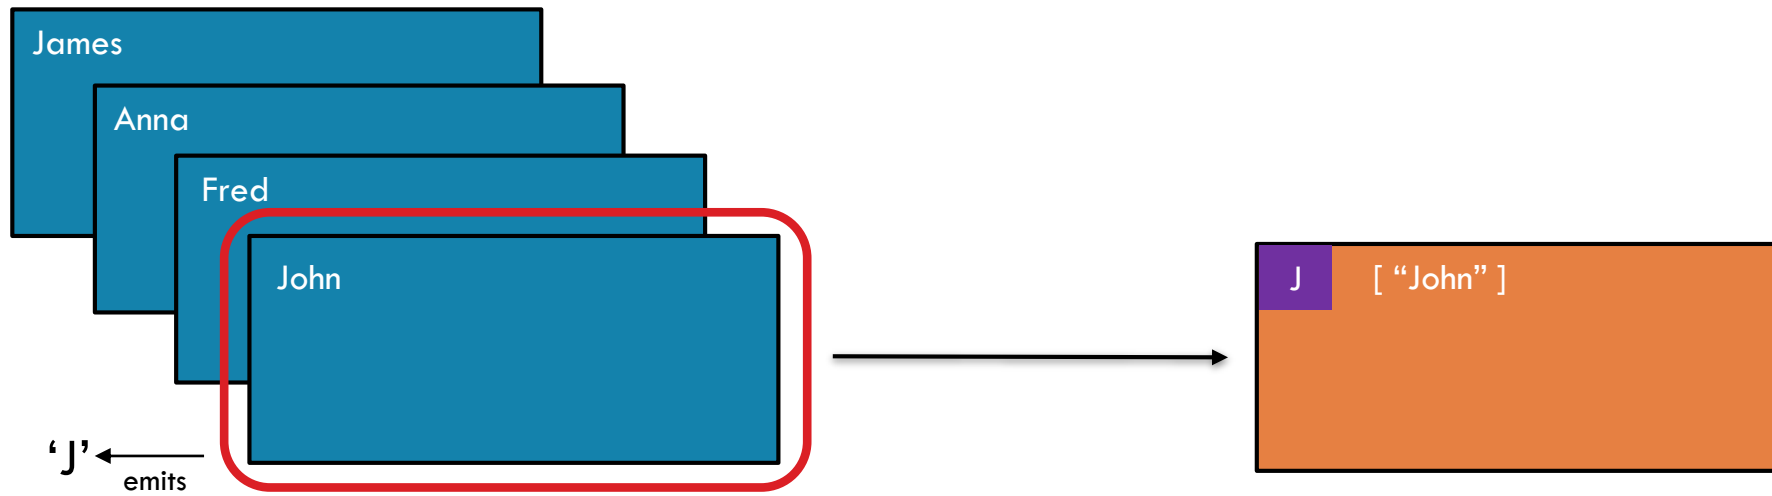




GROUPBY

RDD: **x**

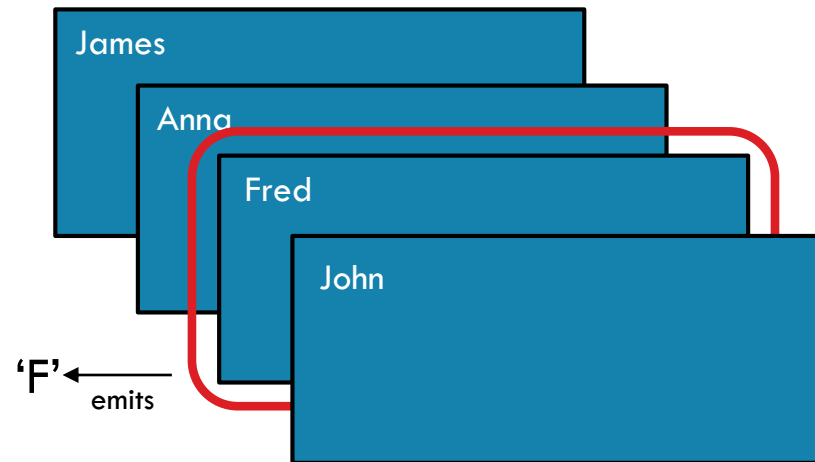
RDD: **y**



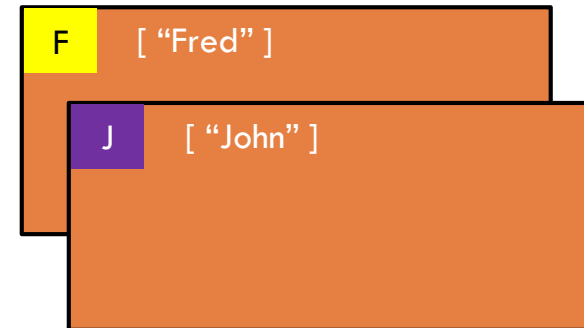


GROUPBY

RDD: **x**




RDD: **y**



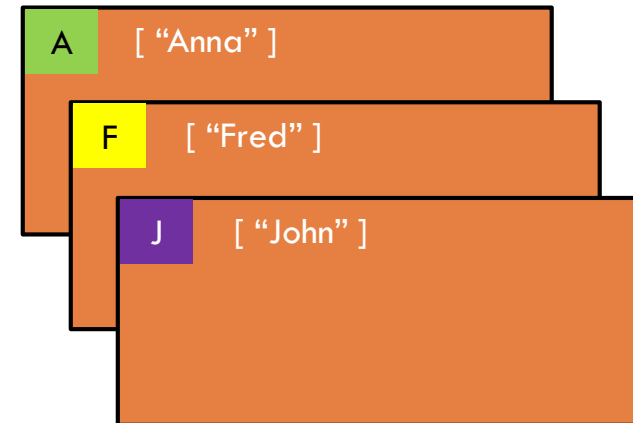


GROUPBY

RDD: **x**



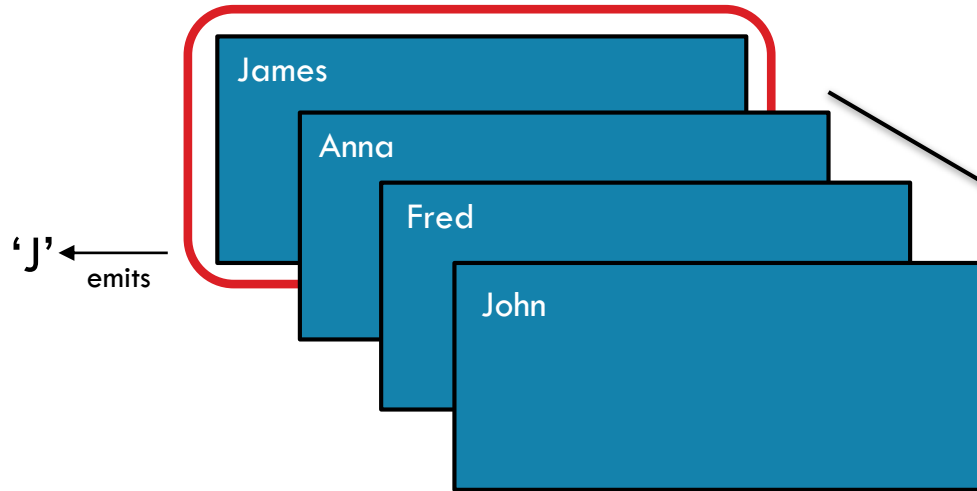
RDD: **y**



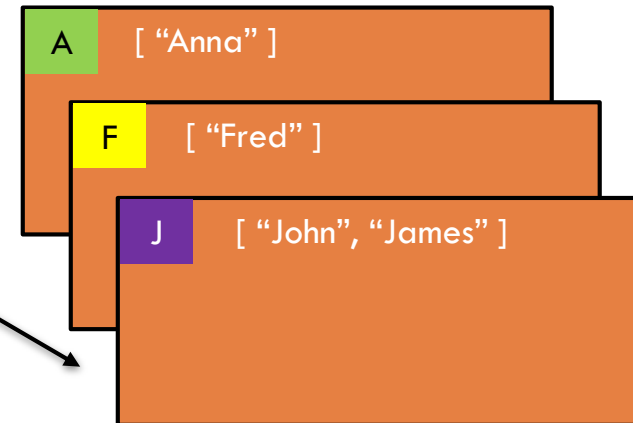


GROUPBY

RDD: **x**

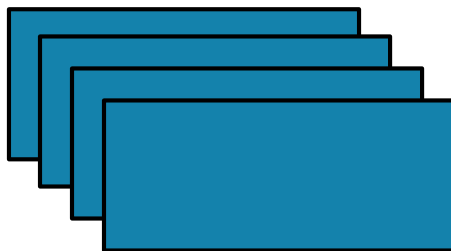


RDD: **y**




GROUPBY

RDD: **x**



RDD: **y**



`groupBy(f, numPartitions=None)`

Group the data in the original RDD. Create pairs where the key is the output of a user function, and the value is all items for which the function yields this key.



```
x = sc.parallelize(['John', 'Fred', 'Anna', 'James'])  
y = x.groupBy(lambda w: w[0])  
print [(k, list(v)) for (k, v) in y.collect()]
```



x: ['John', 'Fred', 'Anna', 'James']

y: [('A', ['Anna']), ('J', ['John', 'James']), ('F', ['Fred'])]




```
val x = sc.parallelize(  
    Array("John", "Fred", "Anna", "James"))  
val y = x.groupBy(w => w.charAt(0))  
println(y.collect().mkString(", "))
```





GROUPBYKEY

Pair RDD: x






GROUPBYKEY

Pair RDD: **x**


RDD: **y**






GROUPBYKEY



Pair RDD: **x**



RDD: **y**



GROUPBYKEY



`groupByKey(numPartitions=None)`

Group the values for each key in the original RDD. Create a new pair where the original key corresponds to this collected group of values.

```
x = sc.parallelize([('B',5),('B',4),('A',3),('A',2),('A',1)])
y = x.groupByKey()
print(x.collect())
print(list((j[0], list(j[1])) for j in y.collect()))
```



x: [('B', 5), ('B', 4), ('A', 3), ('A', 2), ('A', 1)]

y: [('A', [2, 3, 1]), ('B', [5, 4])]

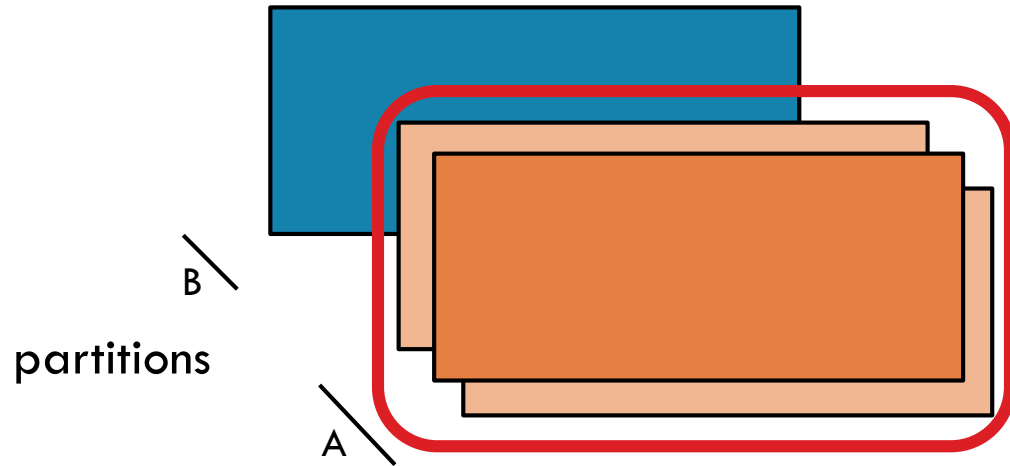
```
val x = sc.parallelize(
    Array(('B',5),('B',4),('A',3),('A',2),('A',1)))
val y = x.groupByKey()
println(x.collect().mkString(", "))
println(y.collect().mkString(", "))
```





MAPPARTITIONS

RDD: **x**



RDD: **y**



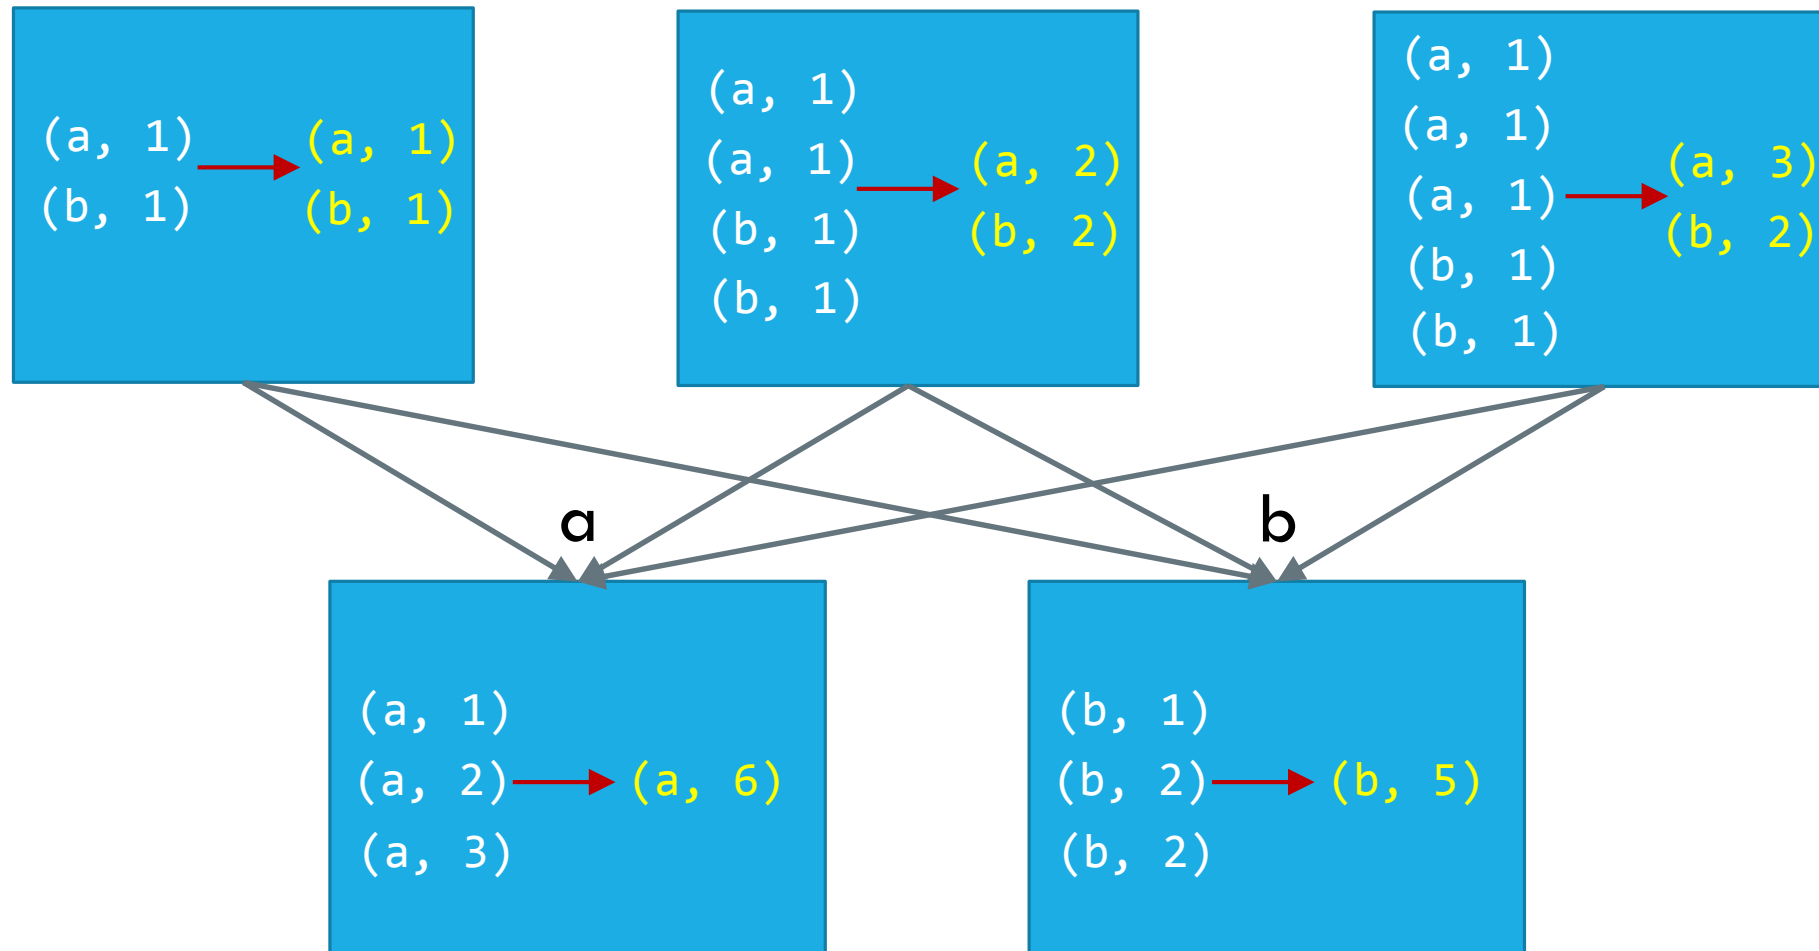
REDUCEBYKEY vs GROUPBYKEY

```
val words = Array("one", "two", "two", "three", "three", "three")  
val wordPairsRDD = sc.parallelize(words).map(word => (word, 1))
```

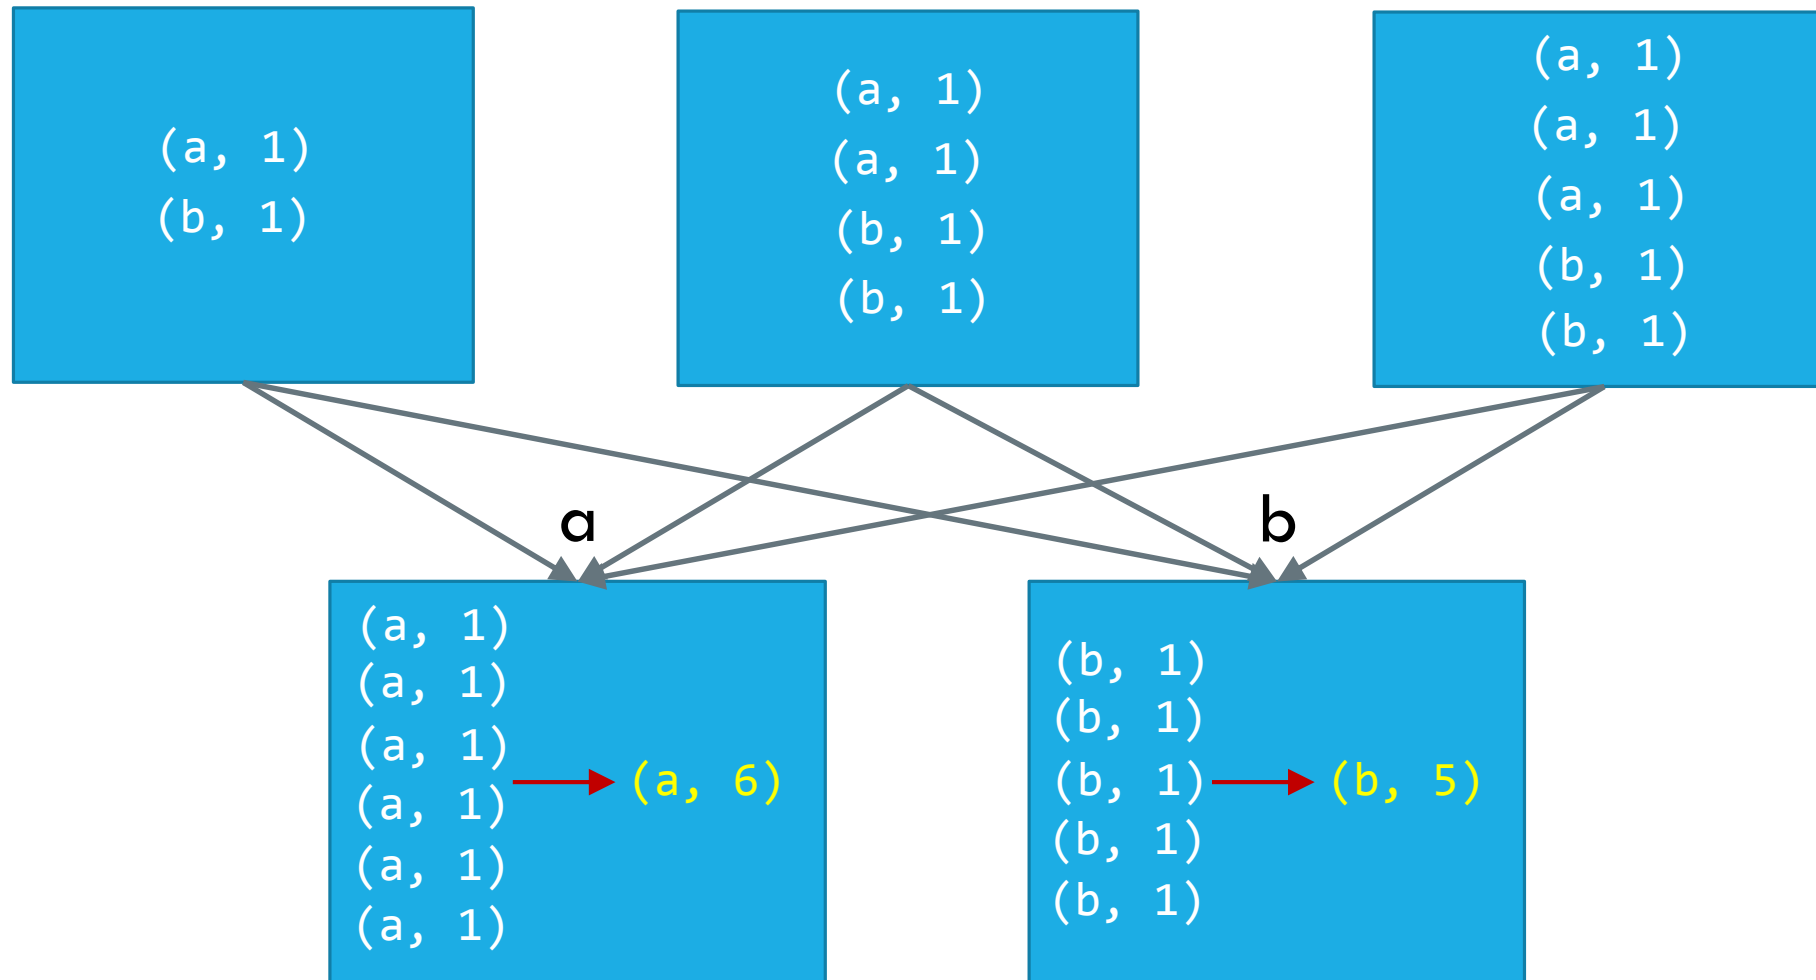
```
val wordCountsWithReduce = wordPairsRDD  
  .reduceByKey(_ + _)  
  .collect()
```

```
val wordCountsWithGroup = wordPairsRDD  
  .groupByKey()  
  .map(t => (t._1, t._2.sum))  
  .collect()
```


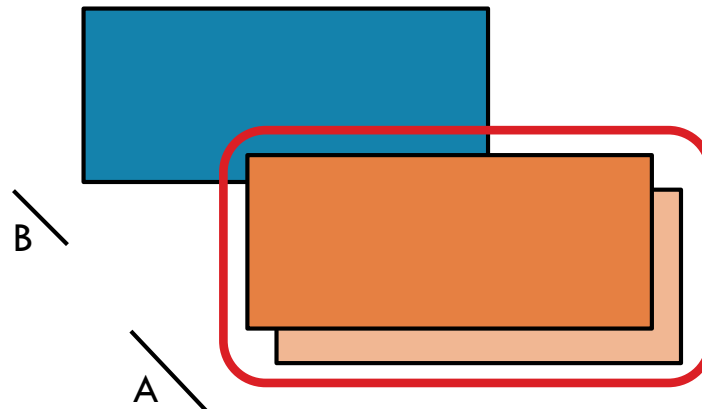
REDUCEBYKEY



GROUPBYKEY



MAPPARTITIONS



`mapPartitions(f, preservesPartitioning=False)`

Return a new RDD by applying a function to each partition of this RDD



```
x = sc.parallelize([1,2,3], 2)
```

```
def f(iterator): yield sum(iterator); yield 42
```

```
y = x.mapPartitions(f)
```


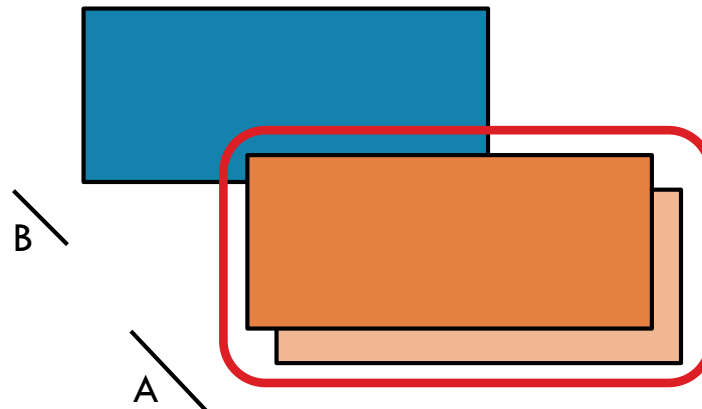
```
# glom() flattens elements on the same partition
print(x.glom().collect())
print(y.glom().collect())
```

x: `[[1], [2, 3]]`

y: `[[1, 42], [5, 42]]`



MAPPARTITIONS



`mapPartitions(f, preservesPartitioning=False)`

Return a new RDD by applying a function to each partition of this RDD

```
val x = sc.parallelize(Array(1,2,3), 2)
```

```
def f(i:Iterator[Int])={ (i.sum,42).productIterator }
```

```
val y = x.mapPartitions(f)
```

```
// glom() flattens elements on the same partition
```

```
val xOut = x.glom().collect()
```

```
val yOut = y.glom().collect()
```



x: Array(Array(1), Array(2, 3))

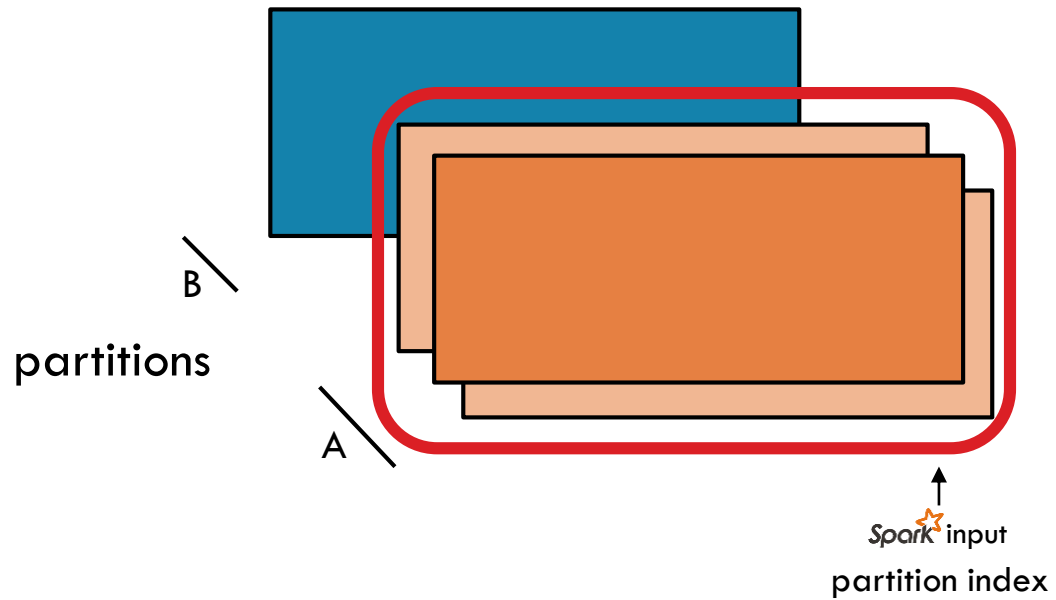
y: Array(Array(1, 42), Array(5, 42))





MAPPARTITIONSWITHINDEX

RDD: **x**

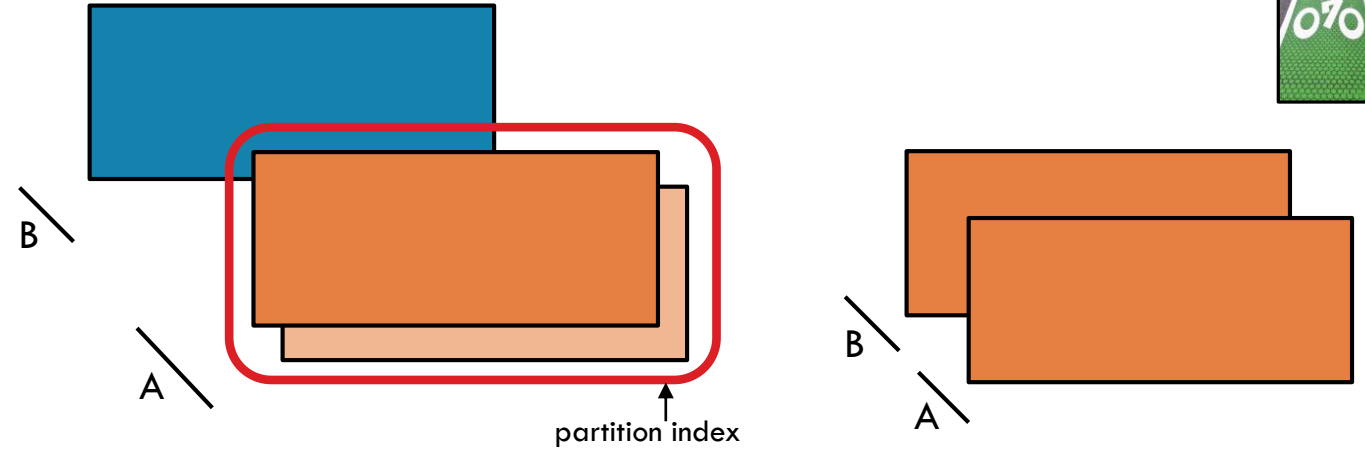


RDD: **y**





MAPPARTITIONSWITHINDEX



`mapPartitionsWithIndex(f, preservesPartitioning=False)`


Return a new RDD by applying a function to each partition of this RDD, while tracking the index of the original partition

```
x = sc.parallelize([1,2,3], 2)
```

```
def f(partitionIndex, iterator): yield (partitionIndex, sum(iterator))
```

```
y = x.mapPartitionsWithIndex(f)
```

```
# glom() flattens elements on the same partition  
print(x.glom().collect())  
print(y.glom().collect())
```

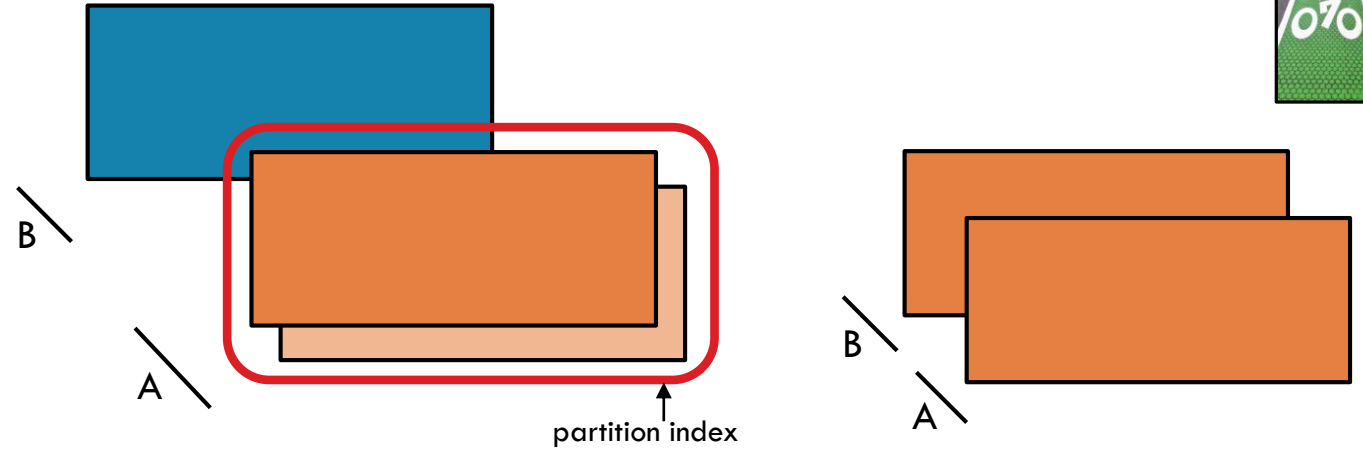


```
x: [[1], [2, 3]]  
y: [[0, 1], [1, 5]]
```





MAPPARTITIONSWITHINDEX



`mapPartitionsWithIndex(f, preservesPartitioning=False)`

Return a new RDD by applying a function to each partition of this RDD, while tracking the index of the original partition.

```
val x = sc.parallelize(Array(1,2,3), 2)
```


```
def f(partitionIndex:Int, i:Iterator[Int]) = {  
  (partitionIndex, i.sum).productIterator  
}
```

```
val y = x.mapPartitionsWithIndex(f)
```

```
// glom() flattens elements on the same partition
```

```
val xOut = x.glom().collect()
```

```
val yOut = y.glom().collect()
```



x: Array(Array(1), Array(2, 3))

y: Array(Array(0, 1), Array(1, 5))

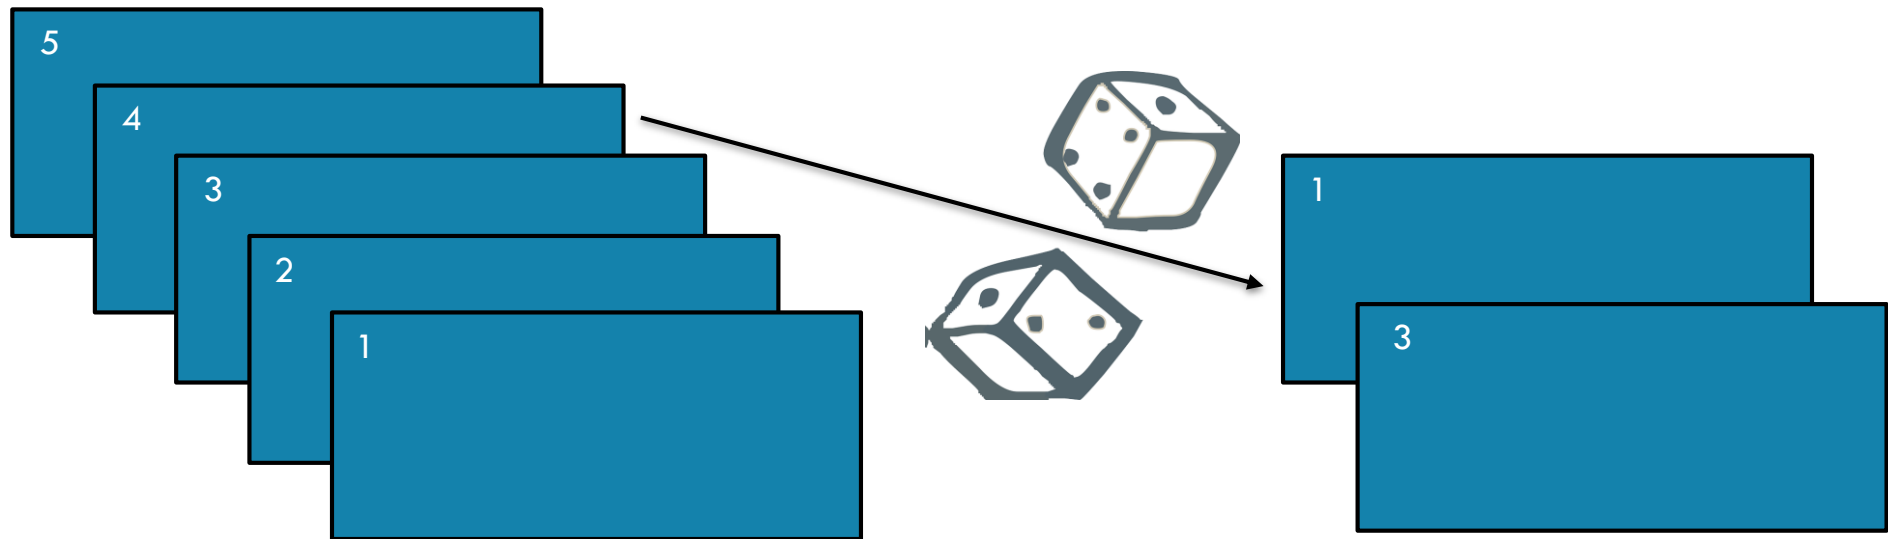




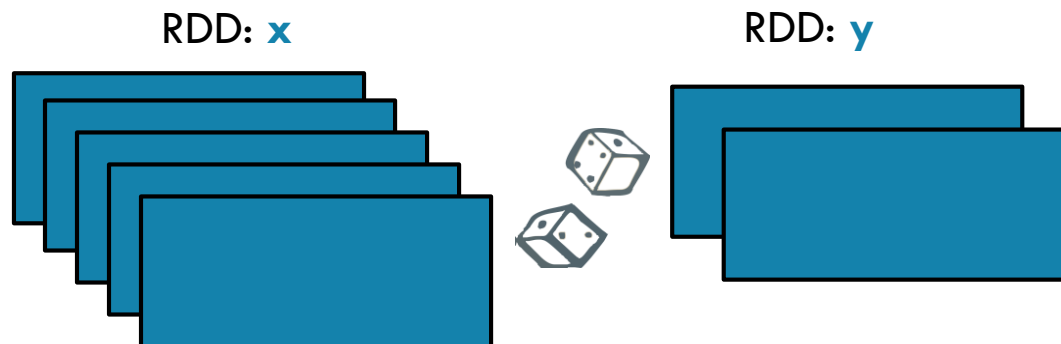
SAMPLE

RDD: **x**

RDD: **y**



SAMPLE



`sample(withReplacement, fraction, seed=None)`

Return a new RDD containing a statistical sample of the original RDD

```
x = sc.parallelize([1, 2, 3, 4, 5])
y = x.sample(False, 0.4, 42)
print(x.collect())
print(y.collect())
```



```
val x = sc.parallelize(Array(1, 2, 3, 4, 5))
val y = x.sample(false, 0.4)
```

```
// omitting seed will yield different output
println(y.collect().mkString(", "))
```

x: [1, 2, 3, 4, 5]

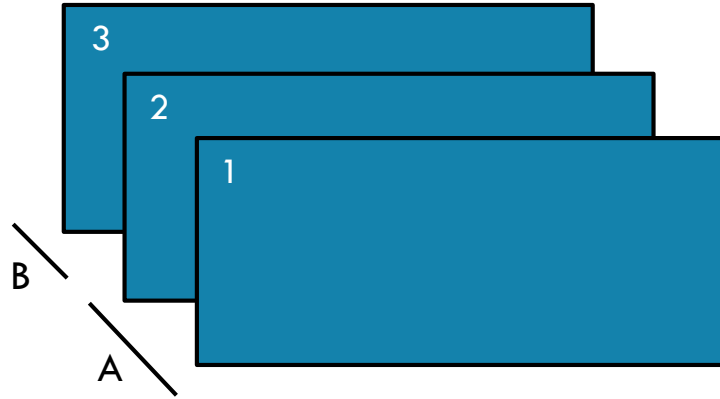
y: [1, 3]



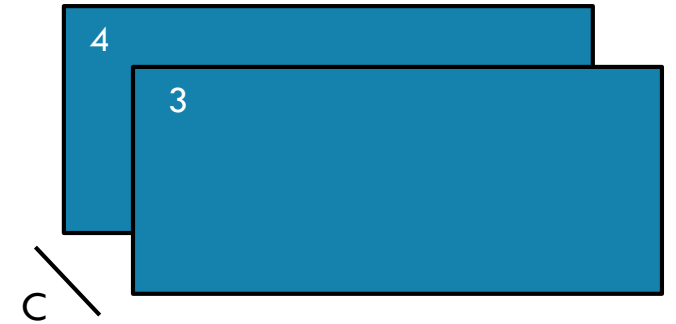


UNION



RDD: **x**



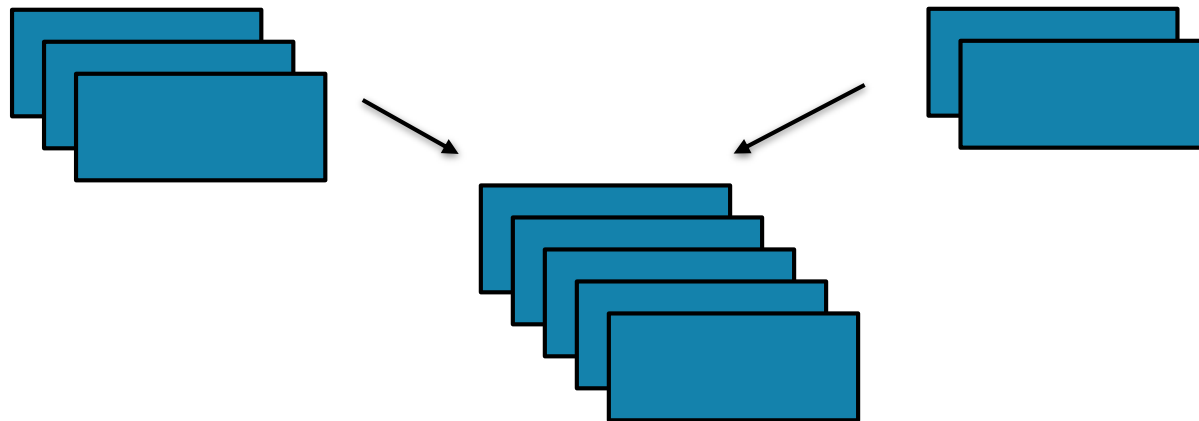
RDD: **y**



RDD: **z**



UNION



Return a new RDD containing all items from two original RDDs. Duplicates are *not* culled.

`union(otherRDD)`



```
x = sc.parallelize([1,2,3], 2)
y = sc.parallelize([3,4], 1)
z = x.union(y)
print(z.glom().collect())
```



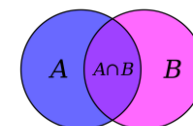
`x:` [1, 2, 3]

`y:` [3, 4]

`z:` [[1], [2, 3], [3, 4]]

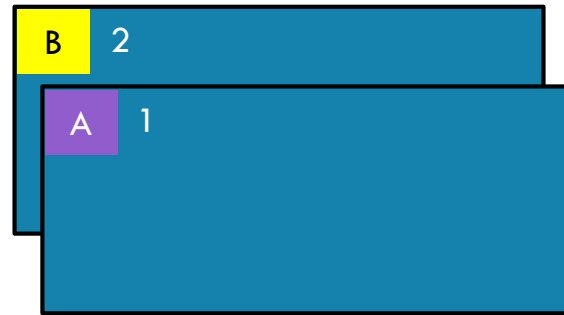


```
val x = sc.parallelize(Array(1,2,3), 2)
val y = sc.parallelize(Array(3,4), 1)
val z = x.union(y)
val zOut = z.glom().collect()
```

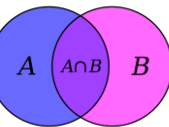


JOIN

RDD: x



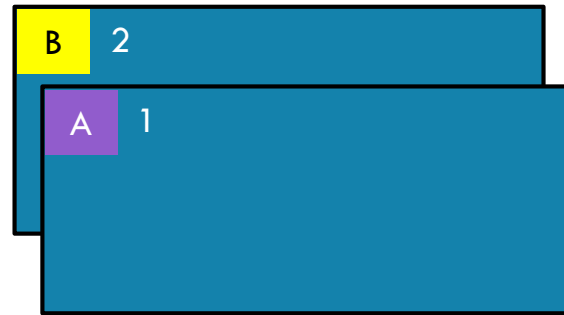
RDD: y



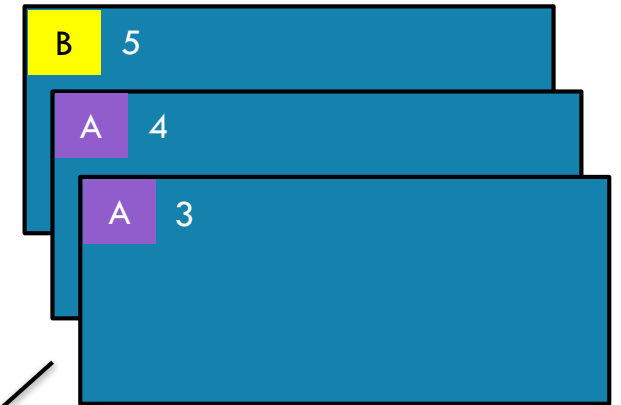


JOIN

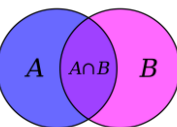
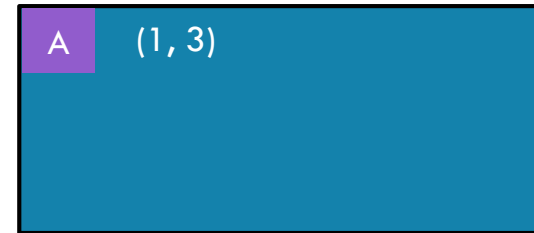
RDD: **x**



RDD: **y**



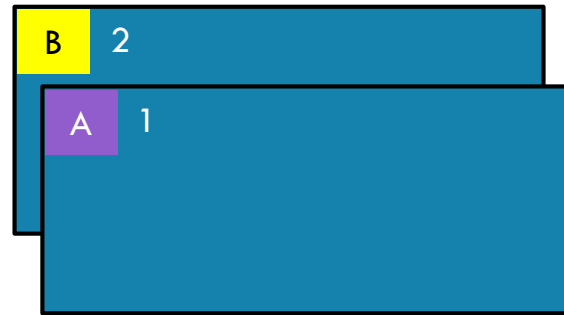
RDD: **z**





JOIN


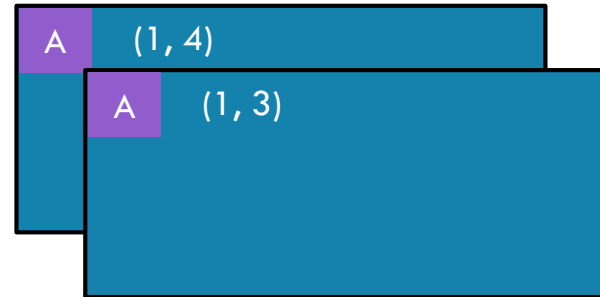
RDD: **x**



RDD: **y**



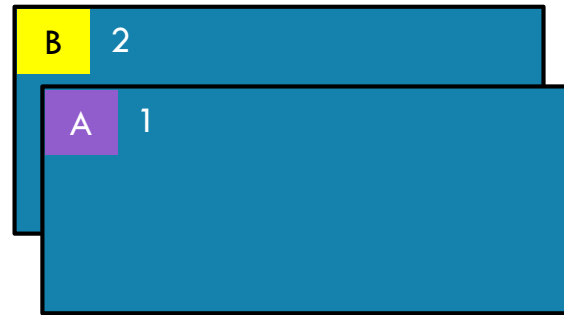
RDD: **z**



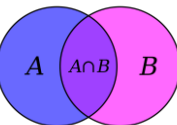
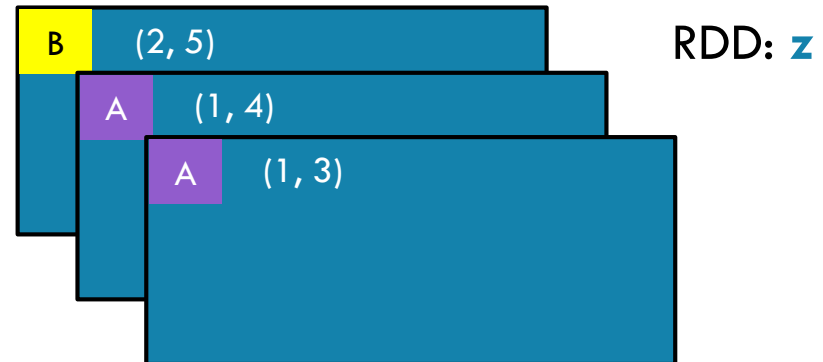



JOIN


RDD: x



RDD: y



JOIN



Return a new RDD containing all pairs of elements having the same key in the original RDDs

`union(otherRDD, numPartitions=None)`



```
x = sc.parallelize([("a", 1), ("b", 2)])
y = sc.parallelize([("a", 3), ("a", 4), ("b", 5)])
z = x.join(y)
print(z.collect())
```




x: [("a", 1), ("b", 2)]

y: [("a", 3), ("a", 4), ("b", 5)]

z: [('a', (1, 3)), ('a', (1, 4)), ('b', (2, 5))]



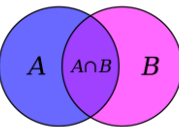
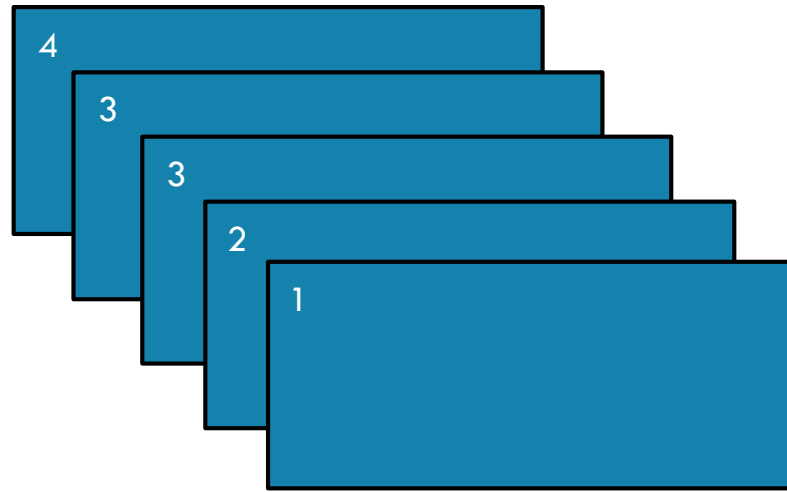
```
val x = sc.parallelize(Array(("a", 1), ("b", 2)))
val y = sc.parallelize(Array(("a", 3), ("a", 4), ("b", 5)))
val z = x.join(y)
println(z.collect().mkString(", "))
```





DISTINCT

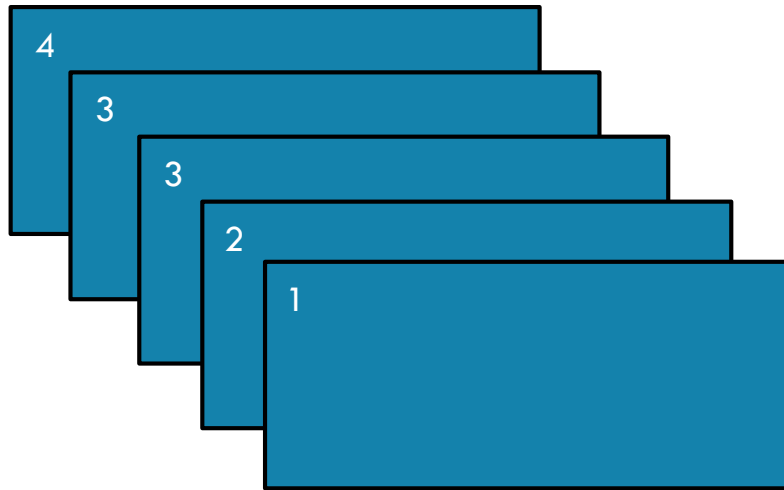
RDD: x







DISTINCT

RDD: x




RDD: y







DISTINCT

RDD: x



RDD: y



DISTINCT



Return a new RDD containing distinct items from the original RDD (omitting all duplicates)
`distinct(numPartitions=None)`



```
x = sc.parallelize([1,2,3,3,4])  
y = x.distinct()  
  
print(y.collect())
```

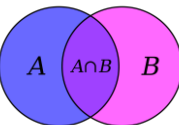


```
val x = sc.parallelize(Array(1,2,3,3,4))  
val y = x.distinct()  
  
println(y.collect().mkString(", "))
```



x: [1, 2, 3, 3, 4]

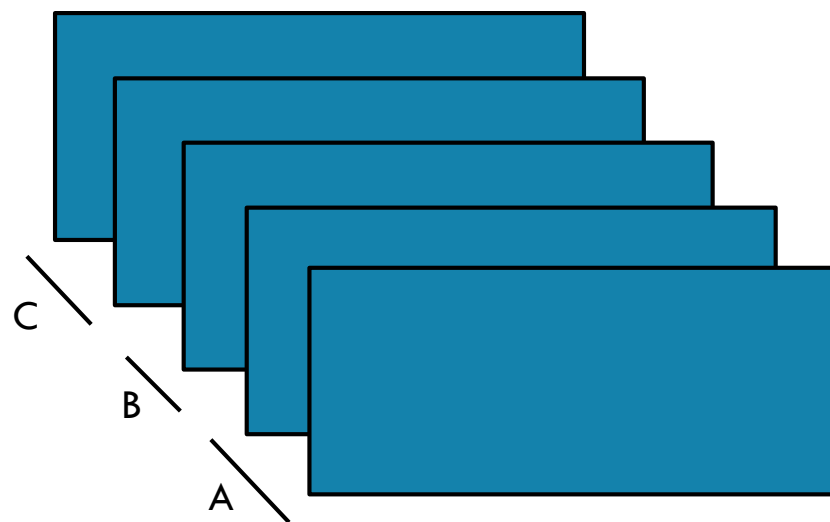
y: [1, 2, 3, 4]





COALESCE

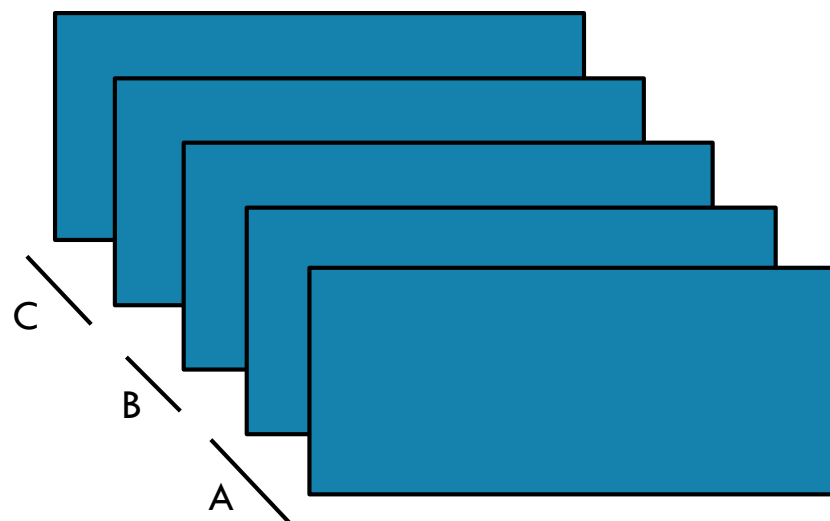
RDD: x



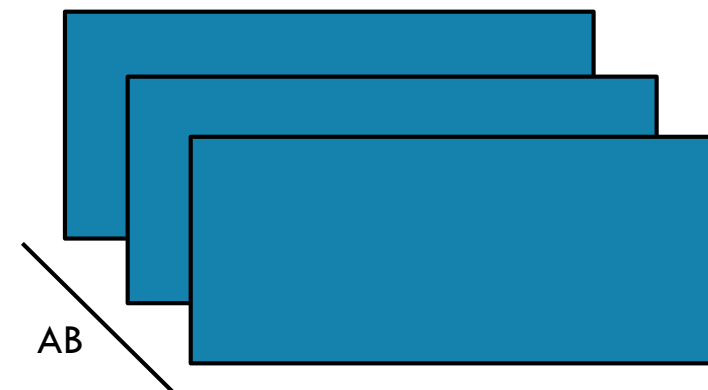


COALESCE

RDD: **x**




RDD: **y**






COALESCE


RDD: **x**



RDD: **y**



COALESCE



Return a new RDD which is reduced to a smaller number of partitions

`coalesce(numPartitions, shuffle=False)`



```
x = sc.parallelize([1, 2, 3, 4, 5], 3)
y = x.coalesce(2)
print(x.glom().collect())
print(y.glom().collect())
```



x: [[1], [2, 3], [4, 5]]

y: [[1], [2, 3, 4, 5]]




```
val x = sc.parallelize(Array(1, 2, 3, 4, 5), 3)
val y = x.coalesce(2)
val xOut = x.glom().collect()
val yOut = y.glom().collect()
```






KEYBY

RDD: **x**




RDD: **y**



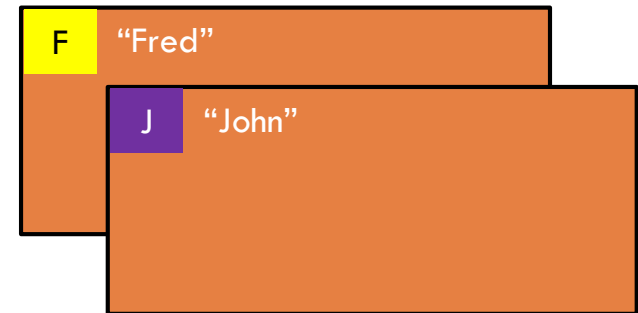


KEYBY

RDD: **x**




RDD: **y**






KEYBY

RDD: **x**




RDD: **y**



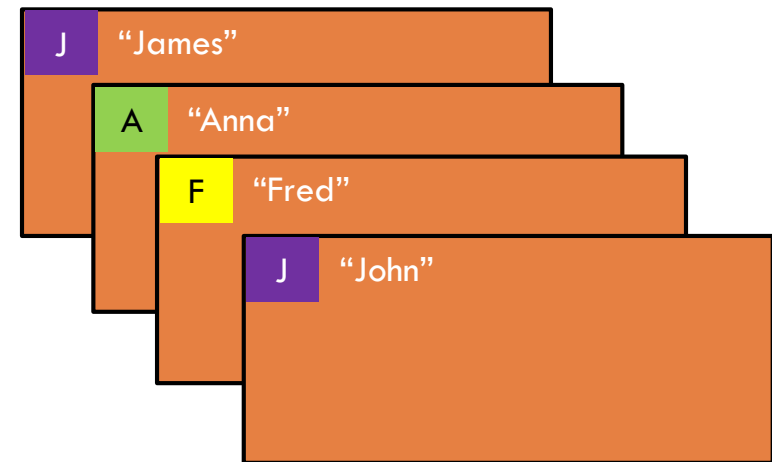


KEYBY

RDD: **x**



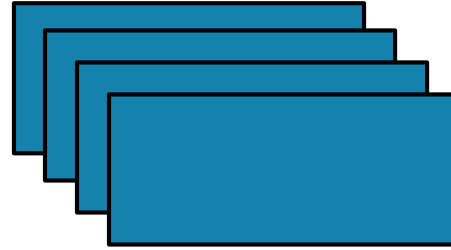
RDD: **y**



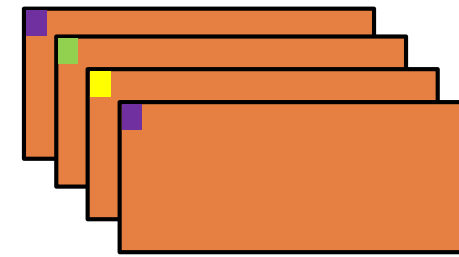


KEYBY

RDD: **x**



RDD: **y**



keyBy(**f**)

Create a Pair RDD, forming one pair for each item in the original RDD. The pair's key is calculated from the value via a user-supplied function.



```
x = sc.parallelize(['John', 'Fred', 'Anna', 'James'])
y = x.keyBy(lambda w: w[0])
print y.collect()
```



x: ['John', 'Fred', 'Anna', 'James']

y: [('J', 'John'), ('F', 'Fred'), ('A', 'Anna'), ('J', 'James')]




```
val x = sc.parallelize(
  Array("John", "Fred", "Anna", "James"))
val y = x.keyBy(w => w.charAt(0))
println(y.collect().mkString(", "))
```





PARTITIONBY

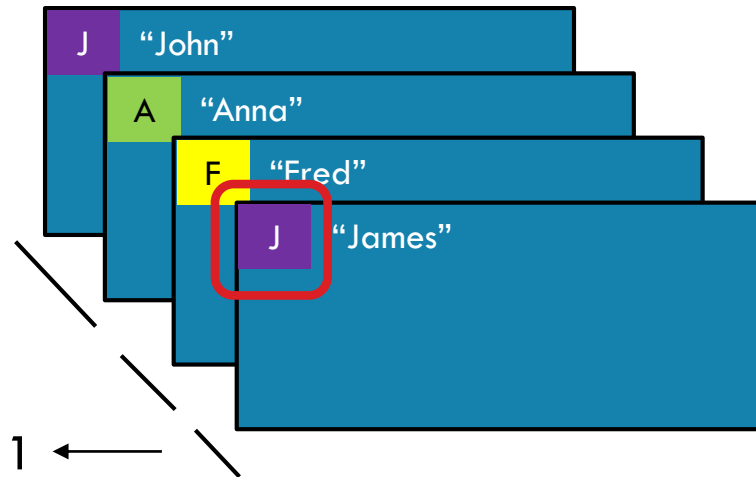
RDD: x



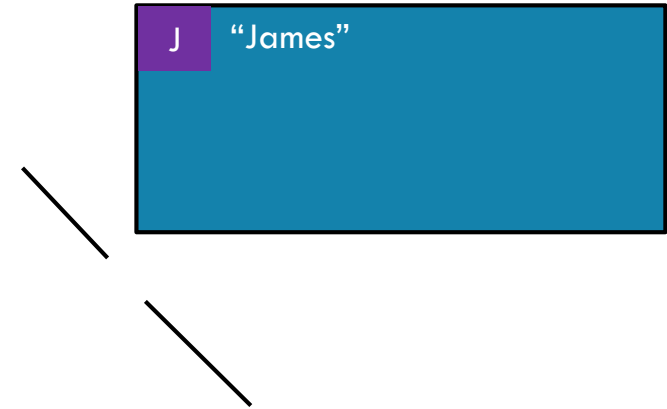


PARTITIONBY

RDD: **x**



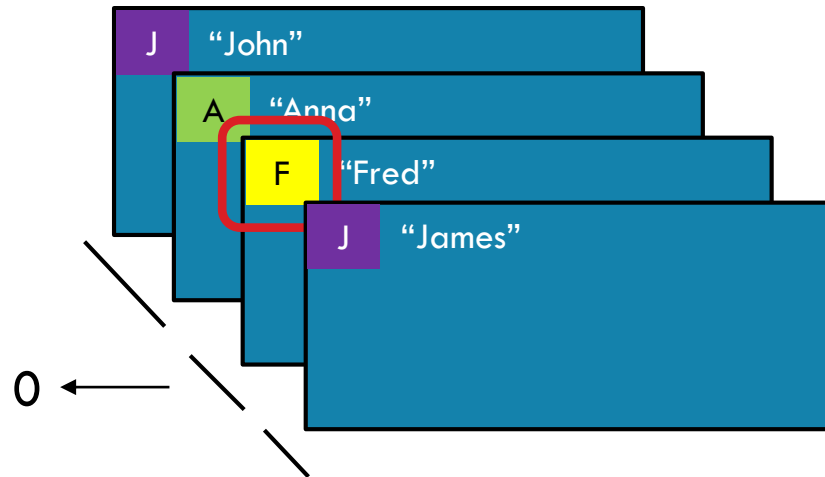
RDD: **y**






PARTITIONBY

RDD: **x**



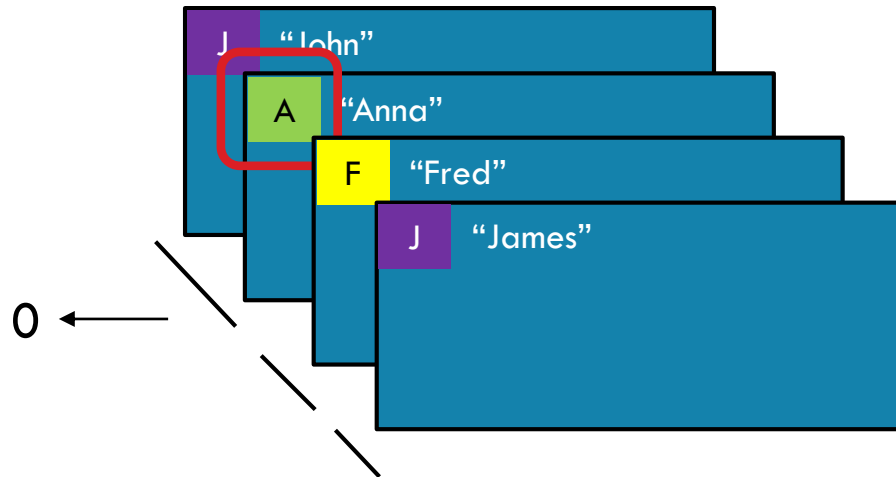
RDD: **y**



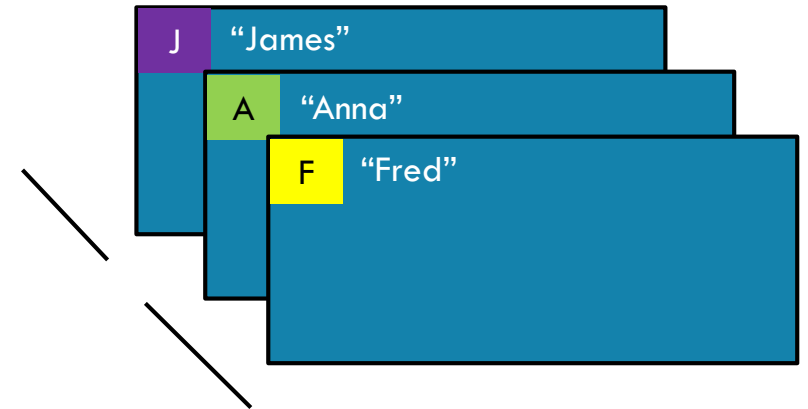


PARTITIONBY

RDD: **x**



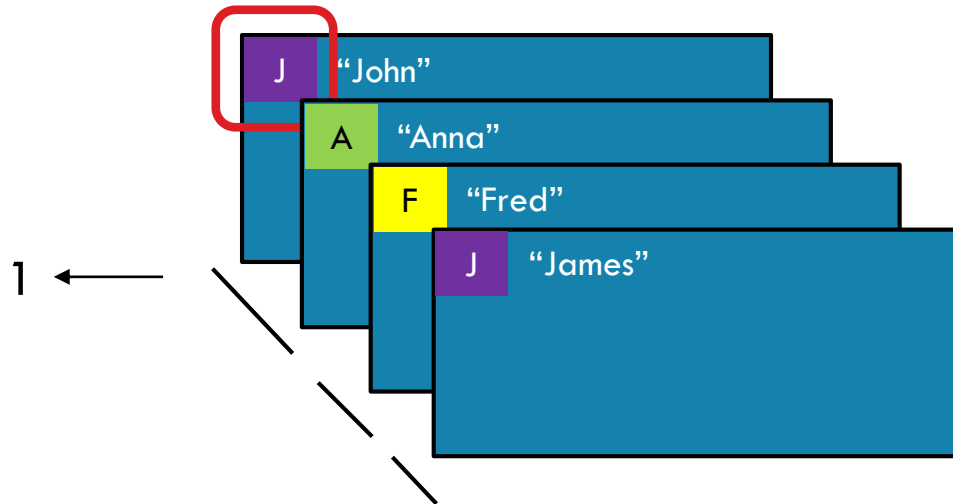
RDD: **y**



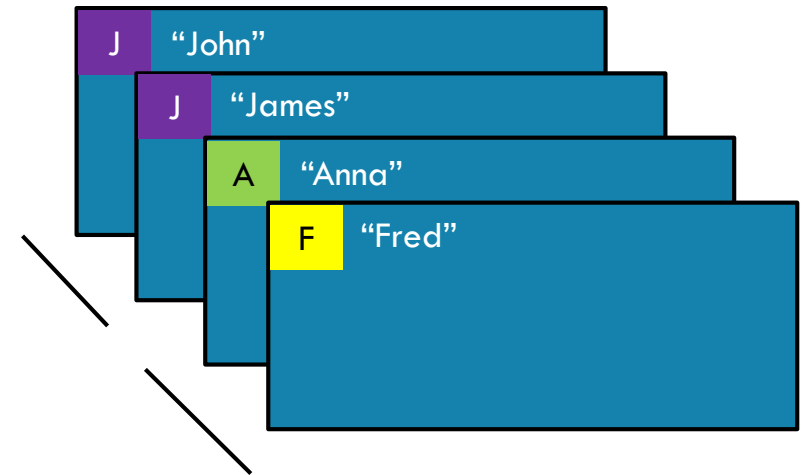


PARTITIONBY


RDD: **x**



RDD: **y**



PARTITIONBY



Return a new RDD with the specified number of partitions, placing original items into the partition returned by a user supplied function

`partitionBy(numPartitions, partitioner=portable_hash)`



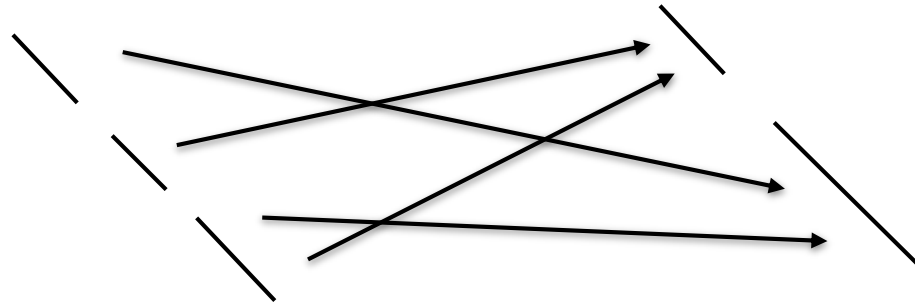
```
x = sc.parallelize([('J','James'),('F','Fred'),  
                    ('A','Anna'),('J','John')], 3)  
  
y = x.partitionBy(2, lambda w: 0 if w[0] < 'H' else 1)  
print x.glom().collect()  
print y.glom().collect()
```



```
x: [[('J', 'James')], [('F', 'Fred')],  
    [('A', 'Anna'), ('J', 'John')]]  
  
y: [[('A', 'Anna'), ('F', 'Fred')],  
    [('J', 'James'), ('J', 'John')]]
```



PARTITIONBY



Return a new RDD with the specified number of partitions, placing original items into the partition returned by a user supplied function.

`partitionBy(numPartitions, partitioner=portable_hash)`

```
import org.apache.spark.Partitioner
val x = sc.parallelize(Array(('J',"James"),('F',"Fred"),
                           ('A',"Anna"),('J',"John")), 3)
```

```
val y = x.partitionBy(new Partitioner() {
  val numPartitions = 2
  def getPartition(k:Any) = {
    if (k.asInstanceOf[Char] < 'H') 0 else 1
  }
})
```

```
val y0ut = y.glom().collect()
```



x: Array(Array((A,Anna), (F,Fred)),
 Array((J,John), (J,James)))

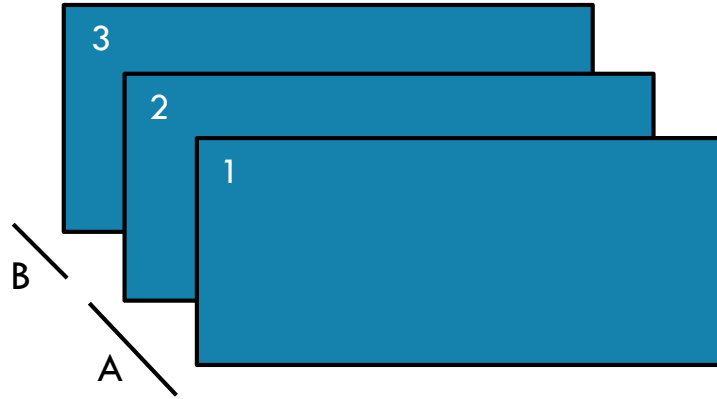
y: Array(Array((F,Fred), (A,Anna)),
 Array((J,John), (J,James)))




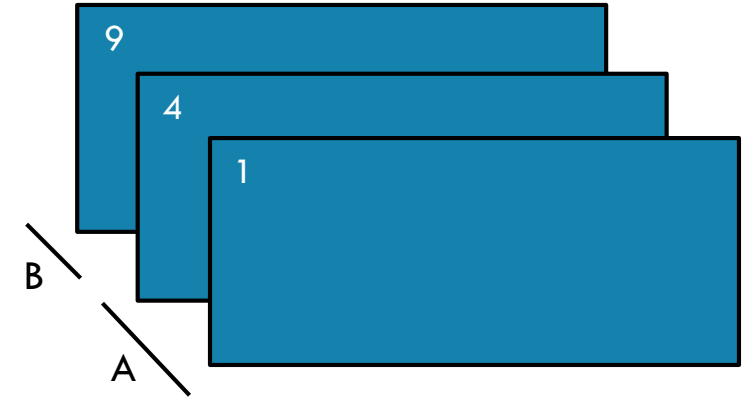


ZIP

RDD: x


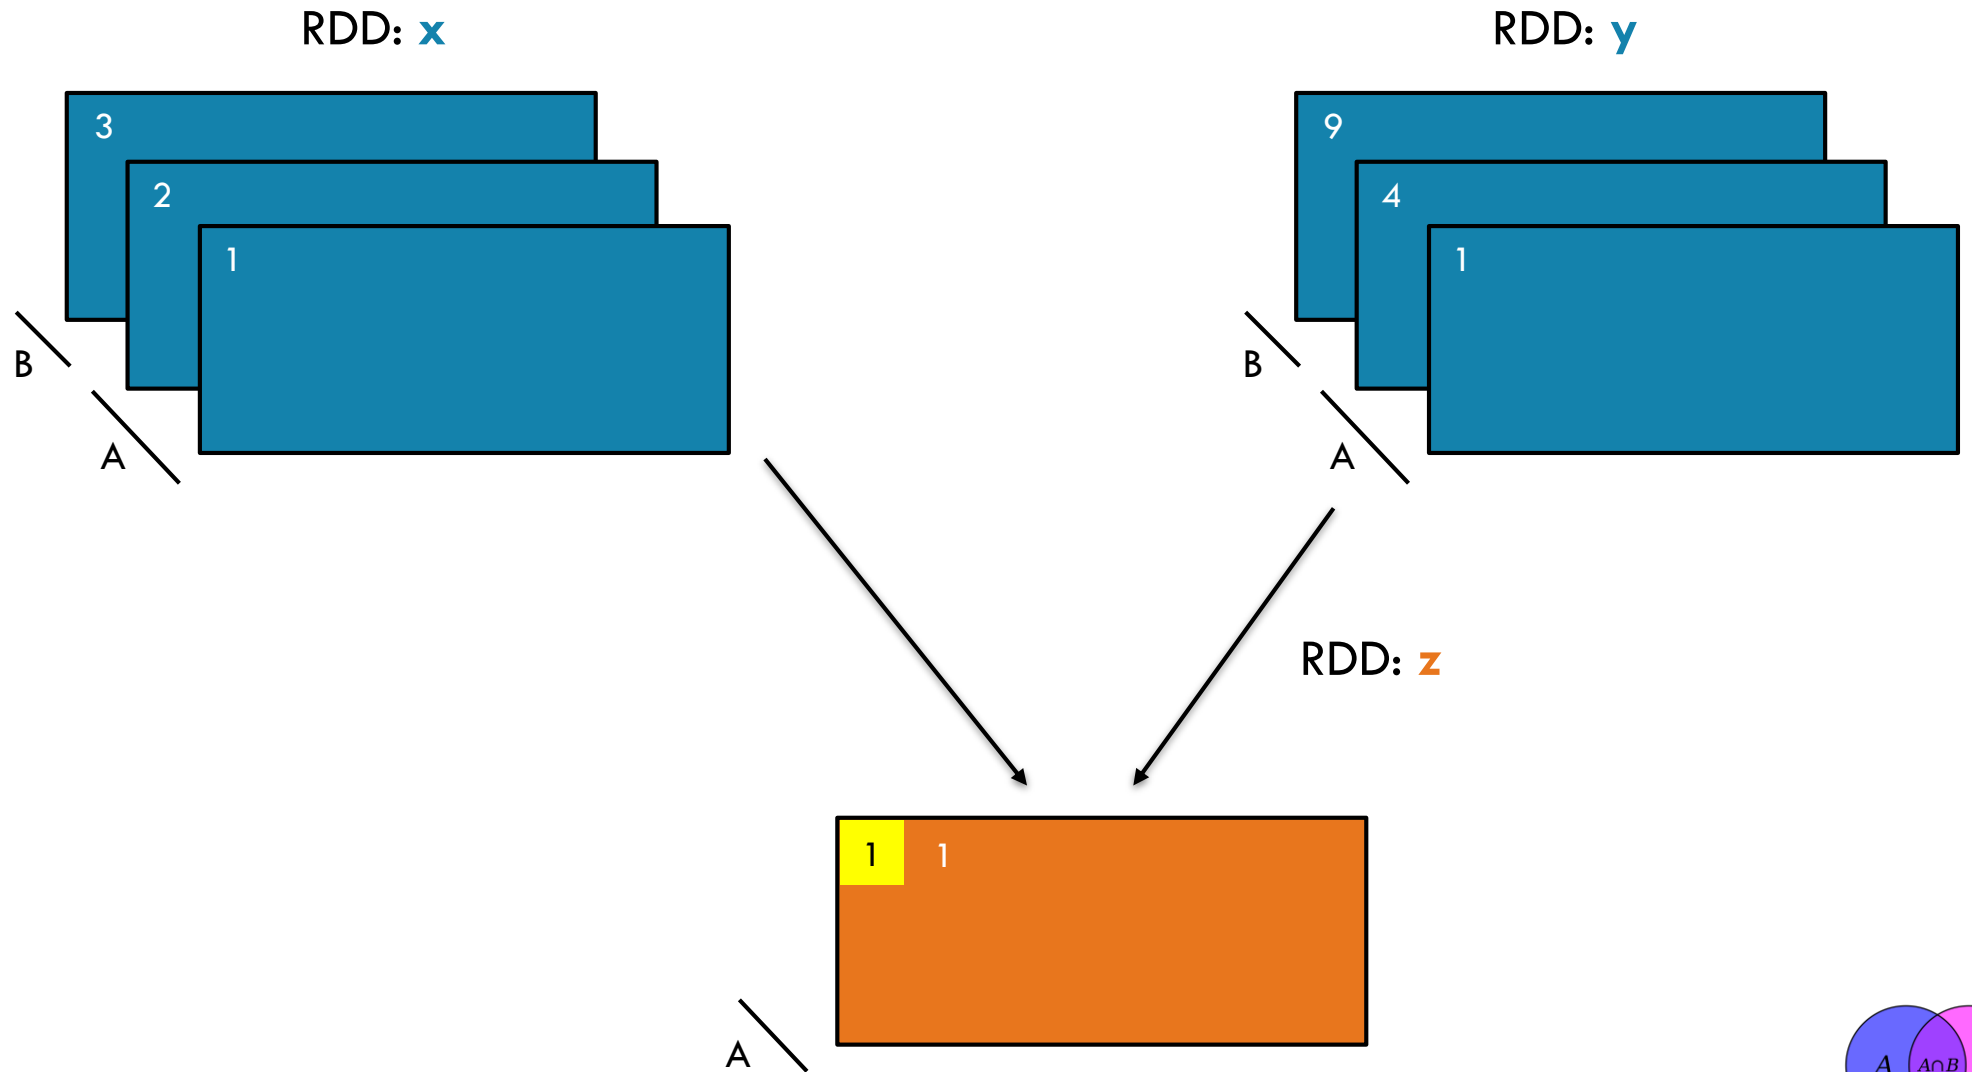


RDD: y



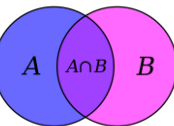
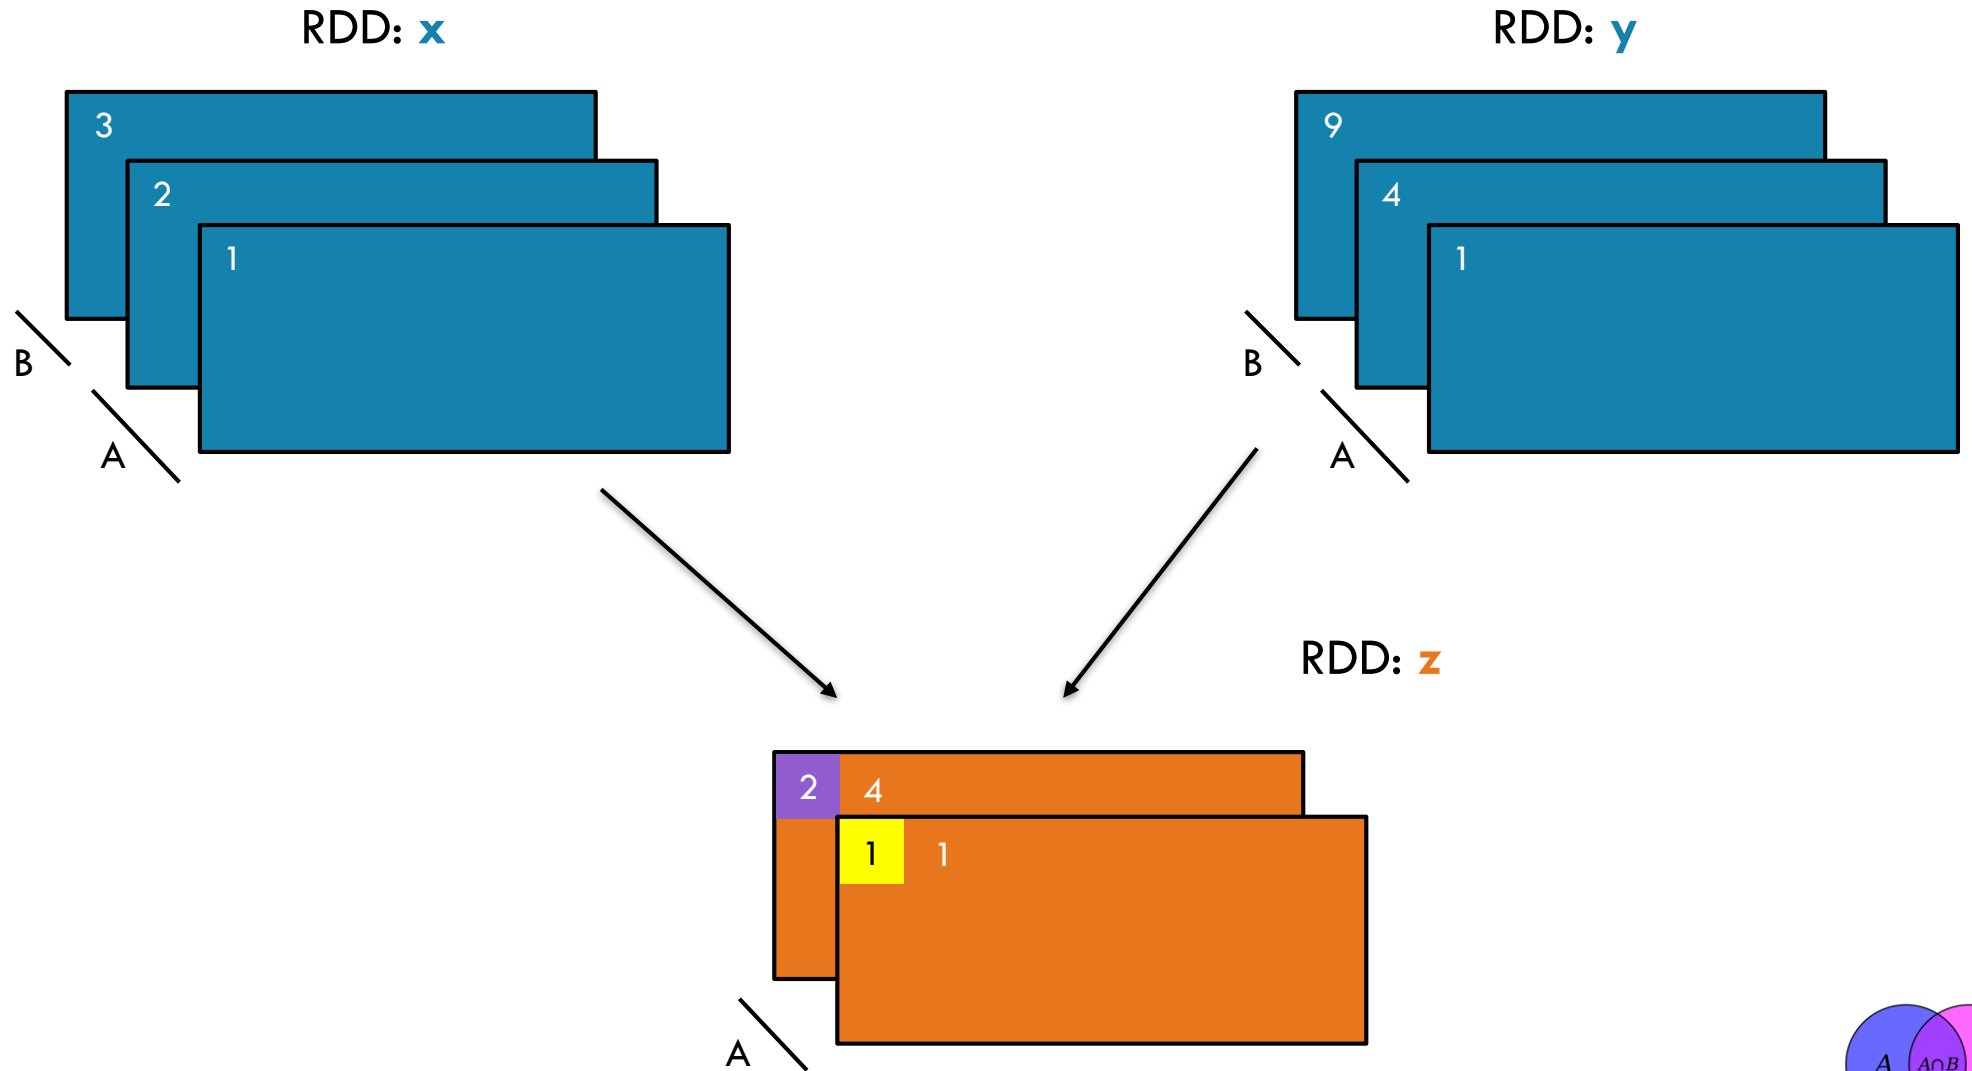


ZIP




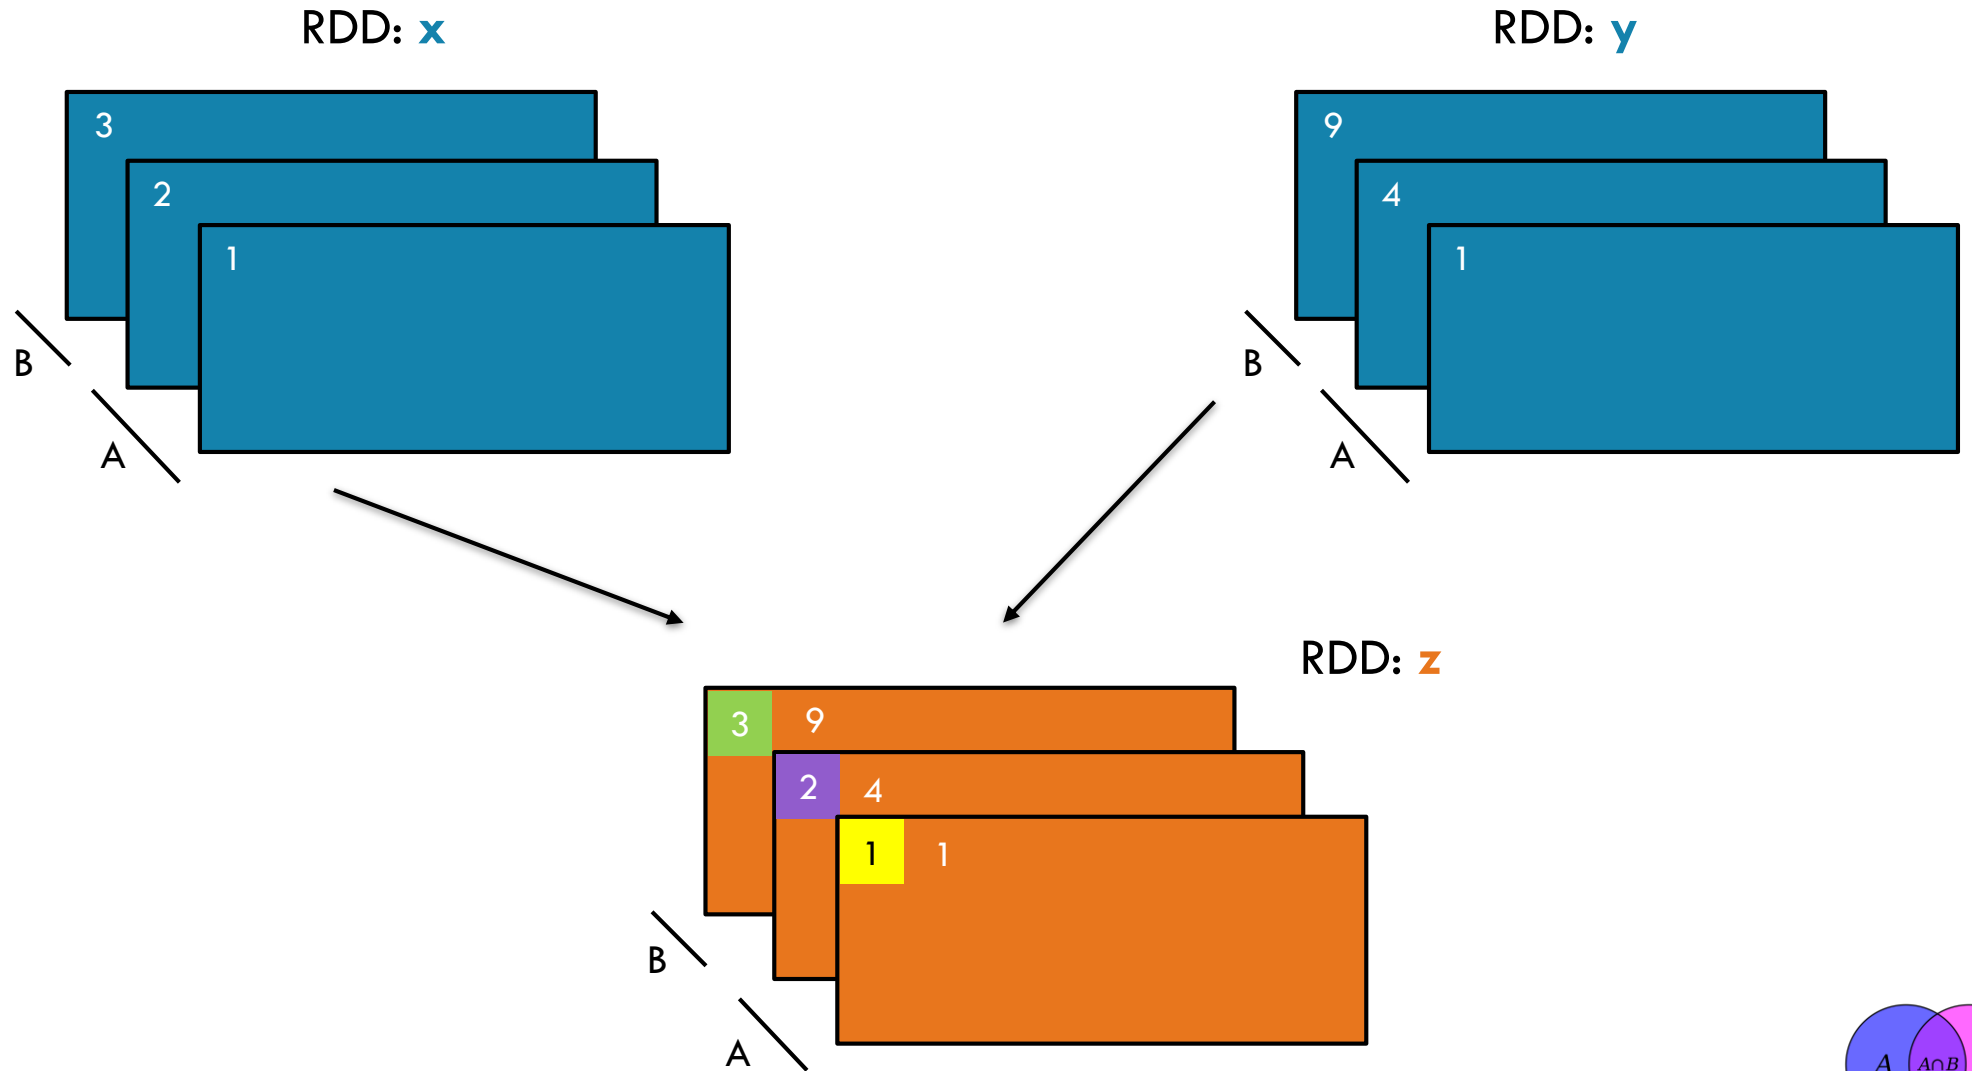


ZIP

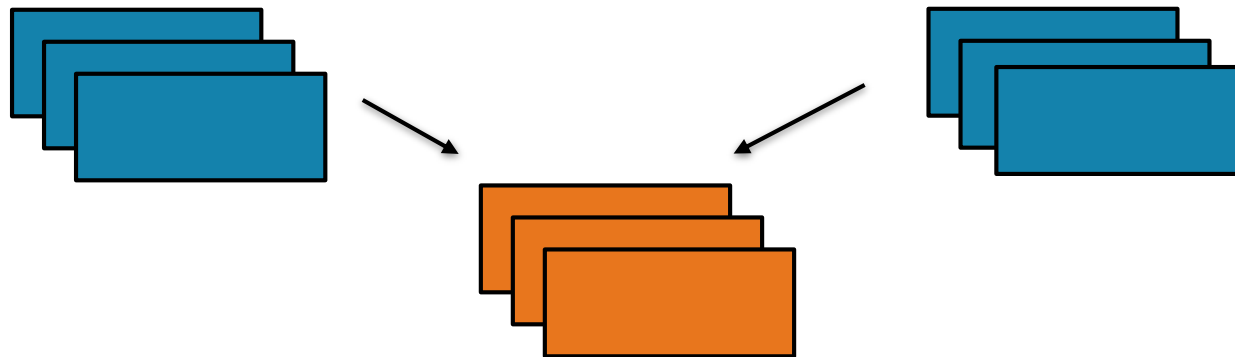




ZIP



ZIP



Return a new RDD containing pairs whose key is the item in the original RDD, and whose value is that item's corresponding element (same partition, same index) in a second RDD

`zip(otherRDD)`



```
x = sc.parallelize([1, 2, 3])
y = x.map(lambda n:n*n)
z = x.zip(y)
```

```
print(z.collect())
```



```
val x = sc.parallelize(Array(1,2,3))
val y = x.map(n=>n*n)
val z = x.zip(y)
```

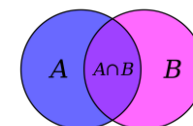
```
println(z.collect().mkString(", "))
```




`x:` [1, 2, 3]


`y:` [1, 4, 9]

`z:` [(1, 1), (2, 4), (3, 9)]





ACTIONS



Core Operations



distributed

occurs across the cluster

VS

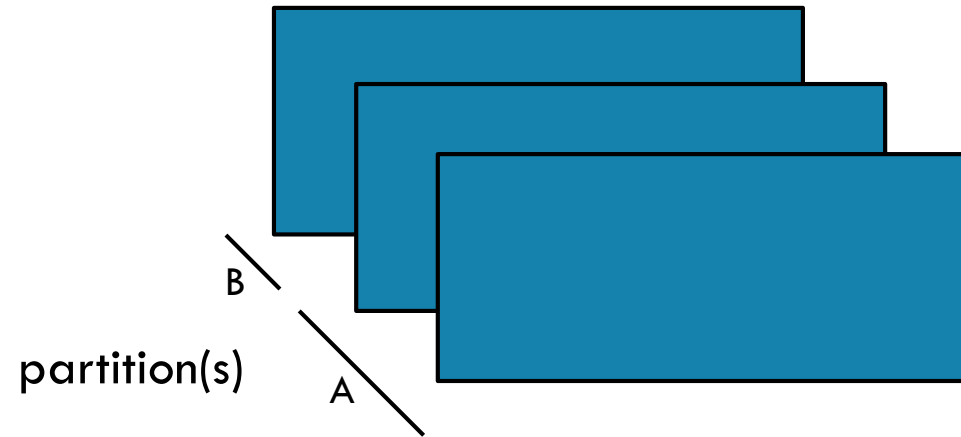
A screenshot of a terminal window. The title bar says "m2-user@p-10-0-12-60". The prompt is "m2-user@p-10-0-12-60 ~\$". The user has entered "dse spark", which has resulted in a "Welcome to" message. Below this is the Spark logo and "version 1.1.0". The terminal shows the Spark version (2.0.0) and the Java version (1.7.0_71). It then shows the creation of a SparkContext and the execution of a Scala REPL session. The REPL session shows the creation of a key-value RDD from a Cassandra table and the execution of a count() operation, resulting in "res0: Long = 4".

driver

result must fit in driver JVM



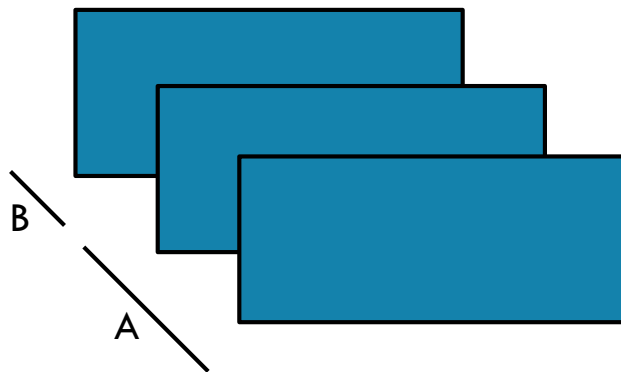
GETNUMPARTITIONS



2



GETNUMPARTITIONS



`getNumPartitions()`

Return the number of partitions in RDD



```
x = sc.parallelize([1,2,3], 2)
y = x.getNumPartitions()

print(x.glom().collect())
print(y)
```



```
val x = sc.parallelize(Array(1,2,3), 2)
val y = x.partitions.size
val xOut = x.glom().collect()
println(y)
```

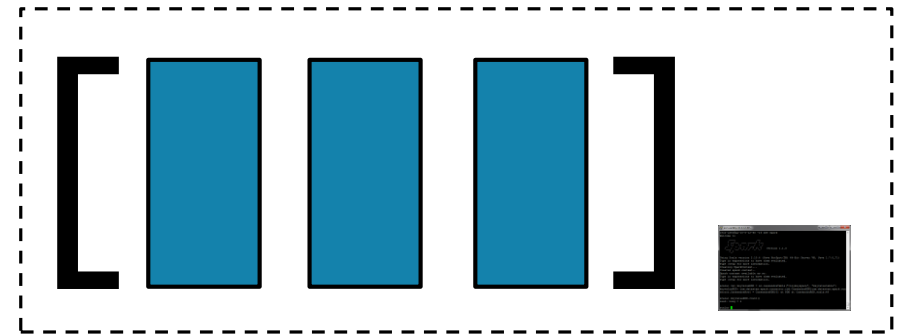
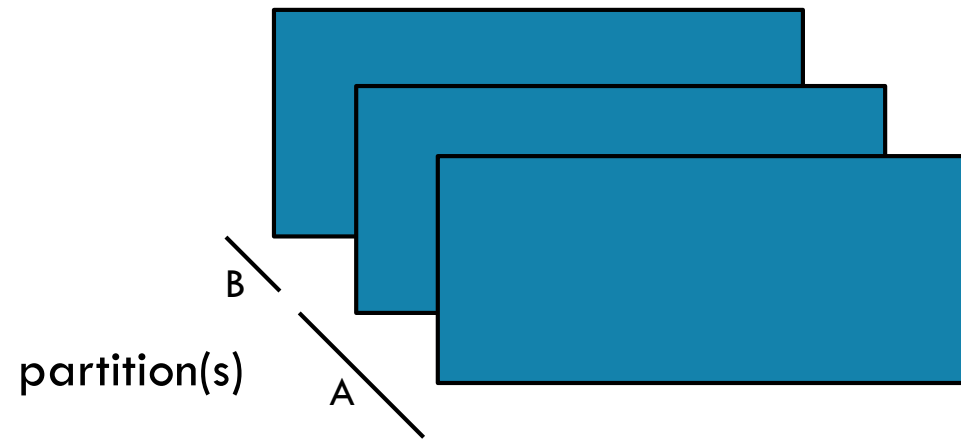


`x:` `[[1], [2, 3]]`

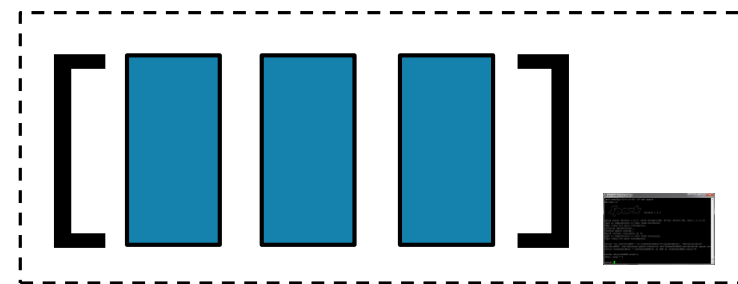
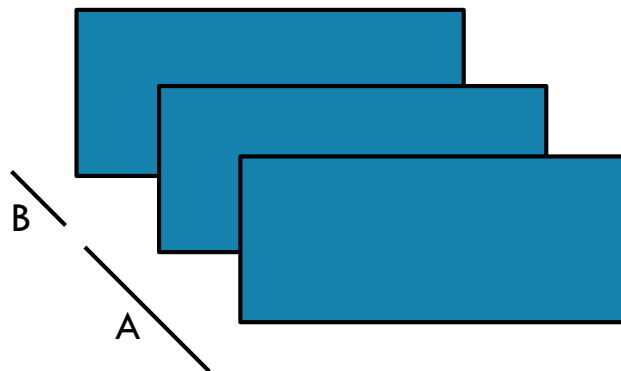
`y:` `2`



COLLECT



COLLECT



`collect()`

Return all items in the RDD to the driver in a single list



```
x = sc.parallelize([1,2,3], 2)
y = x.collect()

print(x.glom().collect())
print(y)
```



```
val x = sc.parallelize(Array(1,2,3), 2)
val y = x.collect()

val xOut = x.glom().collect()
println(y)
```

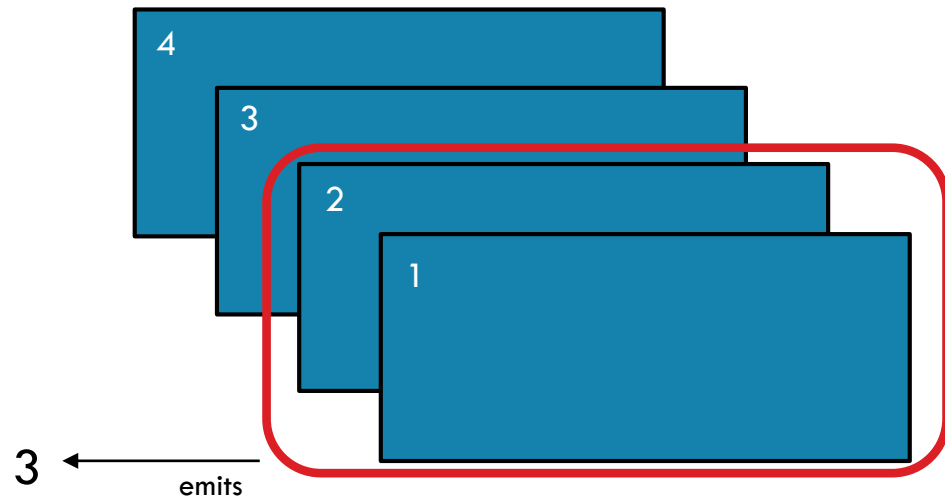


`x:` `[[1], [2, 3]]`


`y:` `[1, 2, 3]`



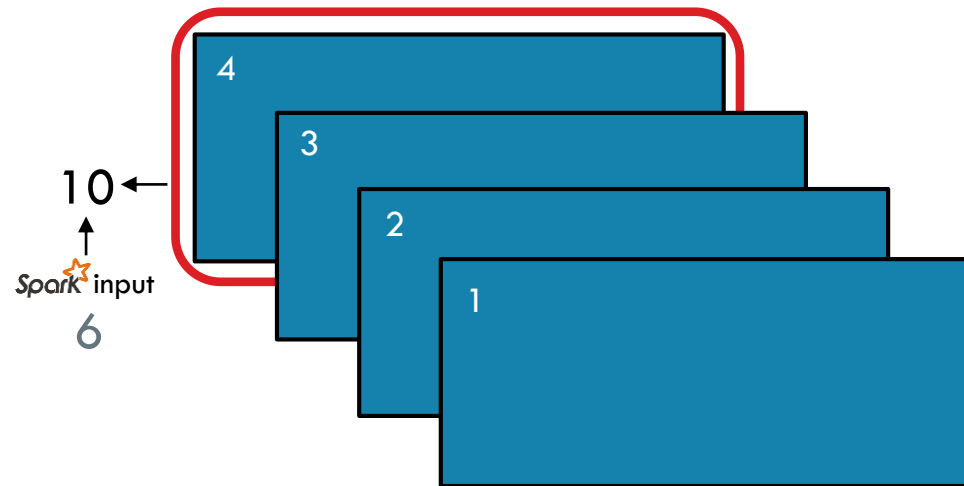
REDUCE




REDUCE



REDUCE



REDUCE



`reduce(f)`

Aggregate all the elements of the RDD by applying a user function pairwise to elements and partial results, and returns a result to the driver



```
x = sc.parallelize([1,2,3,4])
y = x.reduce(lambda a,b: a+b)

print(x.collect())
print(y)
```



```
val x = sc.parallelize(Array(1,2,3,4))
val y = x.reduce((a,b) => a+b)

println(x.collect.mkString(", "))
println(y)
```

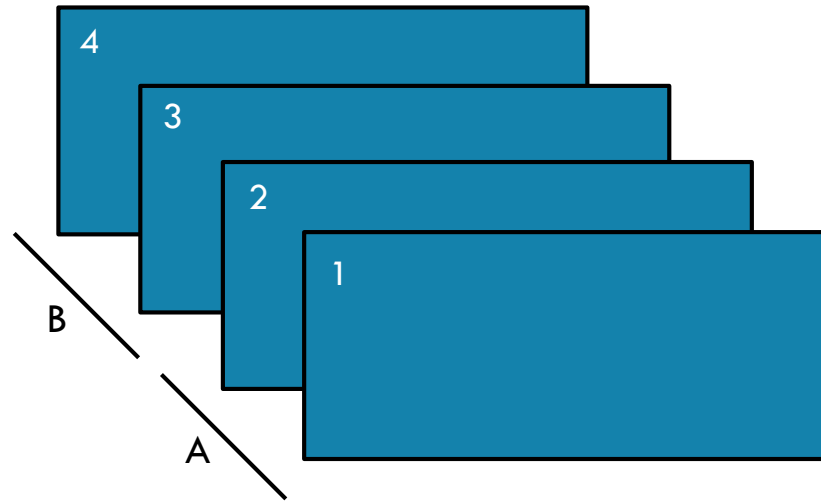


x: [1, 2, 3, 4]

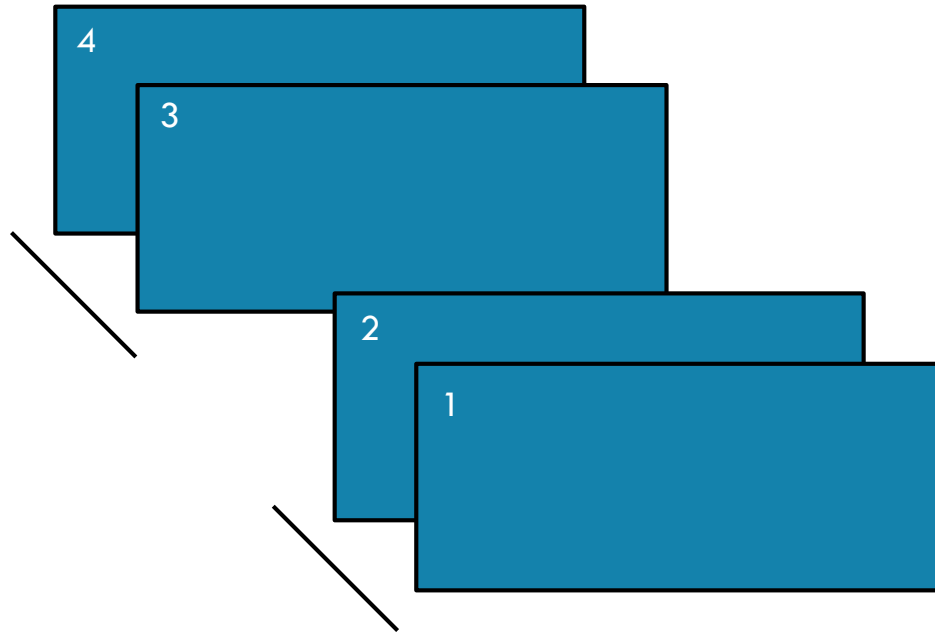
y: 10



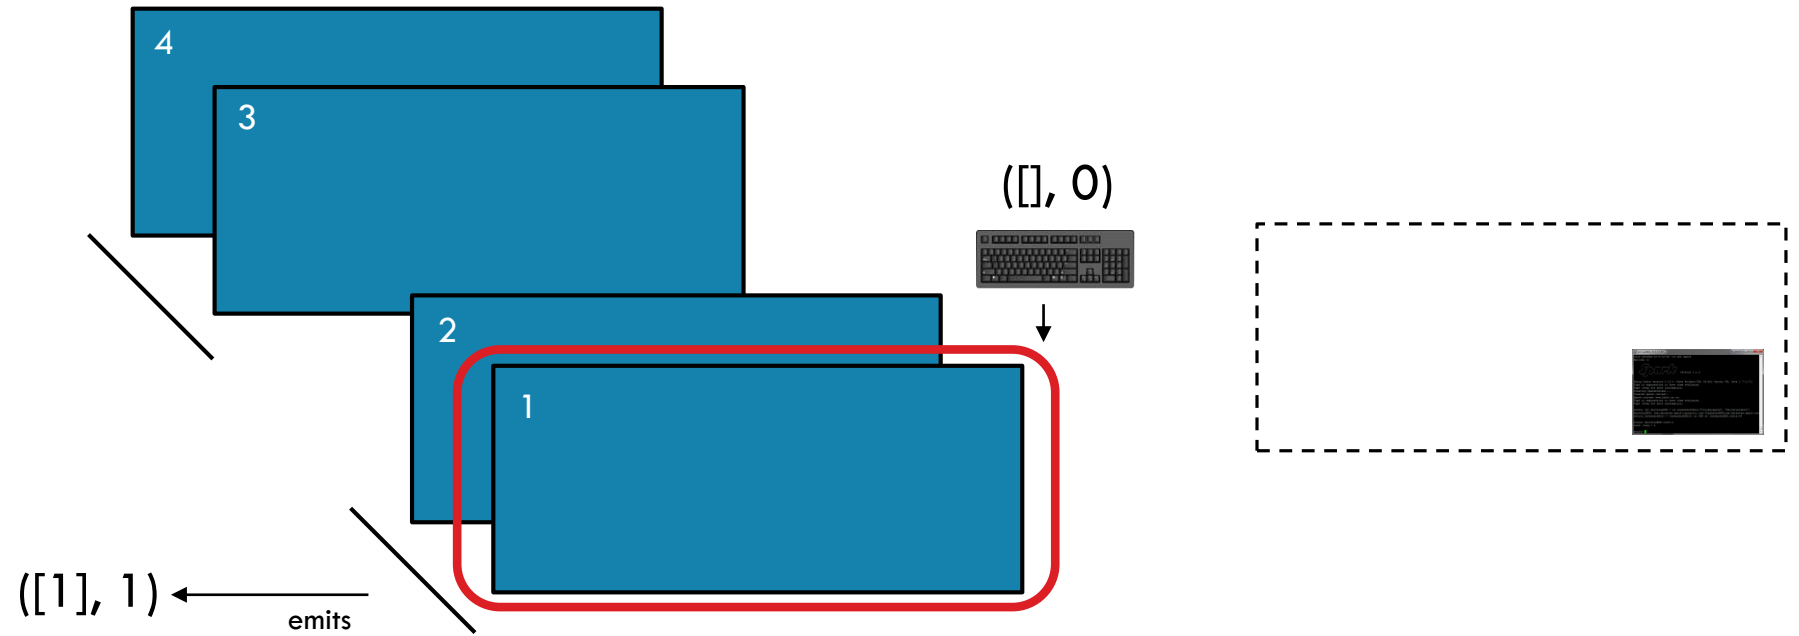
AGGREGATE



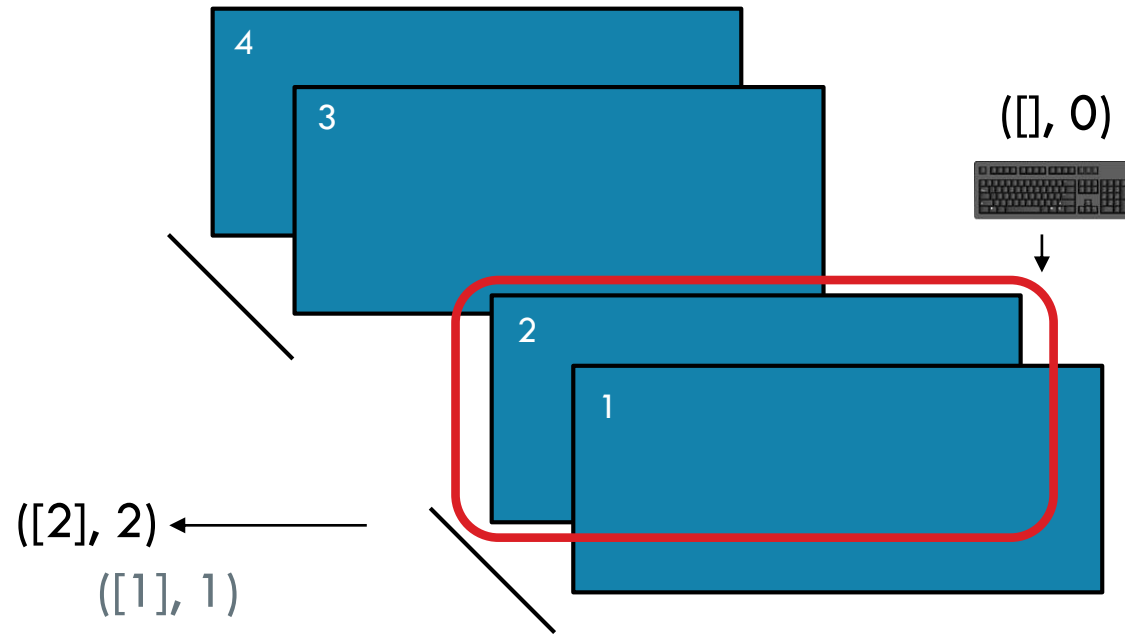
AGGREGATE



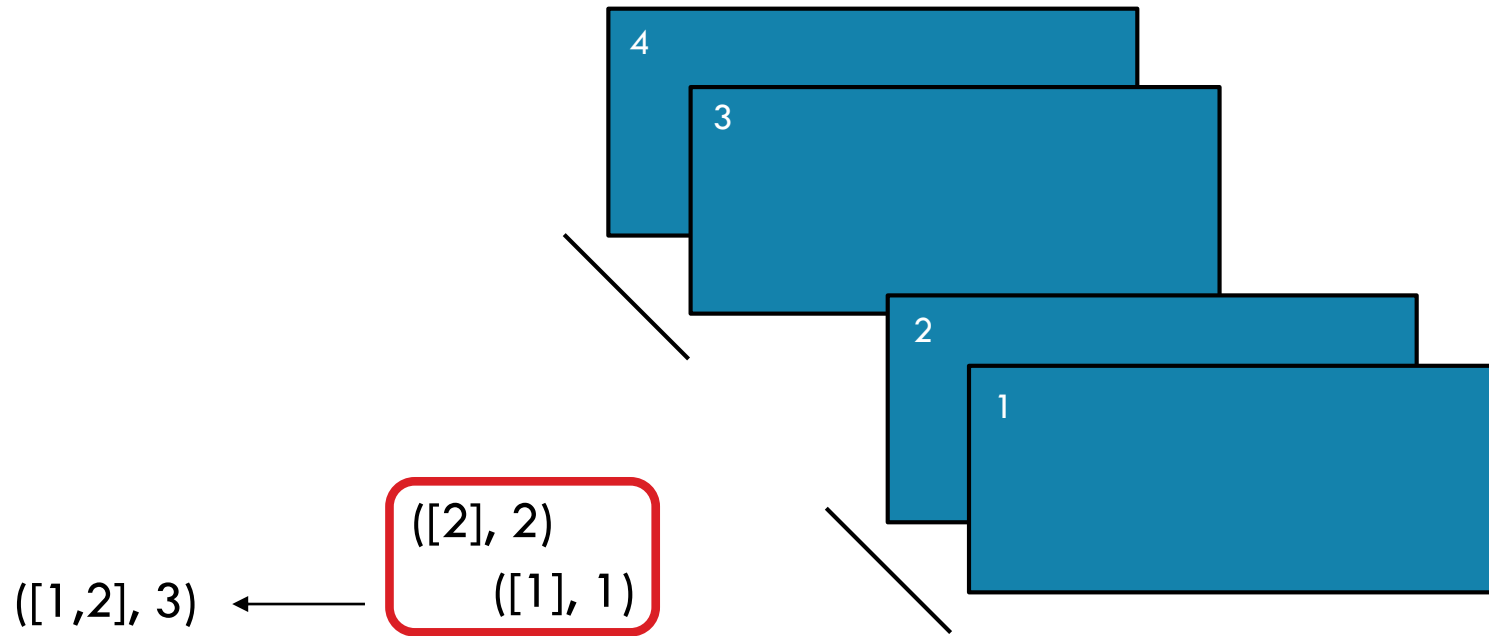
AGGREGATE



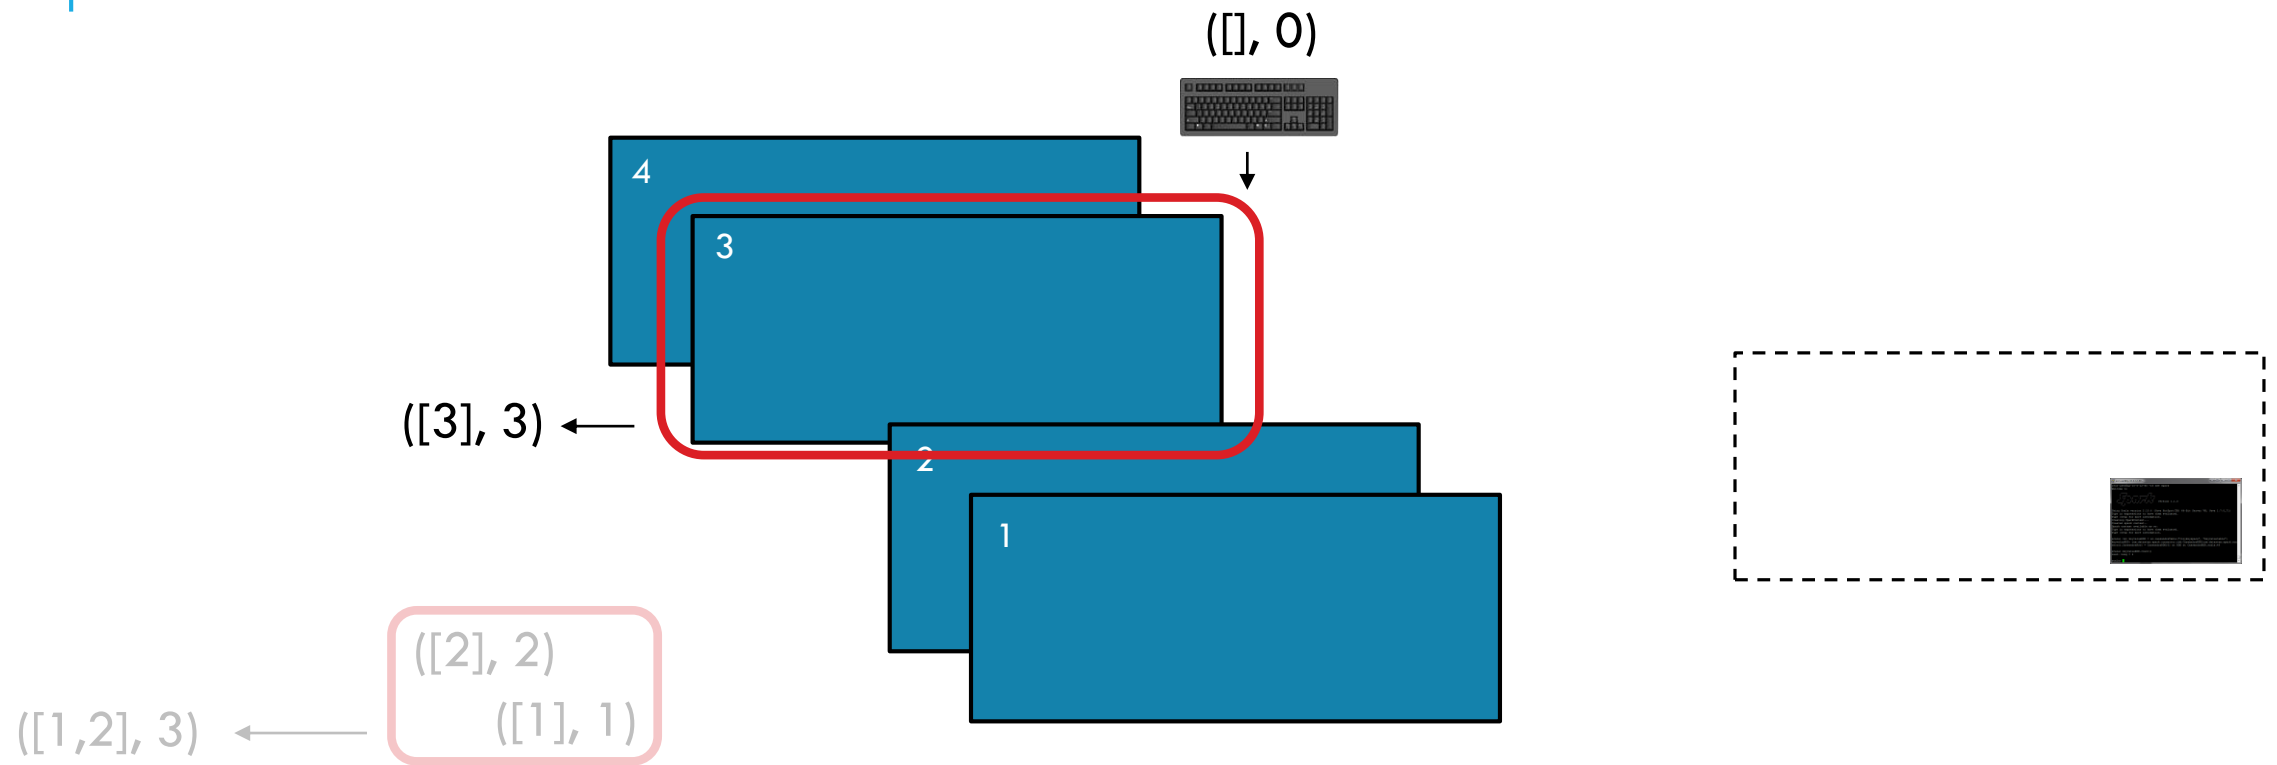
AGGREGATE




AGGREGATE




AGGREGATE




AGGREGATE



AGGREGATE



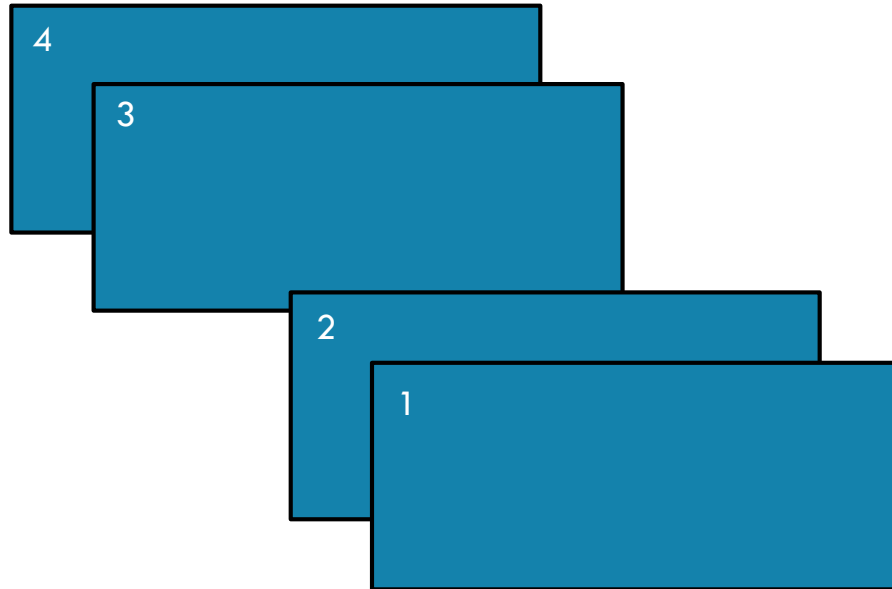
AGGREGATE




AGGREGATE

$([3,4], 7)$

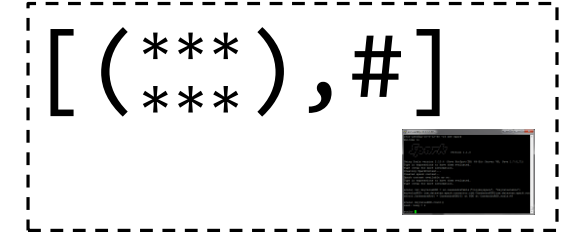
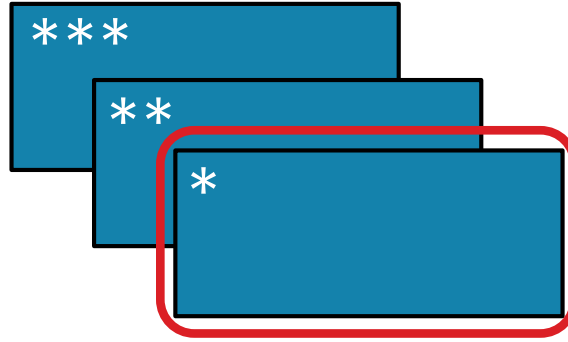
$([1,2], 3)$



AGGREGATE



AGGREGATE



`aggregate(identity, seqOp, combOp)`

Aggregate all the elements of the RDD by:

- applying a user function to combine elements with user-supplied objects,
- then combining those user-defined results via a second user function,
- and finally returning a result to the driver.



```
seqOp = lambda data, item: (data[0] + [item], data[1] + item)
combOp = lambda d1, d2: (d1[0] + d2[0], d1[1] + d2[1])
```

```
x = sc.parallelize([1,2,3,4])
y = x.aggregate([], 0), seqOp, combOp)
print(y)
```





x: [1, 2, 3, 4]

y: ([1, 2, 3, 4], 10)



AGGREGATE



`aggregate(identity, seqOp, combOp)`

Aggregate all the elements of the RDD by:

- applying a user function to combine elements with user-supplied objects,
- then combining those user-defined results via a second user function,
- and finally returning a result to the driver.

```
def seqOp = (data:(Array[Int], Int), item:Int) =>
  (data._1 :+ item, data._2 + item)

def combOp = (d1:(Array[Int], Int), d2:(Array[Int], Int)) =>
  (d1._1.union(d2._1), d1._2 + d2._2)

val x = sc.parallelize(Array(1,2,3,4))
val y = x.aggregate((Array[Int](), 0))(seqOp, combOp)

println(y)
```


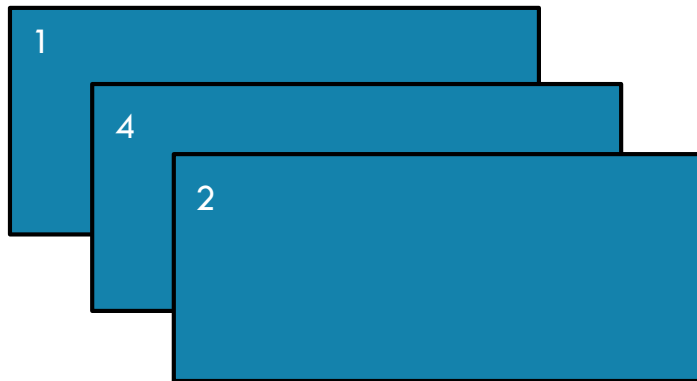


x: [1, 2, 3, 4]

y: (Array(3, 1, 2, 4),10)



MAX



MAX



`max()`

Return the maximum item in the RDD



```
x = sc.parallelize([2,4,1])  
y = x.max()  
  
print(x.collect())  
print(y)
```

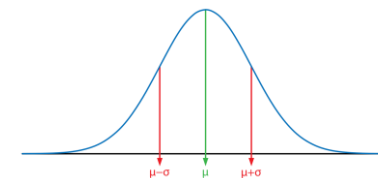


```
val x = sc.parallelize(Array(2,4,1))  
val y = x.max  
  
println(x.collect().mkString(", "))  
println(y)
```





`x:` [2, 4, 1]

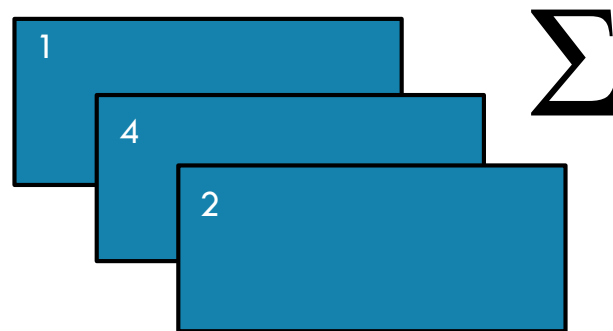
`y:` 4



SUM



SUM



`sum()`

Return the sum of the items in the RDD



```
x = sc.parallelize([2,4,1])
y = x.sum()

print(x.collect())
print(y)
```



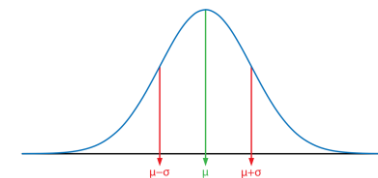
```
val x = sc.parallelize(Array(2,4,1))
val y = x.sum

println(x.collect().mkString(", "))
println(y)
```

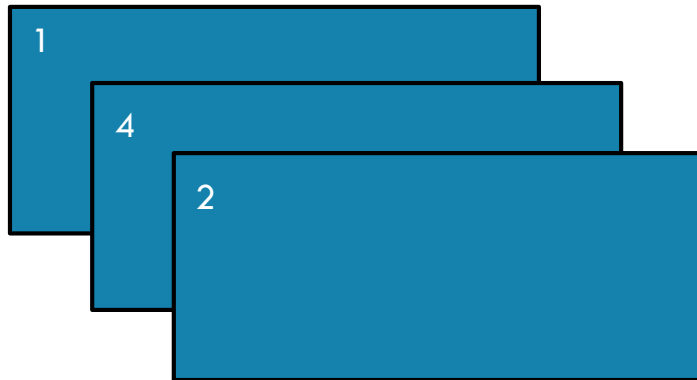


`x:` [2, 4, 1]

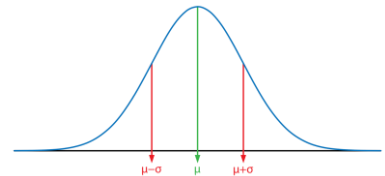
`y:` 7



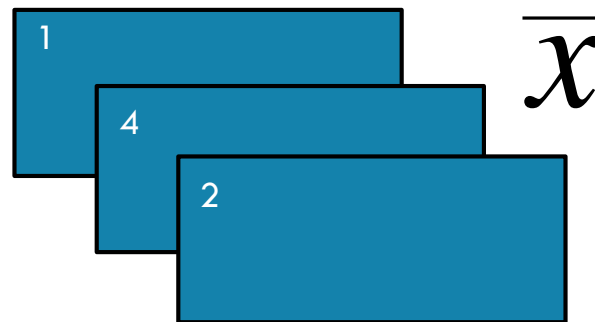
MEAN



2.33333333



MEAN



2.3333333

`mean()`

Return the mean of the items in the RDD



```
x = sc.parallelize([2,4,1])  
y = x.mean()  
  
print(x.collect())  
print(y)
```

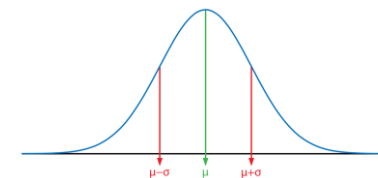


```
val x = sc.parallelize(Array(2,4,1))  
val y = x.mean  
  
println(x.collect().mkString(", "))  
println(y)
```

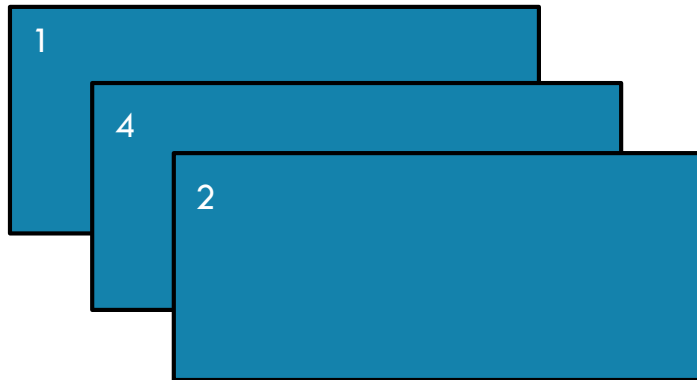


`x:` [2, 4, 1]


`y:` 2.3333333



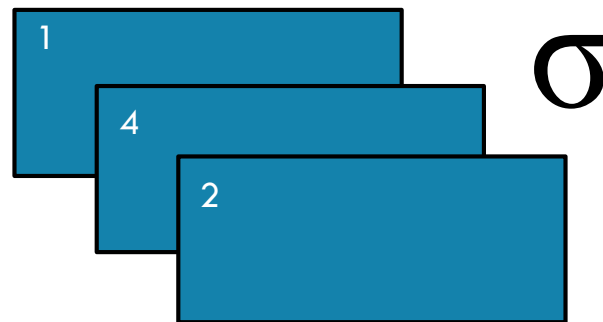
STDEV



1.2472191



STDEV



1.2472191

`stdev()`

Return the standard deviation of the items in the RDD



```
x = sc.parallelize([2,4,1])
y = x.stdev()

print(x.collect())
print(y)
```




```
val x = sc.parallelize(Array(2,4,1))
val y = x.stdev

println(x.collect().mkString(", "))
println(y)
```

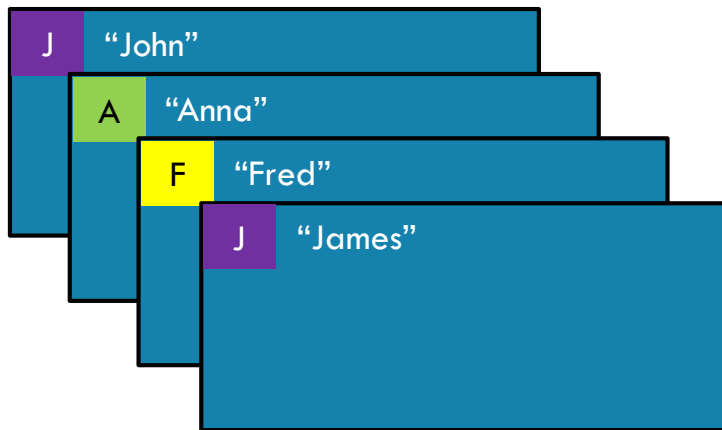


`x:` [2, 4, 1]

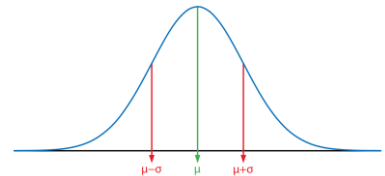
`y:` 1.2472191



COUNTBYKEY



{'A': 1, 'J': 2, 'F': 1}



COUNTBYKEY



countByKey()

Return a map of keys and counts of their occurrences in the RDD



```
x = sc.parallelize([('J', 'James'), ('F','Fred'),  
                  ('A','Anna'), ('J','John')])
```

```
y = x.countByKey()  
print(y)
```



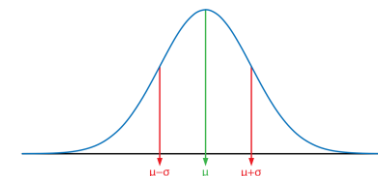
x: [('J', 'James'), ('F','Fred'),
 ('A','Anna'), ('J','John')]

y: {'A': 1, 'J': 2, 'F': 1}




```
val x = sc.parallelize(Array(('J',"James"),('F',"Fred"),  
                           ('A',"Anna"),('J',"John")))
```

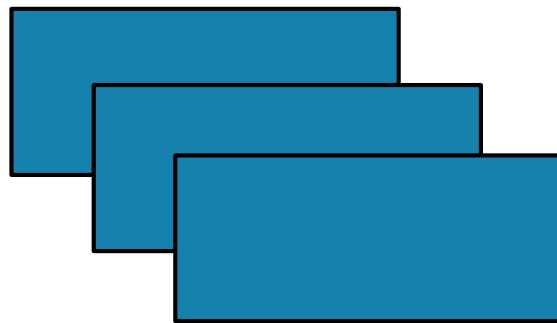
```
val y = x.countByKey()  
println(y)
```



SAVEASTEXTFILE



SAVEASTEXTFILE



`saveAsTextFile(path, compressionCodecClass=None)`

Save the RDD to the filesystem indicated in the path



```
dbutils.fs.rm("/temp/demo", True)
x = sc.parallelize([2,4,1])
x.saveAsTextFile("/temp/demo")
```

```
y = sc.textFile("/temp/demo")
print(y.collect())
```



```
dbutils.fs.rm("/temp/demo", true)
val x = sc.parallelize(Array(2,4,1))
x.saveAsTextFile("/temp/demo")
```

```
val y = sc.textFile("/temp/demo")
println(y.collect().mkString(", "))
```



`x:` [2, 4, 1]

`y:` [u'2', u'4', u'1']



LAB



Q&A

